

---

# Curvature-Calibrated Exploration: Matrix-Free GGN Widths and Conditional Regret Bounds

---

Anonymous Author  
Anonymous Institution

## Abstract

Curvature surrogates in differentiable-model bandits change exploration widths, but greater fidelity to full curvature need not improve online decisions. We use the full generalized Gauss–Newton (GGN)/Fisher operator as a reference, computing its widths with curvature–vector products and conjugate gradients (CG) without materializing a  $d \times d$  matrix. Rewards have conditionally sub-Gaussian noise; the fitted curvature is the Gaussian squared-loss quasi-likelihood. In the small-drift exact-current regime, a sharpened singular-value argument transfers the analysis to the monotone frozen-feature potential. Splitting its spectrum at rank  $r$  yields a full-dimensional regret bound in terms of the top- $r$  log determinant and the residual spectral mass  $\Delta_{T,r}$ , without assuming an exact tangent subspace. Exact rank is a special case; a supplied subspace also permits an  $r$ -dimensional implementation. A more general transfer theorem separates linearization, centering, feature and operator transfer, CG error, and realized width complexity. Predictable Welford and residual statistics instantiate its nonlinear certificates with  $O(d)$  additional path state. These are conditional differentiable-model results, not guarantees for unrestricted neural training. Experiments are reference-operator and failure-mode studies rather than a performance demonstration, and they show no uniform fidelity-based regret ordering in the tested regimes.

## 1 Introduction

Contextual bandits are a standard framework for sequential decision-making under partial feedback, and the reward predictor may be a differentiable parametric model or a pretrained representation composed with a reward head (Zhou et al., 2020; Riquelme et al., 2018). The central algorithmic question is how to build an exploration bonus that faithfully reflects epistemic uncertainty over such a model’s predictions. We analyze the Gaussian/squared-loss curvature model with sub-Gaussian noise and uniformly Lipschitz parameter gradients on a known certified region. The closed bounds also require small feature drift and a controlled spectral tail or an exact-rank special case; unrestricted neural training is not covered.

In the *white-box* setting—gradient and curvature–vector-product access to the model—the canonical bonus is the GGN/Fisher-linearized predictive variance

$$s_t^2(a) = \mathbf{g}_t(a)^\top \mathbf{C}_t^{-1} \mathbf{g}_t(a), \quad (1)$$

where  $\mathbf{g}_t(a) = \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}_t}(\mathbf{x}_t, a)$  is the gradient feature (Jacobian of the mean reward) and  $\mathbf{C}_t$  is the damped Gauss–Newton (GGN) curvature at the current estimate  $\boldsymbol{\theta}_t$ . Exact full-history GGN/Fisher widths are costly at large  $d$  or  $t$ : dense materialization uses  $O(d^2)$  state, while matrix-free CG replays history at every iteration. NeuralUCB retains the full neural-gradient Gram in its theoretical algorithm but diagonalizes it in experiments (Zhou et al., 2020); neural-linear and scalable Laplace methods instead use last-layer or structured curvature approximations (Riquelme et al., 2018; Ritter et al., 2018; Daxberger et al., 2021). Outside online bandits, implicit Jacobian/Fisher products and CG have already been used to compute full-network linearized predictive variances without forming the Gram matrix (Maddox et al., 2021). Each curvature surrogate changes the *spectrum* of  $\mathbf{C}_t$  and hence the value of  $s_t^2(a)$ . The spectral change does not by itself imply a change in the selected action, and therefore in regret; whether it does is an empirical question, which we make measurable and study directly.

Current-parameter relinearization asks a different scientific question from a historical-gradient Gram. The latter preserves the features that were predictable when data were collected and therefore supports a rank-one self-normalized confidence process. The former re-evaluates every historical example under the representation used for the next decision, but destroys that monotone update. Our analysis keeps the predictable frozen Gram for confidence and quantifies when its widths transfer to the current full-history GGN. This separation is why feature drift, rather than curvature fidelity alone, appears explicitly in the regret bound.

**We use full curvature as a reference operator in a falsification study.** Common-trajectory diagnostics isolate surrogate-induced width distortion, while independently executed policies test whether those distortions survive adaptive feedback. We do not attribute a uniform fidelity–regret ordering claim to prior work; instead we use it as a deliberately strong diagnostic null hypothesis. A controlled reversal rejects that hypothesis, whereas agreement does not establish superiority. Because executed policies follow different histories, their regret contrasts are descriptive rather than causal operator effects. We adapt the established matrix-free predictive-variance primitive to online contextual exploration: curvature–vector products (CVPs), available through automatic differentiation, let us solve  $\mathbf{C}_t \mathbf{u} = \mathbf{g}_t(a)$  by conjugate gradients (CG) *without forming  $\mathbf{C}_t$* : only  $O(d)$  additional curvature-state vectors (a constant number of CG vectors) are maintained and no  $d \times d$  curvature matrix or per-sample gradient vector is stored—though exact full-history curvature–vector products still require retaining or replaying the historical raw examples to recompute their Jacobian products, and model parameters, optimizer state, activations, replay data, and any frozen-feature parameter checkpoints are not counted in this  $O(d)$  curvature state. Because the damped GGN is symmetric positive definite, CG converges in exact arithmetic. All CG convergence and residual-to-energy statements below use exact arithmetic unless a verified rounding-error enclosure is supplied. Building on this, we propose **Curvature-Calibrated UCB (CC-UCB)**, a bandit algorithm whose exploration bonus is the predictive variance under the *uncompressed* GGN/Fisher metric, accessed entirely through implicit CG solves.

**The analysis tracks predictable certificate requirements for each approximation source and constructs them for the nonlinear path terms.** The Gaussian/squared-loss theorem (Theorem 1) handles a general predictable SPD curvature sequence with a one-sided transfer factor, a per-round CG factor, and a realized width complexity  $\Lambda_T^C$ . Its exact determinant identity accounts for endpoint growth and determinant-

destroying refreshes but does not itself bound either term. More importantly, the previously external nonlinear inputs are computable: Welford parameter moments and observed collection residuals upper-bound feature transfer and centering, while an analytic Taylor schedule supplies the confidence radius (Lemma 11 and Corollary 4). These schedules are conservative and need known smoothness constants and a certified optimizer residual. In the exact-current small-drift regime, the sharpened two-sided sandwich closes regret through the frozen potential, and a top- $r$  plus spectral-tail decomposition closes its information gain without a known subspace (Corollaries 2 and 3). A supplied exact-rank path is a computational special case, while a persistently excited growing window gives a separate conditional regret and width-solve sample-CVP accounting bound (Corollary 14 and Theorem 5). None of these results certifies unrestricted fine-tuning.

**Contributions.** (1) A full-dimensional, approximate-rank regret closure for exact current GGN under predictable small drift, using a sharpened frozen/current sandwich and a residual spectral-mass term; exact rank is a special case (Corollaries 2 and 3 and Lemma 7). (2) A general realized-width regret reduction with one-sided feature/operator transfer and explicit centering, linearization, and CG terms (Theorem 1 and Lemma 16). (3) Predictable  $O(d)$  additional-state nonlinear transfer and centering certificates (Lemma 11 and Corollary 4). (4) A conditional projected-Gram shortcut for a supplied tangent subspace and a separate growing-window bound with its burn-in and certificate terms exposed (Corollary 14 and Theorem 5). (5) Matrix-free full-GGN/Fisher widths through residual-checked CG/CVP solves, together with executed nonlinear, operator-mechanism, baseline, and systems audits that separate policy outcomes from post-hoc diagnostics (Section 5 and Proposition 31).

## 2 Problem Setup

We present the problem setup for the *Gaussian/squared-loss* instantiation, which is the setting analyzed in Theorem 1. An extension to exponential-family likelihoods is discussed in Appendix H.

**Protocol.** At each round  $t = 1, \dots, T$  the learner observes a context  $\mathbf{x}_t \in \mathcal{X}$  and a finite action set  $\mathcal{A}_t$ . It selects  $a_t \in \mathcal{A}_t$  and observes a stochastic reward

$$r_t = \mu^*(\mathbf{x}_t, a_t) + \eta_t, \quad \eta_t \text{ is } \sigma\text{-sub-Gaussian}, \quad (2)$$

where  $\mu^* : \mathcal{X} \times \mathcal{A} \rightarrow \mathbb{R}$  is the unknown mean reward. The goal is to minimize cumulative regret  $R_T = \sum_{t=1}^T [\mu^*(\mathbf{x}_t, a_t^*) - \mu^*(\mathbf{x}_t, a_t)]$ , where  $a_t^* =$

$\arg \max_{a \in \mathcal{A}_t} \mu^*(\mathbf{x}_t, a)$ . We assume  $|\mathcal{A}_t| \leq A_{\max}$  deterministically for all  $t$ .

### Reward model and parameterization convention.

We fix a pretrained initialization  $\theta_0 \in \mathbb{R}^d$  and parameterize the model by the *displacement*  $\theta \in \mathbb{R}^d$  from it: the deployed predictor is  $\mu_{\theta_0+\theta}(\mathbf{x}, a)$ , and we abbreviate  $\mu_\theta(\mathbf{x}, a) := \mu_{\theta_0+\theta}(\mathbf{x}, a)$  throughout. The initial displacement is  $\theta_1 = \mathbf{0}$  (the deployed model starts at the pretrained weights  $\theta_0$ ). With this convention the ridge penalty below is a penalty on the *displacement from the pretrained weights*, not a zero-centered prior on the raw weights, and the realizability radius  $S$  bounds  $\|\theta^*\| = \|\theta_{\text{orig}}^* - \theta_0\|$ , the distance of the true parameter from the pretrained initialization. Define the regularized loss at round  $t$ :

$$\mathcal{L}_t(\theta) := \frac{1}{2\sigma^2} \sum_{s=1}^{t-1} (r_s - \mu_\theta(\mathbf{x}_s, a_s))^2 + \frac{\lambda}{2} \|\theta\|^2. \quad (3)$$

We write  $\theta_t$  for the parameter returned by the warm-started optimizer on  $\mathcal{L}_t$ . Let  $\mathcal{D}_{t-1} := \{(\mathbf{x}_s, a_s, r_s) : s < t\}$  and  $N_t := |\mathcal{D}_{t-1}| = t-1$ . Equation (3) is a *full-history* objective. The curvature replay set  $\mathcal{S}_t \subseteq \{1, \dots, t-1\}$  and its size  $m_t := |\mathcal{S}_t|$  affect  $\mathcal{C}_t$ , but do not replace  $\mathcal{D}_{t-1}$  in (3). Thus a generic full-batch evaluation of  $\nabla \mathcal{L}_t$  replays  $N_t$  observations even when  $m_t \ll N_t$ . In convex settings (e.g. a frozen backbone with a linear head), under exact optimization this coincides with the global minimizer; in nonconvex settings  $\theta_t$  may be a local minimizer or approximate stationary point. The gap between  $\theta_t$  and an idealized global minimizer is absorbed by the action-scaled centering certificate (Assumption 4).

**Gradient features and curvature matrices.** Define the *gradient feature* (Jacobian of the mean reward) at the current parameter estimate:

$$\mathbf{g}_t(\mathbf{x}, a) := \nabla_\theta \mu_{\theta_t}(\mathbf{x}, a) \in \mathbb{R}^d. \quad (4)$$

The algorithm uses the *relinearized curvature matrix*

$$\mathbf{C}_t = \lambda \mathbf{I} + \frac{1}{\sigma^2} \sum_{s=1}^{t-1} \mathbf{g}_t(\mathbf{x}_s, a_s) \mathbf{g}_t(\mathbf{x}_s, a_s)^\top, \quad (5)$$

the *damped GGN curvature* for the Gaussian/squared-loss model at the current parameter  $\theta_t$ . (For this likelihood, the GGN coincides with the true Fisher information; see the remark below.) While  $\mathbf{C}_t$  is the natural object for computing predictive variances at  $\theta_t$ , it re-evaluates all past Jacobians at  $\theta_t$  and therefore does *not* admit a sequential rank-one update—all summands change when  $\theta_t$  changes.

For the regret analysis we introduce the *frozen-feature*

*Gram matrix*:

$$\bar{\mathbf{C}}_t := \lambda \mathbf{I} + \frac{1}{\sigma^2} \sum_{s=1}^{t-1} \mathbf{g}_s(\mathbf{x}_s, a_s) \mathbf{g}_s(\mathbf{x}_s, a_s)^\top, \quad (6)$$

where each  $\mathbf{g}_s(\mathbf{x}_s, a_s) = \nabla_\theta \mu_{\theta_s}(\mathbf{x}_s, a_s)$  is evaluated at the parameter  $\theta_s$  that was current when the sample was collected. This matrix satisfies the recursive update

$$\bar{\mathbf{C}}_{t+1} = \bar{\mathbf{C}}_t + \sigma^{-2} \mathbf{g}_t(\mathbf{x}_t, a_t) \mathbf{g}_t(\mathbf{x}_t, a_t)^\top, \quad (7)$$

and, critically, each realized design vector  $\mathbf{g}_s(\mathbf{x}_s, a_s)$  is measurable with respect to the pre-reward sigma-algebra  $\mathcal{G}_s = \mathcal{H}_s \vee \sigma(a_s)$  of (78), which is the filtration used in the self-normalized bound. Both the additive update (7) and this measurability are structural requirements of the self-normalized martingale bound and the elliptic potential lemma invoked in Section 4. The algorithm computes UCB widths using the algorithmic curvature  $\mathcal{C}_t$  (via CG; the exact case is  $\mathcal{C}_t = \mathbf{C}_t$ ), but the confidence analysis is conducted in terms of  $\bar{\mathbf{C}}_t$ ; the one-sided optimism-transfer factor of Assumption 5, certified by Lemma 16, provides the bridge.

## 3 Method: Curvature-Calibrated Exploration

### 3.1 CC-UCB Algorithm

For a predictable SPD oracle  $\mathcal{C}_t \succeq \lambda \mathbf{I}$ , set  $\omega_t = \bar{\beta}_t + \bar{\psi}_t$  for the original nonlinear center. Algorithm 1 is the analyzed full-enumeration policy. The same fixed curvature operator is reused across separate per-action CG solves; a solution for one action is not reused as the solution for another action.

---

#### Algorithm 1 Curvature-Calibrated UCB (CC-UCB)

---

- 1: Observe context  $\mathbf{x}_t$  and finite action set  $\mathcal{A}_t$ .
  - 2: Construct a fixed predictable SPD oracle  $\mathbf{v} \mapsto \mathcal{C}_t \mathbf{v}$  with  $\mathcal{C}_t \succeq \lambda \mathbf{I}$ .
  - 3: **for** each  $a \in \mathcal{A}_t$  **do**
  - 4:   Compute  $\hat{\mu}_t(a) = \mu_{\theta_t}(\mathbf{x}_t, a)$  and  $\mathbf{g}_t(a) = \nabla_\theta \mu_{\theta_t}(\mathbf{x}_t, a)$ .
  - 5:   If  $\mathbf{g}_t(a) \neq \mathbf{0}$ , solve  $\mathcal{C}_t \mathbf{u} = \mathbf{g}_t(a)$  separately by CG until the certified tolerance holds; otherwise set  $\tilde{\mathbf{u}} = \mathbf{0}$ .
  - 6:   Set  $\tilde{s}_t^2(a) = \mathbf{g}_t(a)^\top \tilde{\mathbf{u}}$ .
  - 7: **end for**
  - 8: Select and play  $a_t \in \arg \max_{a \in \mathcal{A}_t} \{\hat{\mu}_t(a) + \omega_t \sqrt{u_t(a)/(1 - \bar{\varepsilon}_t)} \sqrt{\tilde{s}_t^2(a)}\}$ .
  - 9: Observe  $r_t$ , append  $(\mathbf{x}_t, a_t, r_t)$ , and update  $\theta_{t+1}$ .
- 

The oracle is accessed only through matrix–vector products. At each round the agent computes, for every candidate action  $a$ , the gradient feature  $\mathbf{g}_t(a) =$

$\nabla_{\theta} \mu_{\theta_t}(\mathbf{x}_t, a)$  and the *curvature-calibrated variance proxy*, the generic algorithmic width,

$$s_{\mathcal{C}_t}^2(a) = \mathbf{g}_t(a)^\top \mathcal{C}_t^{-1} \mathbf{g}_t(a). \quad (8)$$

Under the GGN/Fisher-linearized Gaussian approximation with covariance  $\mathcal{C}_t^{-1}$ , the quantity  $s_{\mathcal{C}_t}^2(a)$  is the metric-dependent predictive variance of the linearized reward at  $(\mathbf{x}_t, a)$ . The *exact* full-history instantiation of CC-UCB takes  $\mathcal{C}_t = \mathbf{C}_t$ , the relinearized curvature (5); weighted replay, subsampling, sliding windows, periodic refreshes, and randomized sketches are special implementations of  $\mathcal{C}_t$  (§A.3), covered by the theory only when the corresponding one-sided transfer factor is certified.

Because  $\mathcal{C}_t \succ 0$  by construction, conjugate gradients (CG) is guaranteed to converge on the system  $\mathcal{C}_t \mathbf{u} = \mathbf{g}_t(a)$ , provided the operator is held *fixed* across all iterations of the solve (the randomized sketch/subsample is drawn into  $\Omega_t$  before action selection; §2). For each action with  $\mathbf{g}_t(a) \neq \mathbf{0}$  we run CG to the per-action relative energy-norm error

$$\varepsilon_{t,a} := \frac{\|\mathbf{u}_{t,a} - \tilde{\mathbf{u}}_{t,a}\|_{\mathcal{C}_t}}{\|\mathbf{u}_{t,a}\|_{\mathcal{C}_t}}, \quad \mathbf{u}_{t,a} := \mathcal{C}_t^{-1} \mathbf{g}_t(a), \quad (9)$$

yielding the computable proxy  $\tilde{s}_t^2(a) = \mathbf{g}_t(a)^\top \tilde{\mathbf{u}}_{t,a}$ ; if  $\mathbf{g}_t(a) = \mathbf{0}$  we set  $\tilde{s}_t^2(a) = 0$  and skip the solve (leaving the residual ratio  $\|r_I\|/\|\mathbf{g}\|$  undefined), and by convention put  $\varepsilon_{t,a} := 0$  in this case, so that  $\varepsilon_{t,a}$  is defined for every  $a \in \mathcal{A}_t$  and (67) holds trivially there. Lemma 22 shows  $\tilde{s}_t^2(a)$  approximates  $s_{\mathcal{C}_t}^2(a)$  within a multiplicative  $(1 \pm \varepsilon_{t,a})$  factor. To ensure the UCB bonus *upper-bounds* the confidence width  $\omega_t \tilde{s}_t(a)$  despite the approximation, the algorithm inflates the width by  $\sqrt{u_t(a)/(1 - \bar{\varepsilon}_t)}$  (line 8), where  $\bar{\varepsilon}_t \in (0, 1)$  is a *per-round* known tolerance enforced uniformly over all actions,

$$\max_{a \in \mathcal{A}_t} \varepsilon_{t,a} \leq \bar{\varepsilon}_t < 1, \quad (10)$$

and  $u_t(a) \geq 1$  is the one-sided transfer factor (Assumption 5).

## 4 Theoretical Analysis

We state a regret bound for CC-UCB (Algorithm 1, with full action enumeration  $K=|\mathcal{A}_t|$ ) under the Gaussian/squared-loss setup of Section 2 and Assumptions 1 to 5. The analysis does not constitute a full exponential-family or generalized linear bandit proof; the extension in Appendix H is a curvature-definition extension only. The analysis handles a *general predictable SPD algorithmic curvature sequence*  $\mathcal{C}_t$  (45). Confidence is proved for the frozen-feature Gram  $\tilde{\mathbf{C}}_t$  (Eq. 6), which has the sequential-update and predictability properties required by the self-normalized

martingale bound; the *one-sided* optimism-transfer factor  $u_t(a)$  (Assumption 5) bridges from  $\tilde{\mathbf{C}}_t$  to  $\mathcal{C}_t$  in the exploration direction only (no lower drift sandwich); the CG approximation enters through the per-round factor  $\alpha_t$ ; and the nonmonotonicity of  $\mathcal{C}_t$  is retained in the realized width complexity. Appendix Appendix A.5 gives an exact determinant bookkeeping identity, not a primitive sublinear bound.

The theorem is a modular transfer result. The path statistics below instantiate its nonlinear inputs from pre-action quantities; smallness still requires a restricted parameter path and is not claimed for unrestricted fine-tuning.

**Core conditions and links.** Assume conditionally  $\sigma$ -sub-Gaussian noise,  $\|\mathbf{g}_t(\mathbf{x}_s, a)\| \leq G$ , realizability by  $\|\theta^*\| \leq S$ , and a local expansion with round- $t$  remainder at most  $\varepsilon_{\text{lin}}(t)$ . For the original center assume the action-scaled discrepancy is at most  $\bar{\psi}_t \tilde{s}_t(a)$ . The operator  $\mathcal{C}_t$  is predictable, SPD, satisfies  $\mathcal{C}_t \succeq \lambda \mathbf{I}$  and  $\mathcal{C}_1 = \lambda \mathbf{I}$ , and remains fixed throughout every CG solve. Define

$$\begin{aligned} \tilde{s}_t^2(a) &= \mathbf{g}_t(a)^\top \tilde{\mathbf{C}}_t^{-1} \mathbf{g}_t(a), \quad s_{\mathcal{C}_t}^2(a) = \mathbf{g}_t(a)^\top \mathcal{C}_t^{-1} \mathbf{g}_t(a), \\ \tilde{s}_t^2(a) &= \mathbf{g}_t(a)^\top \tilde{\mathbf{u}}_t(a). \end{aligned} \quad (11)$$

The analysis requires only the one-sided transfer and uniform CG relations

$$\begin{aligned} \tilde{s}_t^2(a) &\leq u_t(a) s_{\mathcal{C}_t}^2(a), \\ (1 - \bar{\varepsilon}_t) s_{\mathcal{C}_t}^2(a) &\leq \tilde{s}_t^2(a) \leq (1 + \bar{\varepsilon}_t) s_{\mathcal{C}_t}^2(a), \end{aligned} \quad (12)$$

for every candidate action, where  $u_t(a) \geq 1$ ,  $\bar{\varepsilon}_t < 1$ , and  $\alpha_t = \sqrt{(1 + \bar{\varepsilon}_t)/(1 - \bar{\varepsilon}_t)}$ . Sufficient feature-drift and approximate-operator constructions are given in Lemma 16 and Theorem 29; no lower spectral transfer is needed.

**Dynamic width complexity.** The realized quantity used below is

$$\Lambda_T^{\mathcal{C}} := \sum_{t=1}^T \log(1 + \sigma^{-2} s_{\mathcal{C}_t}^2(a_t)). \quad (13)$$

The bounded-width cap gives

$$\sum_{t=1}^T s_{\mathcal{C}_t}^2(a_t) \leq (\sigma^2 + G^2/\lambda) \Lambda_T^{\mathcal{C}}. \quad (14)$$

For a general predictable nonmonotone SPD sequence, we do not derive a primitive sublinear bound on  $\Lambda_T^{\mathcal{C}}$ . Appendix Appendix A.5 gives an exact determinant bookkeeping identity, but its upper bound is informative only when both the terminal log determinant and determinant-decreasing refresh charge are independently controlled, for example under the rank and spectral-floor hypotheses of Lemma 20.

#### 4.1 Regret Bound

Let  $\mathcal{E}_{\text{conf}}$  denote the simultaneous self-normalized confidence event and let  $\mathcal{E}_{\text{cert}}$  collect the predictable linearization, centering, transfer, condition-number, CG, and randomized-operator certificates. The appendix defines each component and its failure allocation; write their total failure probability as  $\delta_{\text{cert}}$ .

**Theorem 1** (Dynamic transfer regret bound for CC-UCB (Gaussian/squared-loss)). *Under the conditions above and Assumptions 1 to 5, suppose  $|\mathcal{A}_t| \leq A_{\max}$ , the optimal action is present, and Algorithm 1 enumerates all actions using a predictable schedule  $\bar{\beta}_t \geq \beta_t$  and uniform per-round CG tolerance. Let  $E_T = \sum_{t=1}^T \varepsilon_{\text{lin}}(t)$ ,  $\omega_t = \bar{\beta}_t + \bar{\psi}_t$ ,  $\alpha_t = \sqrt{(1 + \bar{\varepsilon}_t)/(1 - \bar{\varepsilon}_t)}$ , and  $S_T = \sum_{t=1}^T \alpha_t^2 u_t(a_t) \omega_t^2$ . On  $\mathcal{E}_{\text{conf}} \cap \mathcal{E}_{\text{cert}}$ ,*

$$R_T \leq 2\sqrt{(\sigma^2 + G^2/\lambda) \Lambda_T^C S_T} + 2E_T, \quad (15)$$

with probability at least  $1 - \delta - \delta_{\text{cert}}$ . Deterministic or almost-sure certificates contribute zero to  $\delta_{\text{cert}}$ . For the corrected center, the same statement holds without Assumption 4 after replacing  $\omega_t$  by  $\bar{\beta}_t$ .

**Corollary 2** (Exact-current regret through the frozen potential). *Adopt Theorem 1 with exact current curvature  $\mathcal{C}_t = \mathbf{C}_t$ . Suppose a predictable certificate satisfies  $\chi_t \leq \bar{\chi}_t < 1$  for every  $t$ , and use the valid transfer factor  $u_t = (1 + \bar{\chi}_t)^2$ . Then, on the event of Theorem 1,*

$$R_T \leq 2\left\{(\sigma^2 + G^2/\lambda) \gamma_T \times \sum_{t=1}^T \alpha_t^2 \omega_t^2 \left(\frac{1 + \bar{\chi}_t}{1 - \bar{\chi}_t}\right)^2\right\}^{1/2} + 2E_T. \quad (16)$$

For the corrected center of Corollary 8, the same statement holds without Assumption 4 after replacing  $\omega_t$  by  $\bar{\beta}_t$ . This removes the nonmonotone dynamic-width term only in the small-drift, exact-current regime.

*Proof.* The sharpened sandwich in Corollary 17 and Loewner inversion give

$$s_t(a) \leq \frac{\bar{s}_t(a)}{1 - \bar{\chi}_t}.$$

The per-round proof of Theorem 1, together with  $\sqrt{u_t} = 1 + \bar{\chi}_t$ , therefore gives

$$\text{regret}_t \leq 2\alpha_t \omega_t \frac{1 + \bar{\chi}_t}{1 - \bar{\chi}_t} \bar{s}_t(a_t) + 2\varepsilon_{\text{lin}}(t).$$

Sum and apply Cauchy–Schwarz once. The frozen rank-one update and the usual elliptic-potential inequality give  $\sum_t \bar{s}_t^2(a_t) \leq (\sigma^2 + G^2/\lambda) \gamma_T$ , proving the display. The corrected-center claim follows from the prediction identity in Corollary 8.  $\square$

For an integer  $0 \leq r \leq d$ , define

$$\begin{aligned} \Delta_{T,r} &:= \sum_{i>r} \nu_i, & r_T &:= \min\{r, T\}, \\ \Gamma_{T,r}^{\text{tail}} &:= r_T \log\left(1 + \frac{TG^2}{r_T \lambda \sigma^2}\right) + \frac{\Delta_{T,r}}{\lambda}, \end{aligned} \quad (17)$$

where  $\nu_1 \geq \dots \geq \nu_d \geq 0$  are the eigenvalues of  $\bar{\mathbf{C}}_{T+1} - \lambda \mathbf{I}$  and the first term is defined to be zero when  $r_T = 0$ . Lemma 7 proves  $\gamma_T \leq \Gamma_{T,r}^{\text{tail}}$ .

**Corollary 3** (Approximate-rank exact-current GGN). *Under Corollary 2, for every fixed  $r \in \{0, \dots, d\}$ ,*

$$R_T \leq 2\left\{(\sigma^2 + G^2/\lambda) \Gamma_{T,r}^{\text{tail}} \times \sum_{t=1}^T \alpha_t^2 \omega_t^2 \left(\frac{1 + \bar{\chi}_t}{1 - \bar{\chi}_t}\right)^2\right\}^{1/2} + 2E_T. \quad (18)$$

*This is a realized-complexity statement: the policy may retain the existing observable  $\hat{\gamma}_{t-1}$  schedule, but a terminal realized  $\Delta_{T,r}$  is not used in any action-time radius.*

Operationally, if before round  $t$  a supplied deterministic or  $\mathcal{H}_t$ -measurable envelope  $\bar{\Delta}_{t-1,r}$  upper-bounds the spectral tail of  $\bar{\mathbf{C}}_t - \lambda \mathbf{I}$ , define

$$\begin{aligned} r_{t-1} &:= \min\{r, t-1\}, \\ \bar{\Gamma}_{t-1,r}^{\text{tail}} &:= r_{t-1} \log\left(1 + \frac{(t-1)G^2}{r_{t-1} \lambda \sigma^2}\right) + \frac{\bar{\Delta}_{t-1,r}}{\lambda}. \end{aligned}$$

again taking the first term as zero when  $r_{t-1} = 0$ . Replacing  $\gamma_{t-1}$  in the constructive confidence radius by this predictable upper bound is valid. In that version, (18) holds with  $\Gamma_{T,r}^{\text{tail}}$  replaced by any supplied terminal envelope  $\bar{\Gamma}_{T,r}^{\text{tail}}$  and with  $\omega_t$  formed from the resulting radius.

No exact tangent subspace is assumed. A small spectral tail can contain decision-relevant directions and does not supply their basis, so full-dimensional CG may still be necessary. Exact rank is recovered when  $\Delta_{T,r} = 0$ ; the supplied-subspace result Corollary 14 is then an additional computational special case. The corrected-center version again replaces  $\omega_t$  by  $\bar{\beta}_t$ .

*Proof.* Apply Lemma 7 to substitute  $\gamma_T \leq \Gamma_{T,r}^{\text{tail}}$  in (16). The prefix version of the same lemma makes  $\bar{\Gamma}_{t-1,r}^{\text{tail}}$  a predictable upper bound on  $\gamma_{t-1}$ , so Lemma 24 and Corollary 25 justify the operational radius. No eigenspace is used by either argument.  $\square$

In particular, the bound is sublinear whenever  $\Lambda_T^C S_T = o(T^2)$  and  $E_T = o(T)$ .

The exact-curvature instantiation is  $\mathcal{C}_t = \mathbf{C}_t$  with  $u_t = (1 + \bar{\chi}_t)^2$  from Lemma 16; approximate operators are covered through Theorem 29. The proof is in Appendix D.

**Proof sketch.** Self-normalized confidence plus the centering and linearization terms gives  $|\mu^*(a) - \mu_{\theta_t}(a)| \leq \omega_t \bar{s}_t(a) + \varepsilon_{\text{lin}}(t)$ . The one-sided relation and CG lower bound make the computed bonus optimistic, so instantaneous regret is at most  $2w_t(a_t) + 2\varepsilon_{\text{lin}}(t)$ . The CG upper bound gives  $w_t(a_t) \leq \alpha_t \omega_t \sqrt{u_t(a_t) s_{C,t}(a_t)}$ . Summing, applying Cauchy–Schwarz only once, and using (14) yields (15); hence time-varying factors remain inside  $S_T$  rather than being replaced by horizon maxima.

The corrected center replaces  $\omega_t$  by  $\bar{\beta}_t$  and removes the centering certificate, but requires exact or certified frozen-feature access; the formal statement and implementation discussion are in Corollary 8.

## 4.2 Online nonlinear certificates

At the start of round  $t$ , let  $n_t = t - 1$ , and let  $\bar{\theta}_{<t}$  and  $J_{<t} = \sum_{s<t} \|\theta_s - \bar{\theta}_{<t}\|^2$  be maintained by a vector Welford update. With collection residuals  $c_s = \mu_{\theta_s}(x_s, a_s) - r_s$ , define

$$Q_t = J_{<t} + n_t \|\theta_t - \bar{\theta}_{<t}\|^2, \quad R_t^{\text{col}} = \sum_{s<t} c_s^2. \quad (19)$$

Both are pre-action quantities and  $Q_t = \sum_{s<t} \|\theta_t - \theta_s\|^2$  exactly. If prediction gradients are  $L_g$ -Lipschitz on the certified region, set

$$\bar{\chi}_t = \frac{L_g \sqrt{Q_t}}{\sigma \sqrt{\lambda}}, \quad u_t = \kappa_{+,t} (1 + \bar{\chi}_t)^2, \quad (20)$$

where  $\kappa_{+,t} = 1$  for exact full curvature and for an unrescaled current-parameter subset, and otherwise is a supplied one-sided operator certificate. If  $\|\nabla \mathcal{L}_t(\theta_t)\| \leq \zeta_t$ , an observable centering schedule is

$$\begin{aligned} \bar{\psi}_t &= \frac{\zeta_t + \bar{M}_t}{\sqrt{\lambda}}, \\ \bar{M}_t &= \sigma^{-2} \left[ L_g \sqrt{R_t^{\text{col}} Q_t} + \frac{3}{2} G L_g Q_t + \frac{1}{2} L_g^2 Q_t^{3/2} \right]. \end{aligned} \quad (21)$$

For a known radius- $R$  trust region, the final numerator term improves to  $L_g^2 R Q_t$ . Finally, for an  $L_\mu$ -Lipschitz prediction gradient define  $\bar{\epsilon}_t^{\text{lin}} = \frac{L_\mu}{2} (S + \|\theta_t\|)^2$ ,  $\bar{F}_t = \sum_{s<t} (\bar{\epsilon}_s^{\text{lin}})^2$ , and

$$\begin{aligned} \bar{\beta}_t &= \sqrt{\hat{\gamma}_{t-1} + 2 \log(1/\delta)} + \sqrt{\lambda} S + \sigma^{-1} \sqrt{\bar{F}_t}, \\ \hat{\gamma}_{t-1} &= \sum_{s<t} \log(1 + q_s), \quad q_s = \frac{u_s(a_s) \tilde{s}_s^2(a_s)}{\sigma^2 (1 - \bar{\varepsilon}_s)}. \end{aligned} \quad (22)$$

**Corollary 4** (Operational CC-UCB). *Under the structural hypotheses of Theorem 1, with its certificate inputs instantiated by valid smoothness constants, the optimizer residual bound  $\zeta_t$ , uniform CG certificates, and a valid  $\kappa_{+,t}$ , (19)–(22) are predictable and satisfy every*

*nonlinear input of the theorem. Thus Theorem 1 holds with  $E_T$  replaced by  $\bar{E}_T = \sum_{t \leq T} \bar{\epsilon}_t^{\text{lin}}$  and the observable CG upper bound on  $\Lambda_T^C$  from Corollary 21.*

*Proof.* Lemma 11 proves  $\bar{\chi}_t \geq \chi_t$ ,  $\bar{\psi}_t \geq \psi_t$ ,  $\bar{F}_t \geq F_t$ , and  $\bar{E}_T \geq E_T$ . The one-sided operator relation and Corollary 26 give  $u_t$  and  $\hat{\gamma}_{t-1} \geq \gamma_{t-1}$ , hence  $\bar{\beta}_t \geq \beta_t$ . All quantities use the summary through round  $t - 1$  and the already formed  $\theta_t$ ; uniform candidate-action CG accuracy is checked before selection. Substitution in Theorem 1 and Corollary 21 proves the claim.  $\square$

The proof, filtration-order pseudocode, and a scalar-tanh specialization with  $G \leq B_\phi$  and  $L_\mu = L_g \leq 4B_\phi^2/(3\sqrt{3})$  are in Lemma 11 and Corollary 15.

For an abstract near-linearity scale  $W$ , suppose  $\|\theta_t\|, \|\theta^*\| \leq R$ ,  $L_\mu \leq c_\mu/\sqrt{W}$ ,  $L_g \leq c_g/\sqrt{W}$ ,  $\zeta_t \leq \zeta_0/\sqrt{W}$ , and  $R_t^{\text{col}} \leq C_R t P_T$  for a supplied  $P_T \geq 1$ . Then

$$\begin{aligned} \bar{E}_T &\leq \frac{2c_\mu R^2 T}{\sqrt{W}}, \quad \bar{F}_{T+1} \leq \frac{4c_\mu^2 R^4 T}{W}, \\ \bar{\chi}_t &\leq \frac{2c_g R \sqrt{t-1}}{\sigma \sqrt{\lambda W}}. \end{aligned} \quad (23)$$

and Corollary 12 gives an explicit trust-region bound on  $\bar{\psi}_t$ . If a supplied fixed active subspace gives rank bound  $r \geq 1$ , define

$$\Gamma_{T,r} = r_T \log \left( 1 + \frac{T G^2}{r_T \lambda \sigma^2} \right), \quad r_T = \min\{r, T\}. \quad (24)$$

and set  $\Gamma_{0,r} := 0$  for the round-one confidence schedule. The bound  $r$  must be known before selection; a terminal realized rank is only post-hoc. Then  $\gamma_T \leq \Gamma_{T,r}$  and Corollary 14 gives the fully explicit conditional projected-Gram bound

$$R_T \leq 2\alpha_I \sqrt{\rho_+/\rho_-} \Omega_{T,r} \sqrt{(\sigma^2 + G^2/\lambda) T \Gamma_{T,r}} + 2\bar{E}_T, \quad (25)$$

where  $x_T = 2c_g R \sqrt{T-1}/(\sigma \sqrt{\lambda W})$ ,  $\rho_- = (1 - x_T)^2$ ,  $\rho_+ = (1 + x_T)^2$ , and  $\Omega_{T,r}$  is displayed with every  $T, W, r, P_T, \lambda, \sigma$  dependence in the appendix. At fixed  $r$  and problem constants,  $W = \Omega(T^2 P_T)$  makes the path inflation bounded and gives  $R_T = O(r \sqrt{T} \log T) + O(T/\sqrt{W})$ . This is a low-rank tangent-space or parameter-efficient adaptation regime, not generic full-network training.

**Theorem 5** (Conditional growing-window bound under fixed-subspace excitation). *Adopt the hypotheses, events, and valid predictable schedules of Theorem 1, let  $m_t = \min\{t-1, \lceil t^q \rceil\}$  for  $q \in (0, 1]$ , and let  $C_t$  be the unrescaled current-parameter Gram over the last  $m_t$  observations. Thus Proposition 28 gives  $\kappa_{+,t} = 1$  and  $u_t = (1 + \bar{\chi}_t)^2$ . Suppose candidate gradients lie in*

a fixed subspace  $U$ , and let  $P_U$  be its orthogonal projector. For a supplied burn-in index  $\tau_w \in \{1, \dots, T+1\}$ , assume the predictable window-excitation condition

$$\mathcal{C}_t \succeq \lambda \mathbf{I} + \kappa_w m_t P_U$$

holds for a known  $\kappa_w > 0$  and every  $t \geq \tau_w$ . Then, on the same event of probability at least  $1 - \delta - \delta_{\text{cert}}$  as Theorem 1,

$$\begin{aligned} \Lambda_T^{\mathcal{C}} &\leq L_q(T; \tau_w), \\ L_q(T; \tau_w) &:= (\tau_w - 1) \log \left( 1 + \frac{G^2}{\lambda \sigma^2} \right) \\ &\quad + \frac{G^2}{\sigma^2} \sum_{t=\tau_w}^T \frac{1}{\lambda + \kappa_w m_t}, \\ L_q(T; \tau_w) &= \begin{cases} O(\tau_w + T^{1-q}), & q < 1, \\ O(\tau_w + \log T), & q = 1. \end{cases} \end{aligned} \quad (26)$$

For  $\tau_w = 1$ , more explicitly,  $L_q(T; 1) \leq (G^2/\sigma^2)[\lambda^{-1} + \kappa_w^{-1} h_q(T)]$ , where, for  $T \geq 2$ ,  $h_q(T) = 1 + ((T-1)^{1-q} - 1)/(1-q)$  for  $q < 1$  and  $h_1(T) = 1 + \log(T-1)$ , while  $h_q(1) := 0$  in both cases. Writing  $r_U = \dim U$ , the excitation premise can first hold nontrivially only once  $m_t \geq r_U$ , because  $\text{rank}(\mathcal{C}_t - \lambda \mathbf{I}) \leq m_t$ ; thus  $\tau_w$  includes the required rank- $r_U$  burn-in. If  $H_T = \max_{t \leq T} \alpha_t \sqrt{u_t(a_t)} \omega_t$ , then

$$R_T \leq 2H_T \sqrt{(\sigma^2 + G^2/\lambda) T L_q(T; \tau_w)} + 2E_T. \quad (27)$$

If  $H_T = \tilde{O}(1)$ ,  $E_T = \tilde{O}(T^{1-q/2})$ , and the burn-in obeys  $\tau_w = \tilde{O}(T^{1-q})$  for  $q < 1$  (or  $\tau_w = \tilde{O}(1)$  for  $q = 1$ ), then (27) simplifies to  $R_T = \tilde{O}(T^{1-q/2})$ ; otherwise its displayed  $L_q$  and  $E_T$  terms remain. This is conditional on the fixed subspace, excitation, and all confidence, transfer, optimizer, and CG certificates above. Certified worst-case width-solve work is  $O(KT^{1+3q/2} \log(1/\bar{\varepsilon}))$  sample-CVPs when a fixed uniform target  $\bar{\varepsilon} \in (0, 1)$  is enforced; this count excludes model optimization, candidate-gradient computation, and certificate construction.

*Proof.* Before  $\tau_w$ ,  $\mathcal{C}_t \succeq \lambda \mathbf{I}$  bounds each logarithm in  $\Lambda_T^{\mathcal{C}}$  by  $\log(1 + G^2/(\lambda \sigma^2))$ . Thereafter, for  $g_t(a) \in U$ , inversion of the excitation inequality gives  $s_{\mathcal{C},t}^2(a) \leq G^2/(\lambda + \kappa_w m_t)$ ; summing  $\log(1+x) \leq x$  proves (26). Since  $S_T \leq TH_T^2$ , substitution in Theorem 1 proves (27). The condition-number bound  $\kappa(\mathcal{C}_t) \leq 1 + m_t G^2/(\lambda \sigma^2)$  makes CG use  $O(\sqrt{m_t} \log(1/\bar{\varepsilon}))$  iterations, each replaying  $m_t$  samples; summing  $Km_t^{3/2}$  proves the cost claim.  $\square$

Under the hypotheses of Corollary 14, if the window policy uses the predictable rank envelope  $\Gamma_{t-1,r}$  instead of the realized  $\hat{\gamma}_{t-1}$  in its confidence radius, then

$H_T \leq \alpha_I(1+x_T)\Omega_{T,r}^{\text{win}}$ , where  $\Omega_{T,r}^{\text{win}}$  is the explicit  $\Omega_{T,r}$  in Corollary 14 with  $A_{\text{inf}} = 1$ . Hence (27) contains no uncontrolled information-gain placeholder. The excitation premise is restrictive and standard in persistently excited identification; contextual diversity alone does not imply it. For  $\dim U = r > 1$  one must choose  $\tau_w$  after a rank- $r$  burn-in, as the formal early charge records. The predictable rank schedule in Corollary 14 makes  $H_T$  polylogarithmic under the same bounded-path conditions. Illustrative achieved upper-bound exponents are reported in Appendix A.4; they are not a Pareto frontier or a lower bound. When  $U$  is supplied, the projected Gram in (43) can instead be solved directly.

## 5 Experiments

All regret rows are independently executed policies. Common random numbers are used only where pairing is valid. Exact teacher and dense replay quantities are post-action audits; common-trajectory operator comparisons never report causal regret. Tuning and evaluation seeds are disjoint. The derived full-grid aggregate retains seed-level checkpoints and complete hash provenance; release scope and missing raw tiers are stated in Appendix J.

The grid contains three ambient dimensions, three active ranks, eight methods, and 50 fixed seeds. For the column-orthonormal embedding  $E_d$ ,

$$(E_d g)^\top (\lambda I_d + E_d A_t E_d^\top)^{-1} (E_d g) = g^\top (\lambda I_r + A_t)^{-1} g.$$

Thus non-diagonal widths are ambient-invariant in exact arithmetic. Empirically, the maximum cross- $d$  regret range is zero; the dynamic-width range is  $5.7 \times 10^{-14}$  for non-diagonal methods and  $2.2501 \times 10^3$  for the embedding-dependent diagonal method. Ambient-CVP runtime is excluded. Every scaling method uses the same full-history optimizer from (3);  $m_t$  changes only the curvature operator. The implementation also constructs full-history dense audit objects for window rows, so its runtime is not a measurement of the window-only cost model.

The policy predeclares  $\zeta_t = 10^{-6}$  and stops after at most 40 optimizer iterations; it does not enforce that envelope. A failed field therefore means the executed radius does not instantiate the theorem on that round, so the path is an uncertified diagnostic. Five-point slopes and fit residuals are in Table 4.

Full CG is a one-step equivalence test: each right-hand side is an active coordinate and the Gram is diagonal in that basis. Width errors are about  $10^{-15}$ , yet near-tie perturbations can change an action and the subsequent adaptive trajectory. The rotated-SPD audit in Appendix C, not this scaling row, tests nontrivial Krylov convergence.




Figure 1: Finite-horizon theorem-scaling diagnostic. The first panel is the event-clean  $d = 128, r = 4$  reference; the remaining panels use  $d = 2048$ . Curves are 50-seed means and bands are pointwise percentile-bootstrap 95% intervals at five preregistered prefixes. Panel labels report the conjunction of float64 theorem-event fields for exact, CG, and the three window methods. Greedy is uncertified. No exponent comparison tests Theorem 5 unless every premise and audit field for that method/rank cell passes.

### 5.1 Predictable nonlinear execution

We first use the genuinely nonlinear Gaussian model  $\mu_\theta(x, a) = \tanh(\phi(x, a)^\top \theta)$  with  $d = 6$ , four bounded actions,  $T = 100$ ,  $\|\theta^*\| \leq 0.35$ , and a radius-0.5 projected optimizer. The analytic constants in Corollary 15, exact full-objective gradient norm, the path statistics in (19), and the analytic condition bound  $1 + (t - 1)G^2/(\lambda\sigma^2)$  are computed before action selection. No teacher, future reward, exact  $\chi_t$ , exact  $\psi_t$ , or exact Taylor error enters the policy. A fixed protocol is evaluated on 50 seeds (160–209); an eight-cell one-factor sensitivity grid is confined to ten tuning seeds.

Every original- and corrected-center trajectory passes every recorded transfer, centering, linearization, information, CG, confidence, optimism, and regret event: zero observed failures over 5,000 rounds per center. Mean regret is 7.86 [7.66, 8.07] for the original center and 5.69 [5.46, 5.92] for the corrected center. The schedules are deliberately conservative: at the final round, original  $\chi_t$  averages 0.0216 versus  $\bar{\chi}_t = 4.89$ , and  $\psi_t$  averages 0.0520 versus  $\bar{\psi}_t = 128.1$ . The mean observable RHS/regret ratios are 6,956.7 and 1,328.9,

respectively, so the bounds are numerically vacuous. Because the residual and dense audit checks are ordinary float64 point computations, these executions are classified *post-hoc theorem-event verified with fully predictable schedules*, not verified-enclosure theorem certified.

### 5.2 Reference-operator falsification and matched baselines

An eight-cell phase map runs seven operators on 30 independent policies per cell and replays the exact-full path for detailed width/rank diagnostics. Full CG matches dense full regret; diagonal has lower regret in all eight cells, block in six, while full beats window and stale refresh in all eight and rank-3 Lanczos in seven. Thus fidelity does not order regret. The separate linear-Gram witness changes the diagonal policy’s action, but is existential rather than a dominance result (Figure 4 and table 6).

Scalar calibration changes conclusions but does not induce a uniform ordering. Across the 48 surrogate-cell comparisons per protocol, Holm correction over the full 144-comparison family leaves 35 significant under



| Cell ( $d/r$ ) | Exact | CG    | $q = 1/2$ | $q = 2/3$ | $q = 1$ | Frozen | Diag.  | Greedy |
|----------------|-------|-------|-----------|-----------|---------|--------|--------|--------|
| 128/4          | 1.80  | 1.85  | 98.30     | 17.23     | 1.80    | 153.60 | 153.60 | 0.82   |
| 2048/4         | 1.80  | 1.85  | 98.30     | 17.23     | 1.80    | 153.60 | 153.60 | 0.82   |
| 2048/8         | 7.50  | 7.82  | 108.61    | 106.44    | 7.50    | 108.61 | 108.61 | 2.13   |
| 2048/16        | 43.95 | 45.24 | 76.80     | 76.80     | 43.95   | 76.80  | 76.80  | 6.27   |

| Theorem-event status and failed round-event categories |                  |                   |                   |                                       |
|--------------------------------------------------------|------------------|-------------------|-------------------|---------------------------------------|
| Method                                                 | $d = 128, r = 4$ | $d = 2048, r = 4$ | $d = 2048, r = 8$ | $d = 2048, r = 16$                    |
| Exact current                                          | PASS             | PASS              | PASS              | FAIL [ $O=49, P_e=46, P_\lambda=45$ ] |
| Full CG                                                | PASS             | FAIL [ $O=1$ ]    | PASS              | FAIL [ $O=49, P_e=46, P_\lambda=45$ ] |
| Window $q = 1/2$                                       | PASS             | PASS              | FAIL [ $O=1$ ]    | FAIL [ $O=49, P_e=46, P_\lambda=45$ ] |
| Window $q = 2/3$                                       | PASS             | PASS              | PASS              | FAIL [ $O=49, P_e=46, P_\lambda=45$ ] |
| Window $q = 1$                                         | PASS             | PASS              | PASS              | FAIL [ $O=49, P_e=46, P_\lambda=45$ ] |

PASS means that all recorded float64 theorem-event fields passed for that method and cell.  
Failure counts are round-event fields:  $O$  is optimizer residual,  $P_e$  is  $\psi$ -excitation, and  $P_\lambda$  is  $\psi$ -lambda.

Table 1: Preregistered scaling summary. Panel (a) reports terminal mean regret for every executed method. Panel (b) decomposes failed float64 theorem-event fields by method:  $O$  is the prescribed optimizer-residual envelope,  $P_e$  is the path bound using the excitation floor, and  $P_\lambda$  is the path bound using damping. Pass means zero observed failures for the displayed fields, not a verified numerical enclosure. Any nonzero entry precludes interpreting that cell’s slope as a test of the corresponding conditional theorem.




Figure 2: Coverage-matched operator mechanism map. Entries are mean method-minus-current-full-GGN cumulative pseudo-regret in all eight preregistered condition/rotation/gap cells (50 evaluation seeds). Scalar coefficients are either identical, selected on tuning seeds 2000–2019 for 95% one-step coverage, or selected there to match mean bonus magnitude; evaluation seeds 2100–2149 are untouched during selection. Policies execute independently, so differences are descriptive rather than causal operator effects.

the identical coefficient (13 favor a surrogate, 22 favor full), 20 under coverage matching (8/12), and 27 under mean-bonus matching (11/16). Evaluation coverage under the tuning-selected 95% protocol averages 94.6% across method–cell means and ranges from 85.6% to 99.2%. Current and historical full Grams coincide in this fixed-feature linear environment; their separate labels audit the protocol rather than representation drift. The frozen coefficients are in Table 8, and full coverage/decision heatmaps are in Figure 5.

The balanced nonlinear benchmark uses separate tuning and 30 evaluation seeds; local NeuralUCB is the historical-time full-gradient-Gram comparison, alongside NeuralTS, NeuralLinear, last-layer, diagonal, linear, greedy, and context-free baselines. These are local matched implementations, not verified reproductions.

At  $T = 200$ , CC-UCB full GGN-CG has regret 21.94 [20.64, 23.23]; diagonal has 16.11 [14.62, 17.61], LinUCB 10.58 [9.82, 11.33], NeuralLinear 10.01 [9.33, 10.68], and frozen last-layer UCB 7.44 [6.90, 7.99]. Context-free controls have regret 90.35 and 89.40; the sanity check passes, but the experiment is negative evidence for full-curvature superiority.

A separate balanced-MNIST pipeline fixes disjoint supervised-pretraining, tuning, and online-evaluation splits and the requested local baseline set. Its canonical-data smoke profile passes, but the full tuning and evaluation grid has not run, so we report no MNIST outcome. We do not substitute a different image dataset under the MNIST label or use smoke output as paper evidence.

---

### 5.3 Audits and systems scope

All 280 bounded linear traces pass float64 checks, but lack verified enclosures; their bounds are vacuous. Synthetic CPU operators reach  $d = 8192$ , but use unmatched stored-feature primitives. Small-model policies likewise materialize Jacobians for dense audits, validating algebra rather than on-the-fly  $O(d)$  Krylov state. A Buck-managed `torch.func` pipeline now executes real JVP/VJP operators without retaining an explicit sample-by-parameter Jacobian; only smoke and development-pilot cells have run. The preregistered evaluation grid is not reported, so no comparative method-speed or deployment claim follows.

**Real-data baseline check.** Covertypes fails its baseline check: context-free controls beat every contextual method at every horizon (Appendix C.1).

## 6 Conclusion

CC-UCB supplies a matrix-free full-GGN/Fisher reference with explicit transfer, CG, and width terms. The strongest full-dimensional result combines predictable small feature drift with a top- $r$  plus spectral-tail bound on the frozen potential; it does not require a known tangent subspace. Exact rank recovers a projected-Gram special case where full-dimensional CG is unnecessary. Neither result is an unrestricted neural-training guarantee. Empirically, the executed studies test width and action-order preservation rather than demonstrate performance: the tanh bound is vacuous, the legacy scaling grid exposes failed envelopes, coverage and bonus matching still produce mixed operator comparisons, full curvature is not uniformly preferred, and Covertypes fails its baseline check.

**Limitations.** Theory assumes Gaussian squared loss, enumeration, realizability, smoothness, and, for the original center, an optimizer certificate; approximate operators and screening need further certificates. Controlling feature drift and the spectral tail during unrestricted training remains open. Float64 checks lack enclosures, bounds are vacuous, the full autodiff evaluation grid is unrun, and adaptive comparisons are non-causal.

---

## References

- Yasin Abbasi-Yadkori, Dávid Pál, and Csaba Szepesvári. Improved algorithms for linear stochastic bandits. *Advances in Neural Information Processing Systems*, 24, 2011.
- Shipra Agrawal and Navin Goyal. Thompson sampling for contextual bandits with linear payoffs. In *International Conference on Machine Learning (ICML)*, pages 127–135, 2013.
- Jock A. Blackard. Covertypes [dataset]. UCI Machine Learning Repository, 1998. DOI: 10.24432/C50K5N. Distributed under CC BY 4.0.
- Ha Manh Bui, Enrique Mallada, and Anqi Liu. Variance-aware linear UCB with deep representation for neural contextual bandits. In *Proceedings of the 28th International Conference on Artificial Intelligence and Statistics*, volume 258 of *Proceedings of Machine Learning Research*, pages 4213–4221. PMLR, 2025. URL <https://proceedings.mlr.press/v258/bui25a.html>.
- Peter G. Chang, Gerardo Durán-Martín, Alexander Y. Shestopaloff, Matt Jones, and Kevin Murphy. Low-rank extended Kalman filtering for online learning of neural networks from streaming data. In *Proceedings of the 2nd Conference on Lifelong Learning Agents*, volume 232 of *Proceedings of Machine Learning Research*, 2023. URL <https://proceedings.mlr.press/v232/chang23a.html>.
- Varsha Dani, Thomas P. Hayes, and Sham M. Kakade. Stochastic linear optimization under bandit feedback. In *Proceedings of the 21st Annual Conference on Learning Theory (COLT)*, pages 355–366, 2008.
- Erik Daxberger, Agustinus Kristiadi, Alexander Immer, Runa Eschenhagen, Matthias Bauer, and Philipp Hennig. Laplace Redux—effortless Bayesian deep learning. In *Advances in Neural Information Processing Systems*, volume 34, pages 20089–20103, 2021.
- Rohan Deb, Yikun Ban, Shiliang Zuo, Jingrui He, and Arindam Banerjee. Contextual bandits with online neural regression. In *The Twelfth International Conference on Learning Representations*. OpenReview.net, 2024. URL <https://openreview.net/forum?id=5ep85sakT3>.
- Gerardo Durán-Martín, Aleyna Kara, and Kevin Murphy. Efficient online Bayesian inference for neural bandits. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 6002–6021, 2022.
- Avishek Ghosh, Sayak Ray Chowdhury, and Aditya Gopalan. Misspecified linear bandits. In *AAAI Conference on Artificial Intelligence*, pages 3761–3767, 2017.
- Diego Granziol and Khurshid Juarev. Hessian spectral analysis at foundation model scale. *arXiv preprint arXiv:2602.00816*, 2026.
- David Gross and Vincent Nesme. Note on sampling without replacing from a finite collection of matrices. *arXiv preprint arXiv:1001.2738*, 2010.
- Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.
- Parnian Kassraie and Andreas Krause. Neural contextual bandits without regret. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 240–278, 2022.
- Pang Wei Koh and Percy Liang. Understanding black-box predictions via influence functions. In *International Conference on Machine Learning (ICML)*, pages 1885–1894, 2017.
- Agustinus Kristiadi, Matthias Hein, and Philipp Hennig. Being Bayesian, even just a bit, fixes overconfidence in ReLU networks. In *International Conference on Machine Learning (ICML)*, pages 5436–5446, 2020.
- Ilya Kuzborskij, Leonardo Cella, and Nicolò Cesa-Bianchi. Efficient linear bandits through matrix sketching. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 177–185, 2019.
- Tor Lattimore, Csaba Szepesvári, and Gellért Weisz. Learning with good feature representations in bandits and in RL with a generative model. In *International Conference on Machine Learning (ICML)*, pages 5662–5670, 2020.
- David J. C. MacKay. A practical Bayesian framework for backpropagation networks. *Neural Computation*, 4(3):448–472, 1992.
- Wesley Maddox, Shuai Tang, Pablo Moreno, Andrew Gordon Wilson, and Andreas Damianou. Fast adaptation with linearized neural networks. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 2737–2745, 2021.
- James Martens. New insights and perspectives on the natural gradient method. *Journal of Machine Learning Research*, 21(146):1–76, 2020.
- James Martens and Roger Grosse. Optimizing neural networks with Kronecker-factored approximate curvature. In *International Conference on Machine Learning (ICML)*, pages 2408–2417, 2015.
- Razvan Pascanu and Yoshua Bengio. Revisiting natural gradient for deep networks. *arXiv preprint arXiv:1301.3584*, 2013.
- Barak A. Pearlmutter. Fast exact multiplication by the Hessian. *Neural Computation*, 6(1):147–160, 1994.

- 
- Carlos Riquelme, George Tucker, and Jasper Snoek. Deep Bayesian bandits showdown: An empirical comparison of Bayesian deep networks for Thompson sampling. In *International Conference on Learning Representations (ICLR)*, 2018.
- Hippolyt Ritter, Aleksandar Botev, and David Barber. A scalable Laplace approximation for neural networks. In *International Conference on Learning Representations (ICLR)*, 2018.
- Sudeep Salgia. Provably and practically efficient neural contextual bandits. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 29800–29844. PMLR, 2023. URL <https://proceedings.mlr.press/v202/salgia23a.html>.
- Nicol N. Schraudolph. Fast curvature matrix-vector products for second-order gradient descent. *Neural Computation*, 14(7):1723–1738, 2002.
- William R. Thompson. On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25(3–4):285–294, 1933.
- Joel A. Tropp. An introduction to matrix concentration inequalities. *Foundations and Trends in Machine Learning*, 8(1–2):1–230, 2015.
- Shashanka Ubaru, Jie Chen, and Yousef Saad. Fast estimation of  $\text{tr}(f(a))$  via stochastic Lanczos quadrature. *SIAM Journal on Matrix Analysis and Applications*, 38(4):1075–1099, 2017.
- Michal Valko, Nathan Korda, Rémi Munos, Ilias Flaoumas, and Nello Cristianini. Finite-time analysis of kernelised contextual bandits. In *Uncertainty in Artificial Intelligence (UAI)*, pages 654–663, 2013.
- Dongxie Wen, Hanyan Yin, Xiao Zhang, Peng Zhao, Lijun Zhang, and Zhewei Wei. Revisiting matrix sketching in linear bandits: Achieving sublinear regret via dyadic block sketching. In *International Conference on Learning Representations (ICLR)*, 2026.
- Pan Xu, Zheng Wen, Handong Zhao, and Quanquan Gu. Neural contextual bandits with deep representation and shallow exploration. In *International Conference on Learning Representations (ICLR)*, 2022a.
- Pan Xu, Hongkai Zheng, Eric Mazumdar, Kamyar Azizzadenesheli, and Anima Anandkumar. Langevin Monte Carlo for contextual bandits. In *International Conference on Machine Learning (ICML)*, pages 24830–24850, 2022b.
- Weitong Zhang, Dongruo Zhou, Lihong Li, and Quanquan Gu. Neural Thompson sampling. In *International Conference on Learning Representations (ICLR)*, 2021.
- Dongruo Zhou, Lihong Li, and Quanquan Gu. Neural contextual bandits with UCB-based exploration. In *International Conference on Machine Learning (ICML)*, pages 11492–11502, 2020.

## A Deferred main-text derivations

This appendix collects the formal statements, proofs, and derivations deferred from the compact main text; all cross-referenced results are stated here.

For the sequential frozen-feature Gram, the usual information gain is

$$\gamma_T := \log \frac{\det(\bar{\mathbf{C}}_{T+1})}{\det(\lambda \mathbf{I})} = \sum_{t=1}^T \log(1 + \sigma^{-2} \bar{s}_t^2(a_t)). \quad (28)$$

**Lemma 6** (Rank and statistical-dimension information closure). *For  $A \succeq 0$ ,  $r \geq \text{rank}(A)$ , and  $q_A = \|A\|_{\text{op}}/\lambda$ ,*

$$\log \det(\mathbf{I} + A/\lambda) \leq \min \left\{ r \log \left( 1 + \frac{\text{tr } A}{r\lambda} \right), k(q_A) d_{\text{eff},\lambda}(A) \right\}, \quad (29)$$

where  $d_{\text{eff},\lambda}(A) = \text{tr}[A(A + \lambda \mathbf{I})^{-1}]$ ,  $k(0) = 1$ , and  $k(q) = (1 + q) \log(1 + q)/q$  for  $q > 0$ ; the rank term is defined as zero when  $r = 0$  (which forces  $A = 0$ ). Consequently, if  $\text{rank}(\bar{\mathbf{C}}_{T+1} - \lambda \mathbf{I}) \leq r$ , then  $\gamma_T \leq \Gamma_{T,r}$  from (24).

*Proof.* Writing  $x_i = \lambda_i(A)/\lambda$ , concavity of  $\log(1 + x)$  after padding the spectrum with zeros gives the rank–trace term. Also

$$\log(1 + x) = \frac{(1 + x) \log(1 + x)}{x} \frac{x}{1 + x}.$$

The first factor is increasing because its derivative is  $(x - \log(1 + x))/x^2 \geq 0$ ; bounding it at  $q_A$  and summing proves the statistical-dimension term. For the frozen Gram,  $\text{tr}(\bar{\mathbf{C}}_{T+1} - \lambda \mathbf{I}) \leq TG^2/\sigma^2$ , which gives (24) (with  $r_T = \min\{r, T\}$ ).  $\square$

The second term is useful only with a predictable or deterministic envelope on  $d_{\text{eff},\lambda}$ ; a realized value is post-hoc. It is not the empirical trace effective rank  $\text{tr}(A)/\|A\|_{\text{op}}$ , which does not in general control log determinant.

**Lemma 7** (Spectral-tail information closure). *Let*

$$A_T := \bar{\mathbf{C}}_{T+1} - \lambda \mathbf{I} = \sigma^{-2} \sum_{t=1}^T \mathbf{g}_t \mathbf{g}_t^\top \succeq 0$$

have eigenvalues  $\nu_1 \geq \dots \geq \nu_d \geq 0$ . For an integer  $r \in \{0, \dots, d\}$ , let  $\Delta_{T,r} := \sum_{i>r} \nu_i$  and  $r_T := \min\{r, T\}$ . Then

$$\gamma_T \leq r_T \log \left( 1 + \frac{TG^2}{r_T \lambda \sigma^2} \right) + \frac{\Delta_{T,r}}{\lambda}, \quad (30)$$

where the first term is defined as zero when  $r_T = 0$ . Thus  $r = 0$  gives  $\gamma_T \leq \text{tr}(A_T)/\lambda$ , while  $r \geq \text{rank}(A_T)$  gives  $\Delta_{T,r} = 0$  and recovers the rank closure. The statement also covers  $T = 0$ , when both sides vanish.

*Proof.* Because  $A_T$  is a sum of  $T$  rank-one matrices,  $\text{rank}(A_T) \leq T$ , and Assumption 2 gives  $\text{tr}(A_T) \leq TG^2/\sigma^2$ . If  $r_T > 0$ , concavity of  $x \mapsto \log(1 + x/\lambda)$  on the leading eigenvalues gives

$$\sum_{i=1}^r \log(1 + \nu_i/\lambda) = \sum_{i=1}^{r_T} \log(1 + \nu_i/\lambda) \leq r_T \log \left( 1 + \frac{\sum_{i=1}^{r_T} \nu_i}{r_T \lambda} \right) \leq r_T \log \left( 1 + \frac{TG^2}{r_T \lambda \sigma^2} \right).$$

The equality uses zero eigenvalues when  $r > T$ . For the tail,  $\log(1 + x) \leq x$  yields  $\sum_{i>r} \log(1 + \nu_i/\lambda) \leq \Delta_{T,r}/\lambda$ . Adding the two parts proves the claim. When  $r_T = 0$ , only the tail argument is used, so no zero denominator is evaluated.  $\square$

A deterministic or history-measurable prefix envelope  $\bar{\Delta}_{t,r} \geq \Delta_{t,r}$  may enter an online confidence schedule through the prefix form of (30). In contrast, a realized terminal  $\Delta_{T,r}$  is an a posteriori path-complexity bound only. It must not choose  $r$ , an eigenspace, a radius, or any other action-time quantity. The existing observable  $\hat{\gamma}_t$  schedule remains valid when only the realized terminal tail is reported.

The dynamic-potential and CG-observable consequences of Theorem 1 are

$$\begin{aligned} R_T &\leq 2\sqrt{(\sigma^2 + G^2/\lambda)\Gamma_T^{\text{dyn}}S_T + 2E_T}, \\ R_T &\leq 2\sqrt{(\sigma^2 + G^2/\lambda)\widehat{\Lambda}^C_T S_T + 2E_T}. \end{aligned} \quad (31)$$

**Corollary 8** (Corrected-center variant, no centering certificate). *Replace the predictive center in Algorithm 1 by*

$$\hat{\mu}_t^{\text{corr}}(a) := \mu_{\theta_t}(\mathbf{x}_t, a) + \mathbf{g}_t(a)^\top (\hat{\boldsymbol{\theta}}_t^{\text{lin}} - \boldsymbol{\theta}_t) \quad (32)$$

and use  $\omega_t = \bar{\beta}_t$ . Under the hypotheses of Theorem 1 except Assumption 4, the bounds (15)–(31) hold.

*Proof.* Local linearization gives  $\mu^*(\mathbf{x}_t, a) - \hat{\mu}_t^{\text{corr}}(a) = \mathbf{g}_t(a)^\top (\boldsymbol{\theta}^* - \hat{\boldsymbol{\theta}}_t^{\text{lin}}) + \rho_t$ . Lemma 24 bounds the first term by  $\beta_t \bar{s}_t(a)$  and  $|\rho_t| \leq \varepsilon_{\text{lin}}(t)$ , which is the first step of the regret proof with  $\omega_t = \bar{\beta}_t$ ; all later steps are unchanged.  $\square$

**Remark** (Cost of the corrected center). *Exact evaluation requires a solve against  $\bar{\mathbf{C}}_t$ . A constant number of  $O(d)$  CG vectors is insufficient by itself: one additionally needs stored historical frozen gradients, raw replay data with the historical parameter checkpoints needed to reconstruct them, a frozen-feature matrix–vector oracle, an explicit low-dimensional representation, or a certified sketch. Maintaining  $\bar{\mathbf{C}}_t$  explicitly costs  $O(d^2)$  memory. Approximate frozen-feature access introduces an error that must be returned to the centering certificate unless separately controlled.*

## A.1 Assumptions

**Assumption 1** (Sub-Gaussian noise). *The noise  $\eta_t$  is conditionally  $\sigma$ -sub-Gaussian given the context and action:  $\mathbb{E}[\exp(s\eta_t) \mid \mathcal{G}_t] \leq \exp(s^2\sigma^2/2)$  for all  $s \in \mathbb{R}$ , where  $\mathcal{G}_t := \mathcal{H}_t \vee \sigma(a_t)$  is the pre-reward sigma-algebra of (78), available immediately after the action is selected and before the reward is revealed.*

**Assumption 2** (Bounded gradient features). *For a constant  $G > 0$ ,*

$$\sup_{t \leq T} \sup_{s \leq t} \sup_a \|\nabla_{\boldsymbol{\theta}} \mu_{\theta_t}(\mathbf{x}_s, a)\| \leq G.$$

*This bounds both the current-round features  $\mathbf{g}_t(\mathbf{x}_t, a)$  and the replayed features  $\mathbf{g}_t(\mathbf{x}_s, a_s)$  appearing in the curvature matrix (5), as well as the frozen features  $\mathbf{g}_s(\mathbf{x}_s, a_s)$  in (6).*

**Assumption 3** (Local linearization). *There is a known convex parameter region  $\Theta$  containing  $\boldsymbol{\theta}^*$  and every realized  $\boldsymbol{\theta}_t$ , on which  $\nabla_{\boldsymbol{\theta}} \mu_{\theta}(\mathbf{x}, a)$  is  $L_\mu$ -Lipschitz uniformly over candidate  $(\mathbf{x}, a)$ . Thus*

$$\mu_{\boldsymbol{\theta}}(\mathbf{x}, a) = \mu_{\boldsymbol{\theta}_t}(\mathbf{x}, a) + \mathbf{g}_t(\mathbf{x}, a)^\top (\boldsymbol{\theta} - \boldsymbol{\theta}_t) + \rho_t(\mathbf{x}, a, \boldsymbol{\theta}), \quad (33)$$

*where the remainder satisfies  $|\rho_t(\mathbf{x}, a, \boldsymbol{\theta})| \leq \frac{L_\mu}{2} \|\boldsymbol{\theta} - \boldsymbol{\theta}_t\|^2$  for  $\boldsymbol{\theta} \in \Theta$ . Evaluating at  $\boldsymbol{\theta} = \boldsymbol{\theta}^*$  gives a time-varying linearization error*

$$\varepsilon_{\text{lin}}(t) := \sup_{a \in \mathcal{A}_t} |\rho_t(\mathbf{x}_t, a, \boldsymbol{\theta}^*)| \leq \frac{L_\mu}{2} \|\boldsymbol{\theta}^* - \boldsymbol{\theta}_t\|^2. \quad (34)$$

*For Theorem 1 to give a sublinear guarantee, we require the cumulative linearization error  $E_T := \sum_{t=1}^T \varepsilon_{\text{lin}}(t) = o(T)$ ; see Theorem 1 for the precise dependence. Concretely,  $E_T = O(1)$  when  $\|\boldsymbol{\theta}^* - \boldsymbol{\theta}_t\| = O(1/t)$  (since  $\varepsilon_{\text{lin}}(t) = O(1/t^2)$  and  $\sum t^{-2} < \infty$ ; achievable in convex or strongly convex regimes), and  $E_T = O(\sqrt{T})$  when  $\|\boldsymbol{\theta}^* - \boldsymbol{\theta}_t\| = O(t^{-1/4})$ . In practice this is enforced by restricting fine-tuning depth, using trust-region updates, or freezing the backbone (which renders the last-layer model convex or near-linear with small  $L_\mu$ ). These rate examples are illustrative only. The present paper does not prove convergence of the nonconvex optimization path  $\{\boldsymbol{\theta}_t\}$  to  $\boldsymbol{\theta}^*$ ; a sublinear guarantee from this bound therefore requires this term to be controlled by an external stability or convergence argument.*

**Assumption 4** (Action-scaled centering certificate). *Lemma 9 below proves the geometric bound  $|\mathbf{g}_t(\mathbf{x}_t, a)^\top (\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{s}_t(a) \psi_t$  for every  $a \in \mathcal{A}_t$ , where*

$$\psi_t := \frac{\zeta_t}{\sqrt{\lambda}} + \|M_t\|_{\bar{\mathbf{C}}_t^{-1}}, \quad (35)$$

$\zeta_t \geq 0$  is an  $\mathcal{H}_t$ -measurable optimizer gradient-residual bound ( $\|\nabla \mathcal{L}_t(\boldsymbol{\theta}_t)\|_2 \leq \zeta_t$ ) and  $M_t$  is the gradient mismatch vector defined there. We assume there is an  $\mathcal{H}_t$ -measurable upper bound  $\bar{\psi}_t \geq \psi_t$  (the centering certificate) such that, for every action  $a \in \mathcal{A}_t$ ,

$$|\mathbf{g}_t(\mathbf{x}_t, a)^\top (\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{\psi}_t \bar{s}_t(a). \quad (36)$$

The discrepancy is thus action-scaled: it multiplies the confidence width  $\bar{s}_t(a)$  rather than entering additively. The certificate absorbs the same three sources as before—optimizer residual  $\zeta_t$ , replayed-Jacobian drift inside  $M_t$ , and higher-order nonlinear terms—but now inside a single per-round scalar  $\bar{\psi}_t$ . In the convex-in-parameters regime (frozen backbone with a linear reward head), under exact optimization  $\hat{\boldsymbol{\theta}}_t^{\text{lin}} = \boldsymbol{\theta}_t$  and  $\bar{\psi}_t = 0$ . The corrected-center variant of Corollary 8 removes  $\bar{\psi}_t$  entirely, at the cost of an explicit frozen-feature solve.

**Remark** (Status of Assumption 4). Lemma 11 supplies a predictable valid  $\bar{\psi}_t$  from the optimizer residual, Welford path moments, and observed residual energy. We do not prove that this conservative quantity is small outside a restricted path. It enters the exploration term through  $\omega_t = \bar{\beta}_t + \bar{\psi}_t$  (Equation (69)), not as an additive  $O(T)$  term.

| Assumption                           | Status                   | Use                                | Checkability                          | Experimental instantiation               |
|--------------------------------------|--------------------------|------------------------------------|---------------------------------------|------------------------------------------|
| Conditional Gaussian noise           | sub-Standard             | Self-normalized confidence         | Distributional                        | Gaussian generator                       |
| Realizability, $\ \theta^*\  \leq S$ | Restrictive              | Bias and Taylor schedules          | Not generally testable                | Known synthetic teacher                  |
| Bounded gradients                    | Standard on compact sets | Width and condition bounds         | Analytic or monitored                 | Analytic tanh bound                      |
| $L_\mu, L_g$ on convex $\Theta$      | Restrictive              | Linearization, transfer, centering | Analytic for tanh; otherwise hard     | Global tanh Hessian bound and projection |
| Optimizer residual $\zeta_t$         | Algorithmic              | Centering certificate              | Pre-action gradient norm              | Exact full-objective gradient norm       |
| Operator and CG certificates         | Algorithmic              | Optimism and solve error           | Residual/Loewner proof required       | Analytic condition bound; float64 audit  |
| Residual-energy envelope             | Rate-only                | Near-linear scaling                | Observable after collection           | Reported, not assumed by policy          |
| Fixed active rank $r$                | Restrictive              | Closes $\gamma_T$                  | Supplied subspace, not terminal audit | Constructed active embedding             |
| Window excitation $\kappa_w$         | Restrictive              | Closes $\Lambda_T^C$               | Analytic construction or assumed      | Cyclic construction after burn-in        |

Table 2: Assumption ledger. Analytic and pre-action quantities are policy available; teacher comparisons and exact whitened quantities are post-hoc only.

The next lemma gives the primitive geometric bound underlying Assumption 4, expressed through the certifiable quantity  $\psi_t$ .

**Lemma 9** (Primitive upper bound for prediction discrepancy). *Let  $z_s = (\mathbf{x}_s, a_s)$ ,  $\mathbf{g}_s = \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}_s}(z_s)$ ,  $\mathbf{g}_{s,t} = \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}_t}(z_s)$ , and  $\delta_{s,t} = \mathbf{g}_{s,t} - \mathbf{g}_s$ . Define*

$$\bar{\mu}_{s,t} := \mu_{\boldsymbol{\theta}_s}(z_s) + \mathbf{g}_s^\top (\boldsymbol{\theta}_t - \boldsymbol{\theta}_s), \quad e_{s,t} := \mu_{\boldsymbol{\theta}_t}(z_s) - \bar{\mu}_{s,t}, \quad \bar{r}_{s,t} := \bar{\mu}_{s,t} - r_s.$$

*Let  $\bar{\mathcal{L}}_t$  be the frozen-feature linearized ridge objective whose minimizer is  $\hat{\boldsymbol{\theta}}_t^{\text{lin}}$ , and define*

$$M_t := \frac{1}{\sigma^2} \sum_{s < t} [\bar{r}_{s,t} \delta_{s,t} + e_{s,t} \mathbf{g}_s + e_{s,t} \delta_{s,t}].$$

*If the nonlinear optimizer satisfies  $\|\nabla \mathcal{L}_t(\boldsymbol{\theta}_t)\|_2 \leq \zeta_t$ , then for every candidate action  $a \in \mathcal{A}_t$ ,*

$$|\mathbf{g}_t(a)^\top (\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{s}_t(a) \left( \frac{\zeta_t}{\sqrt{\lambda}} + \|M_t\|_{\bar{\mathbf{C}}_t^{-1}} \right), \quad \bar{s}_t(a) := \sqrt{\mathbf{g}_t(a)^\top \bar{\mathbf{C}}_t^{-1} \mathbf{g}_t(a)}.$$

The left-hand side is exactly the geometric bound  $|\mathbf{g}_t(a)^\top(\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{s}_t(a)\psi_t$  of Assumption 4 with  $\psi_t = \zeta_t/\sqrt{\lambda} + \|M_t\|_{\bar{\mathbf{C}}_t^{-1}}$  (Equation (35)). Consequently, since  $\|M_t\|_{\bar{\mathbf{C}}_t^{-1}} \leq \|M_t\|_2/\sqrt{\lambda}$ , a valid centering certificate is

$$\bar{\psi}_t := \frac{\zeta_t + \|M_t\|_2}{\sqrt{\lambda}} \geq \psi_t.$$

If, in addition, on a fixed event,  $\boldsymbol{\theta} \mapsto \nabla_{\boldsymbol{\theta}}\mu_{\boldsymbol{\theta}}(z)$  is  $L_g$ -Lipschitz,  $|\bar{r}_{s,t}| \leq B_{t,\delta}$  for all  $s < t$ , and  $\Delta_{s,t} := \|\boldsymbol{\theta}_t - \boldsymbol{\theta}_s\|$ , then

$$\|M_t\|_2 \leq \frac{1}{\sigma^2} \sum_{s < t} \left[ B_{t,\delta} L_g \Delta_{s,t} + \frac{G L_g}{2} \Delta_{s,t}^2 + \frac{L_g^2}{2} \Delta_{s,t}^3 \right].$$

*Proof.* The linearized ridge objective satisfies  $\nabla \bar{\mathcal{L}}_t(\hat{\boldsymbol{\theta}}_t^{\text{lin}}) = 0$  and, by its quadratic form,

$$\nabla \bar{\mathcal{L}}_t(\boldsymbol{\theta}_t) = \bar{\mathbf{C}}_t(\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}}).$$

Thus  $\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}} = \bar{\mathbf{C}}_t^{-1} \nabla \bar{\mathcal{L}}_t(\boldsymbol{\theta}_t)$ . Using  $\mu_{\boldsymbol{\theta}_t}(z_s) - r_s = \bar{r}_{s,t} + e_{s,t}$  and  $\mathbf{g}_{s,t} = \mathbf{g}_s + \delta_{s,t}$ , the gradients of the nonlinear loss and the frozen-feature linearized loss satisfy

$$\nabla \mathcal{L}_t(\boldsymbol{\theta}_t) = \nabla \bar{\mathcal{L}}_t(\boldsymbol{\theta}_t) + M_t, \quad \text{so} \quad \nabla \bar{\mathcal{L}}_t(\boldsymbol{\theta}_t) = \nabla \mathcal{L}_t(\boldsymbol{\theta}_t) - M_t.$$

Therefore, by Cauchy–Schwarz in the  $\bar{\mathbf{C}}_t$ -norm,

$$|\mathbf{g}_t(a)^\top(\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{s}_t(a) \|\nabla \bar{\mathcal{L}}_t(\boldsymbol{\theta}_t)\|_{\bar{\mathbf{C}}_t^{-1}} \leq \bar{s}_t(a) \left( \|\nabla \mathcal{L}_t(\boldsymbol{\theta}_t)\|_{\bar{\mathbf{C}}_t^{-1}} + \|M_t\|_{\bar{\mathbf{C}}_t^{-1}} \right).$$

Since  $\bar{\mathbf{C}}_t \succeq \lambda \mathbf{I}$ ,  $\|\nabla \mathcal{L}_t(\boldsymbol{\theta}_t)\|_{\bar{\mathbf{C}}_t^{-1}} \leq \zeta_t/\sqrt{\lambda}$ , which gives the geometric inequality  $|\mathbf{g}_t(a)^\top(\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{s}_t(a)(\zeta_t/\sqrt{\lambda} + \|M_t\|_{\bar{\mathbf{C}}_t^{-1}}) = \bar{s}_t(a)\psi_t$ , i.e. the bound of (35). For the certificate,  $\bar{\mathbf{C}}_t \succeq \lambda \mathbf{I}$  also gives  $\|M_t\|_{\bar{\mathbf{C}}_t^{-1}} \leq \|M_t\|_2/\sqrt{\lambda}$ , so  $\psi_t \leq (\zeta_t + \|M_t\|_2)/\sqrt{\lambda} =: \bar{\psi}_t \geq \psi_t$ , the proposed certificate. For the final display, Lipschitzness gives  $\|\delta_{s,t}\| \leq L_g \Delta_{s,t}$  and the Taylor remainder bound  $|e_{s,t}| \leq (L_g/2)\Delta_{s,t}^2$ ; applying these together with  $\|\mathbf{g}_s\| \leq G$  (Assumption 2) and  $|\bar{r}_{s,t}| \leq B_{t,\delta}$  term by term in the definition of  $M_t$  gives the claimed bound on  $\|M_t\|_2$ .  $\square$

**Lemma 10** (Bounded-output residual energy). *Fix a horizon  $T \geq 1$  and  $\delta_R \in (0, 1)$ . Suppose  $|\mu_{\boldsymbol{\theta}}(x, a)| \leq B_\mu$  throughout the certified parameter region and  $|\mu^*(x, a)| \leq B_\mu$ . Under Assumption 1, with probability at least  $1 - \delta_R$ , simultaneously for every  $t \leq T + 1$ ,*

$$R_t^{\text{col}} \leq \left[ 8B_\mu^2 + 4\sigma^2 \log\left(\frac{2T}{\delta_R}\right) \right] (t - 1). \quad (37)$$

*This finite-horizon envelope is not claimed to be tight.*

*Proof.* Conditional sub-Gaussianity gives

$$\Pr\left(|\eta_s| > \sigma\sqrt{2\log(2T/\delta_R)} \mid \mathcal{G}_s\right) \leq \delta_R/T.$$

Taking expectations and a union bound over  $s = 1, \dots, T$  yields an event of probability at least  $1 - \delta_R$  on which this bound holds for every noise term. On that event,  $c_s = \mu_{\boldsymbol{\theta}_s}(x_s, a_s) - \mu^*(x_s, a_s) - \eta_s$  and  $(a + b)^2 \leq 2a^2 + 2b^2$  give

$$c_s^2 \leq 2(2B_\mu)^2 + 2\sigma^2 2\log(2T/\delta_R) = 8B_\mu^2 + 4\sigma^2 \log(2T/\delta_R).$$

Summing the first  $t - 1$  terms proves (37); at  $t = 1$  both sides are zero.  $\square$

**Lemma 11** ( $O(d)$ -state path certificates). *Suppose  $\nabla_{\boldsymbol{\theta}}\mu_{\boldsymbol{\theta}}(z)$  is  $L_g$ -Lipschitz and bounded by  $G$  on a convex region containing the realized iterates. Before round  $t$ , let  $n_t = t - 1$ , let  $\bar{\boldsymbol{\theta}}_{<t}$  and  $J_{<t} = \sum_{s < t} \|\boldsymbol{\theta}_s - \bar{\boldsymbol{\theta}}_{<t}\|^2$  summarize past iterates, and let  $R_t^{\text{col}} = \sum_{s < t} c_s^2$  with  $c_s = \mu_{\boldsymbol{\theta}_s}(z_s) - r_s$ . At  $t = 1$  use the initialization  $\bar{\boldsymbol{\theta}}_{<1} = \mathbf{0}$ ,  $J_{<1} = Q_1 = R_1^{\text{col}} = 0$ . Then, for all  $t$ ,*

$$Q_t := J_{<t} + n_t \|\boldsymbol{\theta}_t - \bar{\boldsymbol{\theta}}_{<t}\|^2 = \sum_{s < t} \|\boldsymbol{\theta}_t - \boldsymbol{\theta}_s\|^2$$



and the following  $\mathcal{H}_t$ -measurable quantities are valid:

$$\bar{\chi}_t = \frac{L_g}{\sigma\sqrt{\lambda}}\sqrt{Q_t} \geq \chi_t, \quad (38)$$

$$\bar{M}_t = \frac{1}{\sigma^2} \left[ L_g \sqrt{R_t^{\text{col}} Q_t} + \frac{3GL_g}{2} Q_t + \frac{L_g^2}{2} Q_t^{3/2} \right] \geq \|M_t\|_2. \quad (39)$$

If  $\|\theta_s\|, \|\theta_t\| \leq R$ , the final term in (39) may be replaced by  $L_g^2 R Q_t$ . If a nonnegative  $\mathcal{H}_t$ -measurable optimizer certificate satisfies  $\|\nabla \mathcal{L}_t(\theta_t)\| \leq \zeta_t$ , then

$$\bar{\psi}_t = (\zeta_t + \bar{M}_t)/\sqrt{\lambda} \geq \psi_t.$$

If prediction gradients are  $L_\mu$ -Lipschitz on a convex region containing  $\theta_t$  and  $\theta^*$  and  $\|\theta^*\| \leq S$ , then

$$\bar{\epsilon}_t^{\text{lin}} := \frac{L_\mu}{2} (S + \|\theta_t\|)^2 \geq \epsilon_{\text{lin}}(t).$$

Consequently  $\bar{F}_t = \sum_{s < t} (\bar{\epsilon}_s^{\text{lin}})^2 \geq F_t$  and  $\bar{E}_T = \sum_{t \leq T} \bar{\epsilon}_t^{\text{lin}} \geq E_T$ .

*Proof.* The parallel-axis identity follows after writing  $\theta_t - \theta_s = (\theta_t - \bar{\theta}_{<t}) + (\bar{\theta}_{<t} - \theta_s)$ ; the cross term sums to zero. Lemma 16 and  $\|\nabla \mu_{\theta_t}(z_s) - \nabla \mu_{\theta_s}(z_s)\| \leq L_g \|\theta_t - \theta_s\|$  give (38). Let  $d_{s,t} = \|\theta_t - \theta_s\|$ . In Lemma 9,

$$|\bar{r}_{s,t}| \leq |c_s| + G d_{s,t}, \quad \|\delta_{s,t}\| \leq L_g d_{s,t}, \quad |e_{s,t}| \leq (L_g/2) d_{s,t}^2.$$

Thus each summand defining  $M_t$  has norm at most

$$L_g |c_s| d_{s,t} + \frac{3GL_g}{2} d_{s,t}^2 + \frac{L_g^2}{2} d_{s,t}^3.$$

Cauchy–Schwarz gives  $\sum |c_s| d_{s,t} \leq \sqrt{R_t^{\text{col}} Q_t}$ , while  $\sum d_{s,t}^2 = Q_t$  and  $\sum d_{s,t}^3 \leq Q_t^{3/2}$ , proving (39). Under the trust region,  $d_{s,t} \leq 2R$ , so  $\sum d_{s,t}^3 \leq 2R Q_t$ . The centering claim follows from Lemma 9 and  $\bar{\mathbf{C}}_t \succeq \lambda \mathbf{I}$ . Finally Taylor’s theorem and  $\|\theta^* - \theta_t\| \leq S + \|\theta_t\|$  prove the remainder,  $F_t$ , and  $E_T$  claims.  $\square$

**Filtration-order implementation.** At round  $t$ , the summaries through  $t-1$ ,  $\theta_t$ , and the certificate  $\zeta_t$  are available before selection. The learner computes  $Q_t, \bar{\chi}_t, u_t, \bar{M}_t, \bar{\psi}_t, \bar{\epsilon}_t^{\text{lin}}$  and  $\bar{\beta}_t$ , then solves every candidate system to the uniform tolerance and selects  $a_t$ . After observing  $r_t$ , it adds  $c_t^2$  to  $R_{t+1}^{\text{col}}$ , adds  $\theta_t$  with the Welford update

$$d = \theta_t - \bar{\theta}_{<t}, \quad \bar{\theta}_{\leq t} = \bar{\theta}_{<t} + d/t, \quad J_{\leq t} = J_{<t} + d^\top (\theta_t - \bar{\theta}_{\leq t}),$$

and updates  $\bar{F}_{t+1}$  and  $\hat{\gamma}_t$ . It then forms  $\theta_{t+1}$  and certifies the nonnegative bound  $\zeta_{t+1}$  before the next action. No current reward enters the current score.

In frozen linear-head models  $Q_t$  may be nonzero while  $L_g = L_\mu = 0$ ; then  $\bar{\chi}_t = \bar{M}_t = \bar{\epsilon}_t^{\text{lin}} = 0$ , and exact optimization also gives  $\bar{\psi}_t = 0$ .

**Corollary 12** (Bounded-path near-linear rate). *Fix  $T$  and suppose for every  $t \leq T$*

$$\|\theta_t\|, \|\theta^*\| \leq R, \quad L_\mu \leq c_\mu/\sqrt{W}, \quad L_g \leq c_g/\sqrt{W}, \quad \zeta_t \leq \zeta_0/\sqrt{W}, \quad R_t^{\text{col}} \leq C_R t P_T,$$

where  $P_T \geq 1$  is an explicitly supplied residual-energy factor. For exact current-parameter curvature and a uniform CG tolerance  $\bar{\epsilon}_t \leq \bar{\epsilon} < 1$ , set

$$x_T = \frac{2c_g R \sqrt{T-1}}{\sigma \sqrt{\lambda W}}, \quad \rho_- = (1 - x_T)^2, \quad \rho_+ = (1 + x_T)^2.$$

If  $x_T < 1$ , then  $\rho_- \bar{\mathbf{C}}_t \preceq \mathbf{C}_t \preceq \rho_+ \bar{\mathbf{C}}_t$  for all  $t \leq T$ . Writing

$$\alpha_I = \sqrt{\frac{1 + \bar{\epsilon}}{1 - \bar{\epsilon}}}, \quad A_{\text{inf}} = \alpha_I^2 \frac{\rho_+}{\rho_-},$$

the observable information schedule satisfies  $\gamma_T \leq \hat{\gamma}_T \leq A_{\inf} \gamma_T$ , and

$$R_T \leq 2\alpha_I \sqrt{\rho_+/\rho_-} \omega_{\max,T}^{\text{op}} \sqrt{(\sigma^2 + G^2/\lambda)T\gamma_T} + 2\bar{E}_T \quad (40)$$

holds with

$$\omega_{\max,T}^{\text{op}} := \max_{t \leq T} (\bar{\beta}_t + \bar{\psi}_t)$$

and

$$\begin{aligned} \omega_{\max,T}^{\text{op}} &\leq \sqrt{A_{\inf} \gamma_T + 2 \log(1/\delta)} + \sqrt{\lambda} R + \frac{2c_\mu R^2 \sqrt{T}}{\sigma \sqrt{W}} \\ &\quad + \frac{A_T}{\sqrt{\lambda W}} + \frac{B_T}{\sqrt{\lambda W}}, \\ A_T &= \zeta_0 + \sigma^{-2} \left[ 2c_g R \sqrt{C_R T(T-1)P_T} + 6Gc_g R^2(T-1) \right], \\ B_T &= 4c_g^2 R^3(T-1)/\sigma^2. \end{aligned}$$

In addition, (23) holds and

$$\bar{\psi}_t \leq \frac{1}{\sqrt{\lambda}} \left\{ \frac{\zeta_0}{\sqrt{W}} + \frac{1}{\sigma^2} \left[ \frac{2c_g R \sqrt{C_R t(t-1)P_T} + 6Gc_g R^2(t-1)}{\sqrt{W}} + \frac{4c_g^2 R^3(t-1)}{W} \right] \right\}.$$

*Proof.* The trust region gives  $Q_t \leq 4R^2(t-1)$ ; substitution in Lemma 11 proves all displayed path bounds. Corollary 17 gives the two-sided operator relation. The CG sandwich and  $u_t \leq (1+x_T)^2 = \rho_+$  imply that every logarithm in  $\hat{\gamma}_T$  is at most  $\log(1 + A_{\inf} \sigma^{-2} \bar{s}_t^2(a_t))$ . For  $A \geq 1$ ,  $\log(1 + Ax) \leq A \log(1 + x)$ , proving the information comparison. Substitution into Corollary 34 gives the regret display.  $\square$

**Corollary 13** (Corrected-center near-linear rate). *Fix  $T$  and suppose, for every  $t \leq T$ ,*

$$\|\theta_t\|, \|\theta^*\| \leq R, \quad L_\mu \leq c_\mu/\sqrt{W}, \quad L_g \leq c_g/\sqrt{W},$$

where  $W$  is an abstract near-linearity scale. Use the corrected center, exact current curvature, and a uniform CG tolerance  $\bar{\varepsilon}_t \leq \bar{\varepsilon} < 1$ . Fix  $r \in \{0, \dots, d\}$  and suppose a predictable prefix information envelope  $\mathcal{G}_{n,r} \geq \gamma_n$  is supplied for  $n = 0, \dots, T$ . Valid choices are the rank envelope  $\Gamma_{n,r}$  under a known rank bound, or the spectral-tail envelope in Corollary 3 with supplied prefix tails  $\bar{\Delta}_{n,r}$ . Set

$$\begin{aligned} x_T &:= \frac{2c_g R \sqrt{T-1}}{\sigma \sqrt{\lambda W}}, & \alpha_I &:= \sqrt{\frac{1+\bar{\varepsilon}}{1-\bar{\varepsilon}}}, \\ \mathcal{B}_{T,r}^{\text{corr}} &:= \sqrt{\mathcal{G}_{T,r} + 2 \log(1/\delta)} + \sqrt{\lambda} R + \frac{2c_\mu R^2 \sqrt{T}}{\sigma \sqrt{W}}. \end{aligned}$$

If  $x_T < 1$ , the predictable radius

$$\bar{\beta}_t^{\text{corr}} := \sqrt{\mathcal{G}_{t-1,r} + 2 \log(1/\delta)} + \sqrt{\lambda} R + \frac{2c_\mu R^2 \sqrt{t-1}}{\sigma \sqrt{W}} \quad (41)$$

is valid and the corrected-center policy satisfies

$$\begin{aligned} R_T &\leq 2\alpha_I \frac{1+x_T}{1-x_T} \mathcal{B}_{T,r}^{\text{corr}} \sqrt{(\sigma^2 + G^2/\lambda)T\mathcal{G}_{T,r}} \\ &\quad + \frac{4c_\mu R^2 T}{\sqrt{W}}. \end{aligned} \quad (42)$$

In particular, bounded transfer requires only  $W > 4c_g^2 R^2(T-1)/(\lambda\sigma^2) = \Omega(T)$  up to problem constants, and  $E_T \leq 2c_\mu R^2 T/\sqrt{W}$ . If  $\mathcal{G}_{T,r} = O(r \log T)$ —for example, if  $\bar{\Delta}_{T,r} = O(r \log T)$  at fixed  $\lambda$ —and  $W = \Omega(T)$  with a sufficiently large constant to keep  $x_T$  bounded away from one, then

$$R_T = O(r\sqrt{T} \log T) + O(\sqrt{T})$$

up to confidence and CG constants.

This corollary assumes neither an optimizer-residual rate nor an envelope on  $R_t^{\text{col}}$ . Its cost is frozen-feature access: exact corrected centers require an explicit  $d \times d$  frozen Gram, stored frozen gradients, replay with historical checkpoints, a supplied low-dimensional representation, or a separately certified sketch. The result therefore does not have  $O(d)$  total memory merely because current-curvature products are matrix-free.

*Proof.* The trust region gives  $Q_t \leq 4R^2(t-1)$ , so Lemma 11 yields  $\bar{\chi}_t \leq 2c_g R\sqrt{t-1}/(\sigma\sqrt{\lambda W}) \leq x_T$ ,  $\bar{\epsilon}_t^{\text{lin}} \leq 2c_\mu R^2/\sqrt{W}$ ,  $\bar{F}_t \leq 4c_\mu^2 R^4(t-1)/W$ , and  $\bar{E}_T \leq 2c_\mu R^2 T/\sqrt{W}$ . Hence (41) dominates (72). Apply the corrected-center form of Corollary 2, bound  $\alpha_t \leq \alpha_I$ ,  $(1 + \bar{\chi}_t)/(1 - \bar{\chi}_t) \leq (1 + x_T)/(1 - x_T)$ , and  $\bar{\beta}_t^{\text{corr}} \leq \mathcal{B}_{T,r}^{\text{corr}}$ , then use  $\gamma_T \leq \mathcal{G}_{T,r}$ . Finally substitute  $2\bar{E}_T \leq 4c_\mu R^2 T/\sqrt{W}$ . The stated orders follow directly.  $\square$

**Remark** (Why  $W$  is an abstract near-linearity scale). *The all-weights MLP used in the nonlinear diagnostics has action-a mean  $v_a^\top \tanh(Ux + b) + c_a$  and no  $W^{-1/2}$  forward normalization. With  $\bar{x} = (x, 1)$ , hidden vector  $q_i = (U_i, b_i)$ , and  $z_i = q_i^\top \bar{x}$ , the Hessian block on  $(q_i, v_{a,i})$  is*

$$\begin{bmatrix} v_{a,i} \tanh''(z_i) \bar{x} \bar{x}^\top & \tanh'(z_i) \bar{x} \\ \tanh'(z_i) \bar{x}^\top & 0 \end{bmatrix}.$$

At  $z_i = 0$  and  $v_{a,i} = 0$  its operator norm is  $\|\bar{x}\|$ , independent of hidden width. Thus the executed all-weights parameterization does not imply  $L_\mu, L_g = O(W^{-1/2})$ . Such a conclusion would require a different explicitly normalized parameterization and coordinate constraints, which would also change the ridge penalty and GGN geometry. We therefore do not identify  $W$  with generic neural-network width or derive optimizer-residual scaling from it.

**Corollary 14** (Conditional projected-Gram rate for a supplied subspace). *Under Corollary 12, additionally suppose a subspace  $U \subseteq \mathbb{R}^d$  with  $\dim U \leq r$  is fixed and supplied before interaction, and every candidate or replayed tangent feature through round  $T$  lies in  $U$ . Hence  $\text{rank}(\bar{\mathbf{C}}_t - \lambda \mathbf{I}) \leq r$  for every  $t \leq T+1$ . Let  $r_U = \dim U$ , let  $Q_U \in \mathbb{R}^{d \times r_U}$  have orthonormal columns spanning  $U$ , and define*

$$\mathbf{z}_{t,s} = Q_U^\top \mathbf{g}_t(\mathbf{x}_s, a_s), \quad \mathbf{z}_t(a) = Q_U^\top \mathbf{g}_t(\mathbf{x}_t, a).$$

Then the current-GGN width reduces exactly to

$$\mathbf{g}_t(a)^\top \mathbf{C}_t^{-1} \mathbf{g}_t(a) = \mathbf{z}_t(a)^\top \left( \lambda I_{r_U} + \sigma^{-2} \sum_{s < t} \mathbf{z}_{t,s} \mathbf{z}_{t,s}^\top \right)^{-1} \mathbf{z}_t(a), \quad (43)$$

with the analogous identity for  $\bar{\mathbf{C}}_t$  using  $Q_U^\top \mathbf{g}_s(\mathbf{x}_s, a_s)$ . Thus a supplied basis permits an exact dense  $r_U \times r_U$  Gram with  $O(r_U^2)$  curvature state. Given tangent coordinates, forming the current Gram costs  $O((t-1)r_U^2)$ , factorization  $O(r_U^3)$ , and each width query  $O(r_U^2)$ ; current-parameter replay, model training, and gradient generation can still be  $d$ -dimensional. Full-dimensional CG is not needed in this corollary's covered regime.

Define  $\Gamma_{T,r}$  by (24). Then (25) holds with

$$\begin{aligned} \Omega_{T,r} &:= \sqrt{A_{\text{inf}} \Gamma_{T,r} + 2 \log(1/\delta)} + \sqrt{\lambda} R + \frac{2c_\mu R^2 \sqrt{T}}{\sigma \sqrt{W}} + \frac{A_T}{\sqrt{\lambda W}} + \frac{B_T}{\sqrt{\lambda W}}, \\ A_T &= \zeta_0 + \sigma^{-2} \left[ 2c_g R \sqrt{C_R T(T-1) P_T} + 6Gc_g R^2 (T-1) \right], \\ B_T &= 4c_g^2 R^3 (T-1)/\sigma^2, \end{aligned}$$

and  $A_{\text{inf}}, \rho_\pm$  are exactly those in Corollary 12. In particular, for fixed  $r$  and fixed problem constants, if  $W \geq CT^2 P_T$  for a constant  $C$  satisfying the explicit thresholds in (44), then

$$R_T = O(r\sqrt{T} \log T) + O(T/\sqrt{W}).$$

This rate is conditional on the ex ante subspace  $U$  and every bounded-path, near-linearity, residual, optimizer, transfer, and solver premise of Corollary 12. It neither learns  $U$  nor covers unrestricted representation learning or end-to-end neural-network training; it is an effective projected-feature closure, not a full-dimensional computational advantage for CC-UCB.

*Proof.* Extend  $Q_U$  to an orthogonal basis of  $\mathbb{R}^d$ . In that basis  $\mathbf{C}_t$  is block diagonal with the  $r_U \times r_U$  matrix in (43) and  $\lambda I_{d-r_U}$ ; the query has zero coordinates in the second block. This proves (43), and the same argument proves the frozen-Gram identity. Lemma 6 gives  $\gamma_T \leq \Gamma_{T,r}$ . Every occurrence of  $\gamma_T$  in Corollary 12 is monotone, so direct substitution gives the display. Under the stated near-linearity scaling,  $x_T$ , the centering terms, and the linearization terms are bounded, while  $\Gamma_{T,r} = O(r \log T)$ . Hence  $\Omega_{T,r} = O(\sqrt{r \log T} + 1)$  and the leading product is  $O(r\sqrt{T} \log T)$ .  $\square$

**Remark** (Linear-subclass lower bound). *The exact-rank class contains a standard  $r$ -dimensional stochastic linear bandit. Take a constant context, a known orthonormal basis  $Q_U$ , the finite action set  $z(a) \in \{\pm 1/\sqrt{r}\}^r$ , linear mean  $\mu_\vartheta(a) = z(a)^\top \vartheta$ , and Gaussian noise. Choosing the usual hypercube parameter family with coordinates of order  $\sqrt{r/T}$  keeps the parameter norm bounded when  $T \gtrsim r^2$ . Here  $L_\mu = L_g = E_T = 0$ ,  $\mathbf{C}_t = \hat{\mathbf{C}}_t$ , and all features lie in  $U$ . The stochastic linear-bandit minimax expected-regret lower bound is  $\Omega(r\sqrt{T})$  in this regime (Dani et al., 2008). Hence the leading  $r\sqrt{T}$  dependence of the exact-rank upper bound is sharp up to logarithms on this linear subclass. This does not prove a lower bound for the nonlinear or spectral-tail classes, nor a computational lower bound.*

For explicit targets  $\bar{x}_{\text{drift}} \in (0, 1)$  and  $\Psi, b_F, \eta_E > 0$ , the sufficient conditions

$$\begin{aligned} W &\geq \frac{4c_g^2 R^2 (T-1)}{\lambda \sigma^2 \bar{x}_{\text{drift}}^2}, & W &\geq \frac{4c_\mu^2 R^4 T}{\sigma^2 b_F^2}, & W &\geq \frac{4c_\mu^2 R^4}{\eta_E^2}, \\ W &\geq \max \left\{ \frac{4A_T^2}{\lambda \Psi^2}, \frac{2B_T}{\sqrt{\lambda} \Psi} \right\} \end{aligned} \quad (44)$$

respectively ensure  $x_T \leq \bar{x}_{\text{drift}} < 1$  (and therefore  $\rho_- \geq (1 - \bar{x}_{\text{drift}})^2$  and  $\rho_+ \leq (1 + \bar{x}_{\text{drift}})^2$ ),  $\sqrt{F_{T+1}}/\sigma \leq b_F$ ,  $\bar{E}_T/T \leq \eta_E$ , and  $\bar{\psi}_T \leq \Psi$ . Thus the centering requirement has leading sufficient scaling  $W = \Omega(T^2 P_T)$ ; these conditions do not bound  $\gamma_T$ .

**Corollary 15** (Nonlinear tanh-link Gaussian bandit). *For  $\mu_\theta(x, a) = \tanh(\phi(x, a)^\top \theta)$  with  $\|\phi(x, a)\| \leq B_\phi$ , one may take*

$$G = B_\phi, \quad L_\mu = L_g = \frac{4}{3\sqrt{3}} B_\phi^2.$$

*On any certified trust region, Corollary 4 therefore gives an executable nonlinear Gaussian-reward CC-UCB policy using only the online path statistics, a pre-action optimizer-residual certificate, and the stated operator/CG certificates. Under realizability, both the model and teacher outputs have  $B_\mu = 1$ . Hence Lemma 10 removes the supplied residual-energy factor in Corollary 12: on its event one may take*

$$C_R = 1, \quad P_T = 8 + 4\sigma^2 \log(2T/\delta_R),$$

*and allocate  $\delta_R$  to the certificate failure budget.*

*Proof.*  $\nabla_\theta \mu = \text{sech}^2(\phi^\top \theta) \phi$ , so its norm is at most  $B_\phi$ . The Hessian is  $-2 \tanh(z) \text{sech}^2(z) \phi \phi^\top$ . For  $y = |\tanh z|$ ,  $2y(1-y^2)$  is maximized at  $y = 1/\sqrt{3}$  with value  $4/(3\sqrt{3})$ , proving both Lipschitz constants. Apply Lemma 11 and Corollary 4.  $\square$

**Predictable algorithmic curvature under certified transfer.** The analysis is conducted for a *generic predictable SPD algorithmic curvature sequence*  $\{\mathcal{C}_t\}_{t \geq 1}$  satisfying

$$\mathcal{C}_t \text{ is symmetric,} \quad \mathcal{C}_t \succeq \lambda \mathbf{I} \succ 0, \quad \mathcal{C}_t \text{ is } \mathcal{H}_t\text{-measurable.} \quad (45)$$

This is the operator whose inverse the algorithm actually applies (via CG) to form widths. The exact full-history result is the special case  $\mathcal{C}_t = \mathbf{C}_t$  (relinearized curvature (5)); approximate, subsampled, windowed, refreshed, or sketched operators  $\hat{\mathbf{C}}_t$  are also represented by  $\mathcal{C}_t$ , but are covered only when the required transfer certificate below is available. Write the algorithmic width

$$s_{\mathcal{C}_t}^2(a) := \mathbf{g}_t(a)^\top \mathcal{C}_t^{-1} \mathbf{g}_t(a). \quad (46)$$

**Assumption 5** (One-sided optimism transfer). *For every round  $t \leq T$  there is an  $\mathcal{H}_t$ -measurable factor  $u_t(a) \geq 1$  such that, for every candidate action  $a \in \mathcal{A}_t$ ,*

$$\bar{s}_t^2(a) \leq u_t(a) s_{\mathcal{C}_t, t}^2(a). \quad (47)$$

This is the *only* spectral relation the main theorem needs, and it is one-sided: it dominates the predictable frozen-feature width  $\bar{s}_t(a)$  (the object of the self-normalized confidence bound) by an inflated algorithmic width. No *lower* Loewner factor on  $\mathcal{C}_t$  is required for optimism. The factor may be action-dependent; when a uniform per-round  $u_t$  satisfies (47) for every action we write  $u_t$  and state the action-dependent theorem with  $u_t(a)$ . For the exact-curvature instantiation  $\mathcal{C}_t = \mathbf{C}_t$ , the whitened relative-drift lemma below certifies (47) with the uniform factor  $u_t = (1 + \bar{\chi}_t)^2$ ; for an approximate operator, the composition of Lemma 16 with a one-sided spectral bound  $\hat{\mathbf{C}}_t \preceq \kappa_{+,t} \mathbf{C}_t$  gives  $u_t = \kappa_{+,t}(1 + \bar{\chi}_t)^2$  (Theorem 29).

The following lemma supplies the exact-curvature transfer factor through a *whitened, relative* feature-drift bound: the drift is measured in the  $\bar{\mathbf{C}}_t^{-1}$  metric and accumulated in  $\ell_2$ , and the resulting one-sided bound needs *no* smallness condition.

**Lemma 16** (Whitened relative feature-drift bound). *Let  $\mathbf{g}_s := \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}_s}(\mathbf{x}_s, a_s)$ ,  $\mathbf{g}_{s,t} := \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}_t}(\mathbf{x}_s, a_s)$ ,  $\delta_{s,t} := \mathbf{g}_{s,t} - \mathbf{g}_s$ , and  $\Delta_{s,t} := \|\boldsymbol{\theta}_t - \boldsymbol{\theta}_s\|$ . Using the symmetric (principal) square root  $\bar{\mathbf{C}}_t^{1/2}$ , define the  $d \times (t-1)$  whitened column matrices*

$$\mathbf{G}_t := \left[ \bar{\mathbf{C}}_t^{-1/2} \mathbf{g}_s / \sigma \right]_{s < t}, \quad \mathbf{D}_t := \left[ \bar{\mathbf{C}}_t^{-1/2} \delta_{s,t} / \sigma \right]_{s < t}, \quad \chi_t := \|\mathbf{D}_t\|_{\text{op}}. \quad (48)$$

Writing  $n_t = t - 1$ , also define the  $(d + n_t) \times d$  stacked matrices

$$\mathbf{A}_t := \begin{bmatrix} \sqrt{\lambda} \bar{\mathbf{C}}_t^{-1/2} \\ \mathbf{G}_t^\top \end{bmatrix}, \quad \mathbf{B}_t := \begin{bmatrix} \sqrt{\lambda} \bar{\mathbf{C}}_t^{-1/2} \\ (\mathbf{G}_t + \mathbf{D}_t)^\top \end{bmatrix}. \quad (49)$$

At  $t = 1$  the lower blocks have zero rows. All square roots and inverse square roots in this statement are principal. Then:

- (a)  $\mathbf{G}_t \mathbf{G}_t^\top = \mathbf{I} - \lambda \bar{\mathbf{C}}_t^{-1} \preceq \mathbf{I}$ , hence  $\|\mathbf{G}_t\|_{\text{op}} \leq 1$ .
- (b)  $\bar{\mathbf{C}}_t^{-1/2} (\mathbf{C}_t - \bar{\mathbf{C}}_t) \bar{\mathbf{C}}_t^{-1/2} = \mathbf{G}_t \mathbf{D}_t^\top + \mathbf{D}_t \mathbf{G}_t^\top + \mathbf{D}_t \mathbf{D}_t^\top$ , so  $\|\bar{\mathbf{C}}_t^{-1/2} (\mathbf{C}_t - \bar{\mathbf{C}}_t) \bar{\mathbf{C}}_t^{-1/2}\|_{\text{op}} \leq 2\chi_t + \chi_t^2$ .
- (c) The stacked designs obey

$$\mathbf{A}_t^\top \mathbf{A}_t = \mathbf{I}_d, \quad \mathbf{B}_t^\top \mathbf{B}_t = \bar{\mathbf{C}}_t^{-1/2} \mathbf{C}_t \bar{\mathbf{C}}_t^{-1/2}, \quad \|\mathbf{B}_t - \mathbf{A}_t\|_{\text{op}} = \chi_t. \quad (50)$$

- (d) (**One-sided, no smallness condition.**) For every  $\chi_t \geq 0$ ,

$$\mathbf{C}_t \preceq (1 + \chi_t)^2 \bar{\mathbf{C}}_t, \quad \text{hence} \quad \bar{s}_t^2(a) \leq (1 + \chi_t)^2 s_t^2(a) \quad \forall a. \quad (51)$$

A predictable certificate  $\bar{\chi}_t \geq \chi_t$  therefore satisfies Assumption 5 in the exact case  $\mathcal{C}_t = \mathbf{C}_t$  with the uniform factor  $u_t = (1 + \bar{\chi}_t)^2$ .

- (e) (**Whitened  $\ell_2$  primitive bound.**)

$$\chi_t \leq \|\mathbf{D}_t\|_F = \left[ \sigma^{-2} \sum_{s < t} \|\delta_{s,t}\|_{\bar{\mathbf{C}}_t^{-1}}^2 \right]^{1/2} \leq \frac{L_g}{\sigma \sqrt{\lambda}} \left[ \sum_{s < t} \|\boldsymbol{\theta}_t - \boldsymbol{\theta}_s\|^2 \right]^{1/2}, \quad (52)$$

the last step under the  $L_g$ -Lipschitz-Jacobian assumption  $\|\delta_{s,t}\| \leq L_g \Delta_{s,t}$ .

*Proof.* (a)  $\mathbf{G}_t \mathbf{G}_t^\top = \sigma^{-2} \sum_{s < t} \bar{\mathbf{C}}_t^{-1/2} \mathbf{g}_s \mathbf{g}_s^\top \bar{\mathbf{C}}_t^{-1/2} = \bar{\mathbf{C}}_t^{-1/2} (\sigma^{-2} \sum_{s < t} \mathbf{g}_s \mathbf{g}_s^\top) \bar{\mathbf{C}}_t^{-1/2} = \bar{\mathbf{C}}_t^{-1/2} (\bar{\mathbf{C}}_t - \lambda \mathbf{I}) \bar{\mathbf{C}}_t^{-1/2} = \mathbf{I} - \lambda \bar{\mathbf{C}}_t^{-1}$ . Since  $\lambda \bar{\mathbf{C}}_t^{-1} \succeq 0$ ,  $\mathbf{G}_t \mathbf{G}_t^\top \preceq \mathbf{I}$ , and  $\|\mathbf{G}_t\|_{\text{op}}^2 = \|\mathbf{G}_t \mathbf{G}_t^\top\|_{\text{op}} \leq 1$ .

(b) From  $\mathbf{g}_{s,t} = \mathbf{g}_s + \delta_{s,t}$ ,  $\mathbf{g}_{s,t} \mathbf{g}_{s,t}^\top - \mathbf{g}_s \mathbf{g}_s^\top = \mathbf{g}_s \delta_{s,t}^\top + \delta_{s,t} \mathbf{g}_s^\top + \delta_{s,t} \delta_{s,t}^\top$ . Conjugating each summand by  $\bar{\mathbf{C}}_t^{-1/2}$  and weighting by  $\sigma^{-2}$ , the three groups sum exactly to  $\mathbf{G}_t \mathbf{D}_t^\top$ ,  $\mathbf{D}_t \mathbf{G}_t^\top$ , and  $\mathbf{D}_t \mathbf{D}_t^\top$ . By the triangle inequality and  $\|\mathbf{G}_t\|_{\text{op}} \leq 1$ ,  $\|\mathbf{G}_t \mathbf{D}_t^\top + \mathbf{D}_t \mathbf{G}_t^\top + \mathbf{D}_t \mathbf{D}_t^\top\|_{\text{op}} \leq 2\|\mathbf{G}_t\|_{\text{op}} \|\mathbf{D}_t\|_{\text{op}} + \|\mathbf{D}_t\|_{\text{op}}^2 \leq 2\chi_t + \chi_t^2$ .

(c) By part (a),  $\mathbf{A}_t^\top \mathbf{A}_t = \lambda \bar{\mathbf{C}}_t^{-1} + \mathbf{G}_t \mathbf{G}_t^\top = \mathbf{I}_d$ . Expanding the current replay Gram gives

$$\mathbf{B}_t^\top \mathbf{B}_t = \lambda \bar{\mathbf{C}}_t^{-1} + (\mathbf{G}_t + \mathbf{D}_t)(\mathbf{G}_t + \mathbf{D}_t)^\top = \bar{\mathbf{C}}_t^{-1/2} \mathbf{C}_t \bar{\mathbf{C}}_t^{-1/2}.$$

The two stacks differ only in their bottom block, hence  $\|\mathbf{B}_t - \mathbf{A}_t\|_{\text{op}} = \|\mathbf{D}_t^\top\|_{\text{op}} = \chi_t$ . These identities also hold at  $t = 1$ , with empty bottom blocks and  $\chi_1 = 0$ .

(d) The matrix  $W_t := \bar{\mathbf{C}}_t^{-1/2}(\mathbf{C}_t - \bar{\mathbf{C}}_t)\bar{\mathbf{C}}_t^{-1/2}$  is symmetric with  $\|W_t\|_{\text{op}} \leq 2\chi_t + \chi_t^2$ , hence  $W_t \preceq (2\chi_t + \chi_t^2)\mathbf{I}$ . Congruence by  $\bar{\mathbf{C}}_t^{1/2}$  gives  $\mathbf{C}_t - \bar{\mathbf{C}}_t \preceq (2\chi_t + \chi_t^2)\bar{\mathbf{C}}_t$ , i.e.  $\mathbf{C}_t \preceq (1 + \chi_t)^2 \bar{\mathbf{C}}_t$ ; no bound on  $\chi_t$  is used. Both  $\mathbf{C}_t, \bar{\mathbf{C}}_t \succ 0$ , so inversion reverses the order:  $\mathbf{C}_t^{-1} \succeq (1 + \chi_t)^{-2} \bar{\mathbf{C}}_t^{-1}$ . The quadratic form at  $\mathbf{g}_t(a)$  gives  $s_t^2(a) \geq (1 + \chi_t)^{-2} \bar{s}_t^2(a)$ , i.e.  $\bar{s}_t^2(a) \leq (1 + \chi_t)^2 s_t^2(a)$ . Since  $x \mapsto (1 + x)^2$  is increasing on  $[0, \infty)$ , replacing  $\chi_t$  by any  $\bar{\chi}_t \geq \chi_t$  preserves the inequality, giving  $u_t = (1 + \bar{\chi}_t)^2$ .

(e)  $\chi_t = \|\mathbf{D}_t\|_{\text{op}} \leq \|\mathbf{D}_t\|_F$ , and  $\|\mathbf{D}_t\|_F^2 = \sum_{s < t} \|\bar{\mathbf{C}}_t^{-1/2} \delta_{s,t}/\sigma\|^2 = \sigma^{-2} \sum_{s < t} \delta_{s,t}^\top \bar{\mathbf{C}}_t^{-1} \delta_{s,t} = \sigma^{-2} \sum_{s < t} \|\delta_{s,t}\|_{\bar{\mathbf{C}}_t^{-1}}^2$ . Since  $\bar{\mathbf{C}}_t^{-1} \preceq \lambda^{-1} \mathbf{I}$ ,  $\|\delta_{s,t}\|_{\bar{\mathbf{C}}_t^{-1}}^2 \leq \|\delta_{s,t}\|^2/\lambda \leq L_g^2 \Delta_{s,t}^2/\lambda$ ; summing and taking square roots gives (52).  $\square$

**Corollary 17** (Sharp two-sided feature-drift sandwich). *If  $\chi_t < 1$ , then*

$$(1 - \chi_t)^2 \bar{\mathbf{C}}_t \preceq \mathbf{C}_t \preceq (1 + \chi_t)^2 \bar{\mathbf{C}}_t. \quad (53)$$

*Consequently, if a predictable certificate satisfies  $\chi_t \leq \bar{\chi}_t < 1$ , then  $(1 - \bar{\chi}_t)^2 \bar{\mathbf{C}}_t \preceq \mathbf{C}_t \preceq (1 + \bar{\chi}_t)^2 \bar{\mathbf{C}}_t$ .*

*Proof.* Because  $\mathbf{A}_t^\top \mathbf{A}_t = \mathbf{I}_d$ , all  $d$  singular values of the  $(d + n_t) \times d$  matrix  $\mathbf{A}_t$  equal one. Singular-value perturbation and (50) give, for  $j = 1, \dots, d$ ,

$$|\sigma_j(\mathbf{B}_t) - 1| \leq \|\mathbf{B}_t - \mathbf{A}_t\|_{\text{op}} = \chi_t.$$

For  $\chi_t < 1$ , every singular value of  $\mathbf{B}_t$  therefore lies in  $[1 - \chi_t, 1 + \chi_t]$ . Squaring the singular values and using  $\mathbf{B}_t^\top \mathbf{B}_t = \bar{\mathbf{C}}_t^{-1/2} \mathbf{C}_t \bar{\mathbf{C}}_t^{-1/2}$  yields  $(1 - \chi_t)^2 \mathbf{I} \preceq \bar{\mathbf{C}}_t^{-1/2} \mathbf{C}_t \bar{\mathbf{C}}_t^{-1/2} \preceq (1 + \chi_t)^2 \mathbf{I}$ . Congruence by the principal  $\bar{\mathbf{C}}_t^{1/2}$  proves (53). If  $0 \leq \chi_t \leq \bar{\chi}_t < 1$ , then  $(1 - \bar{\chi}_t)^2 \leq (1 - \chi_t)^2$  and  $(1 + \chi_t)^2 \leq (1 + \bar{\chi}_t)^2$ , proving the certificate version.  $\square$

The main theorem uses only the one-sided bound (51), which holds for *any*  $\chi_t \geq 0$ ; the smallness condition  $\chi_t < 1$  is needed only for the lower half of the two-sided sandwich (53), which in turn is used by the explicit near-linear corollary and by any optional analysis that explicitly assumes a two-sided sandwich. An unrescaled current-parameter window is handled by Proposition 28 through a one-sided relation and does not require this condition. The binding quantity is the *whitened  $\ell_2$  accumulation*  $\chi_t \leq \|\mathbf{D}_t\|_F$  of replayed-Jacobian drift, measured in the  $\bar{\mathbf{C}}_t^{-1}$  metric—a relative, whitened  $\ell_2$  bound that can be sharper than an  $\ell_1$  sum normalized only by  $\lambda$  when the drift occurs in well-observed directions (it credits directions the design has already explored), though it does not dominate uniformly. Together, Lemmas 9 and 16 identify exact whitened diagnostic targets. The operational upper bounds in Lemma 11 replace those targets by the predictable path statistics  $Q_t$  and  $R_t$ , the certified optimizer residual  $\zeta_t$ , and analytic smoothness constants. Thus  $\bar{\psi}_t$  and  $u_t = (1 + \bar{\chi}_t)^2$  require only  $O(d)$  additional state; the exact whitened quantities remain post-hoc diagnostics used to measure the conservatism of these certificates.

Positive definiteness of both  $\mathbf{C}_t$  and  $\bar{\mathbf{C}}_t$  follows from  $\lambda > 0$  and the outer-product form; it is not an independent assumption. All theoretical results hold under these explicit assumptions; we do *not* claim a regret guarantee for the full nonlinear model outside the local-linearization regime.

## A.2 Novelty Relative to Prior Neural Bandits

NeuralUCB/NeuralTS use historical-time full-gradient Grams in their stated algorithms, while NeuralLinear and last-layer Laplace restrict uncertainty to a learned representation or final layer (Zhou et al., 2020; Zhang et al., 2021; Riquelme et al., 2018; Kristiadi et al., 2020). Matrix-free full-network predictive-variance solves also predate this work (Maddox et al., 2021). CC-UCB’s narrower adaptation is to replay every historical Jacobian at the current parameter, separate the predictable frozen confidence Gram from that algorithmic operator, and expose centering, drift, transfer, and residual-checked CG terms. We claim neither the CG/CVP primitive nor a prior assertion of uniform full-metric superiority as new.

The one-sided whitened result in Lemma 16 is the main technical step: optimism transfer needs no lower spectral comparison. The determinant formula is bookkeeping rather than a standalone rate, and the supplied- $U$  closure is exactly the projected Gram in (43), not a concrete neural-training theorem. Appendix B compares NTK, regression, EKF/LO-FI, structured Laplace, and sketching alternatives.

### A.3 Computable Curvature Operator

Applying  $\mathbf{C}_t \mathbf{v} = \lambda \mathbf{v} + \sigma^{-2} \sum_{s < t} \mathbf{g}_t(\mathbf{x}_s, a_s) [\mathbf{g}_t(\mathbf{x}_s, a_s)^\top \mathbf{v}]$  *naively* requires one curvature-vector product (CVP) per historical sample per CG step—cost  $O(It)$  CVPs per round. Each CVP evaluates one Jacobian-vector product (forward mode) followed by one vector-Jacobian product (reverse mode) through the network, at roughly the cost of one gradient evaluation (one JVP plus one VJP)<sup>1</sup>; we do not store per-sample gradient features  $\mathbf{g}_s$ , only a constant number of  $O(d)$  CG vectors for CG, with CVPs computed on the fly. We describe three strategies that limit the per-step cost to  $|\mathcal{S}_t|$  CVPs:

**(i) Subsampled GGN (worst-case spectral baseline).** At each round draw a uniform subsample  $\mathcal{S}_t \subset \{1, \dots, t-1\}$  of size  $n$  (without replacement) and set  $w_s = (t-1)/n$  to obtain the rescaled operator. The following lemma is a worst-case uniform spectral guarantee.

**Lemma 18** (Worst-case uniform spectral guarantee for subsampled GGN). *Fix round  $t$  and set  $m = t-1$ . Let  $Y_s = \sigma^{-2} \mathbf{g}_t(\mathbf{x}_s, a_s) \mathbf{g}_t(\mathbf{x}_s, a_s)^\top$  for  $s = 1, \dots, m$ , so that  $\mathbf{C}_t = \lambda \mathbf{I} + V$  with  $V = \sum_{s=1}^m Y_s$ . Each  $Y_s$  is PSD with  $\|Y_s\|_{\text{op}} \leq G^2/\sigma^2 =: B$ . For  $\delta \in (0, 1)$  and  $1 \leq n \leq m$ , draw  $\mathcal{S}_t$  of size  $n$  uniformly without replacement from  $\{1, \dots, m\}$  and define  $\hat{V} = \frac{m}{n} \sum_{s \in \mathcal{S}_t} Y_s$  and  $\hat{\mathbf{C}}_t = \lambda \mathbf{I} + \hat{V}$ . Write  $L := \log(2d/\delta)$ .*

(i) Deviation bound. *With probability  $\geq 1 - \delta$ ,*

$$\|\hat{V} - V\|_{\text{op}} \leq mB \sqrt{\frac{2L}{n}} + \frac{4mBL}{3n}. \quad (54)$$

(ii) Sufficient sample size. *For any  $\varepsilon \in (0, 1)$ , if*

$$n \geq 8L \left[ \frac{m^2 B^2}{\varepsilon^2 \lambda^2} + \frac{mB}{3\varepsilon \lambda} \right], \quad (55)$$

*then  $\Pr[\|\hat{V} - V\|_{\text{op}} > \varepsilon \lambda] \leq \delta$  and consequently  $(1-\varepsilon)\mathbf{C}_t \preceq \hat{\mathbf{C}}_t \preceq (1+\varepsilon)\mathbf{C}_t$ .*

*Proof.* The  $\{Y_s\}_{s=1}^m$  are fixed PSD self-adjoint matrices (deterministic conditional on the history  $\mathcal{F}_{t-1}$  and the current parameter  $\theta_t$ ). The displayed proof is for *with-replacement* (i.i.d.) sampling. The without-replacement version follows from the domination argument of Gross and Nemes (2010), building on Hoeffding’s reduction (Hoeffding, 1963): the matrix moment-generating function of a sum sampled without replacement from a fixed finite collection of self-adjoint matrices is dominated by that of the corresponding i.i.d. sum, so every tail bound derived from the matrix Laplace-transform method, including the matrix Bernstein inequality of Tropp (2015, Thm. 6.1.1), carries over with the same constants. In our setting the  $\{Y_s\}$  are *fixed* PSD matrices (deterministic conditional on  $\mathcal{F}_{t-1}$  and  $\theta_t$ ) and only the index set  $\mathcal{S}_t$  is random, so the domination argument applies directly. We therefore state and prove the bound for i.i.d. draws  $s_1, \dots, s_n$  uniform on  $\{1, \dots, m\}$ ; the without-replacement guarantee is at least as strong. Define the centered random matrices  $Z_k = (m/n)Y_{s_k} - V/n$  (mean zero, independent). The range satisfies  $\|Z_k\|_{\text{op}} \leq (m/n)B + (1/n)\|V\| \leq 2mB/n =: R$  (using  $\|V\| \leq mB$ ). For the variance parameter, since  $Y_s \preceq B\mathbf{I}$ ,

$$\mathbb{E}[Z_k Z_k^\top] \preceq \frac{m^2}{n^2} \mathbb{E}[Y_{s_1}^2] = \frac{m^2}{n^2} \sum_{s=1}^m Y_s^2 \preceq \frac{m^2}{n^2} B V \preceq \frac{m^2 B^2}{n^2} \mathbf{I},$$

so  $v := \|\sum_{k=1}^n \mathbb{E}[Z_k Z_k^\top]\| \leq m^2 B^2/n$ . By the matrix Bernstein inequality (Tropp, 2015, Thm. 6.1.1),  $\Pr[\|\hat{V} - V\| \geq t] \leq 2d \exp(-t^2/(2v + 2Rt/3))$ . Choosing  $t_0 = \sqrt{2vL} + \frac{2}{3}RL$  makes the matrix-Bernstein tail probability at most  $\delta$ . Substituting  $v \leq m^2 B^2/n$  and  $R = 2mB/n$  gives  $t_0 \leq mB \sqrt{2L/n} + \frac{4mBL}{3n}$ , yielding (54). For (55), require each summand in (54) to be  $\leq \varepsilon \lambda/2$ :  $mB \sqrt{2L/n} \leq \varepsilon \lambda/2$  gives  $n \geq 8m^2 B^2 L/(\varepsilon^2 \lambda^2)$ , and  $4mBL/(3n) \leq \varepsilon \lambda/2$  gives

<sup>1</sup>Implementations based on double-backward Hessian-vector products cost closer to two gradient-equivalents per CVP; the constant does not affect the  $O(\cdot)$  accounting.

$n \geq 8mBL/(3\varepsilon\lambda)$ . Summing the two conditions yields (55); this is a conservative sufficient condition (using  $a + b$  rather than  $\max\{a, b\}$ ), not a sharp threshold. The multiplicative sandwich follows:  $\|\hat{V} - V\| \leq \varepsilon\lambda$  gives  $\mathbf{C}_t - \varepsilon\lambda\mathbf{I} \preceq \hat{\mathbf{C}}_t \preceq \mathbf{C}_t + \varepsilon\lambda\mathbf{I}$ . Since  $\mathbf{C}_t \succeq \lambda\mathbf{I}$ , we have  $\varepsilon\lambda\mathbf{I} \preceq \varepsilon\mathbf{C}_t$ , so  $\mathbf{C}_t - \varepsilon\lambda\mathbf{I} \succeq (1 - \varepsilon)\mathbf{C}_t$  and  $\mathbf{C}_t + \varepsilon\lambda\mathbf{I} \preceq (1 + \varepsilon)\mathbf{C}_t$ .  $\square$

With  $B = G^2/\sigma^2$ , the right-hand side of the sufficient condition (55) is  $O(L \cdot [m^2G^4/(\varepsilon^2\lambda^2\sigma^4) + mG^2/(\varepsilon\lambda\sigma^2)])$ . The first (variance) term dominates unless  $\varepsilon\lambda\sigma^2 \gtrsim mG^2$ , so this worst-case sufficient threshold has a *quadratic leading dependence* on  $m = t-1$  (it is a sufficient upper bound on the required sample size, not an exact or necessary sample complexity). The condition (55) is conservative and sufficient, not necessary, with worst-case leading dependence quadratic in  $m$ . If its threshold exceeds the population size  $m = t-1$ , it provides no nontrivial subsampling guarantee for  $n < m$ . Taking  $n = m$  recovers the full population exactly for the separate reason that every summand is included.

*This is a worst-case baseline guarantee, too conservative for practical use at large  $t$ .* In particular, the bound above does not by itself justify computational savings at large  $t$ ; it should be read as a worst-case spectral guarantee rather than as a practical scaling result. Subsampling reduces computation ( $n \ll m$ ) only when  $m$  is moderate or  $\varepsilon\lambda$  is large relative to  $mB$ ; for large  $t$  the full history may be needed under the uniform Bernstein bound. Under additional spectral structure, leverage-score sampling or Frequent Directions sketches may yield useful approximations with smaller state. For linear bandits, Kuzborskij et al. (2019) give regret bounds in terms of the Frequent Directions spectral-tail error. Translating such structure into a uniform GGN sandwich and a concrete buffer or sketch size is left open. The Bernstein guarantee is not the mechanism by which we expect practical savings at scale. Practical savings require small buffers treated heuristically, spectral decay/effective dimension, ridge-leverage sampling, Frequent Directions, or empirical verification of the spectral sandwich. We leave structure-aware sampling and empirical spectral certification as future work. To obtain a spectral sandwich that holds simultaneously for all  $t \leq T$ , union-bound over  $t$  (replacing  $\delta$  by  $\delta/T$ ). Theorem 29 then exposes the one-sided transfer factor and realized width complexity as operator-dependent terms along the realized path; it does not identify a standalone causal regret inflation.

**(ii) Windowed curvature (unrescaled, current parameter).** Maintain a sliding window  $\mathcal{S}_t \subseteq \{1, \dots, t-1\}$  of the most recent  $n$  interactions with unit weights ( $w_s=1$ ), *re-evaluating the Jacobians at the current parameter  $\theta_t$* . This is a nonnegative-weighted *subset* of the same current-parameter outer products, so  $\hat{\mathbf{C}}_t \preceq \mathbf{C}_t$  (Proposition 28); Theorem 29 then applies with the *deterministic* approximate-to-full factor  $\kappa_{+,t}=1$  and the full-to-frozen factor  $u_t = (1 + \bar{\chi}_t)^2$ —no two-sided sandwich is required, though the overall guarantee remains conditional on that transfer certificate. The cost of the smaller operator is not hidden: it surfaces as larger realized widths  $s_{\mathcal{C},t}^2(a_t)$  and hence a larger dynamic width complexity  $\Lambda_T^{\mathcal{C}}$ . The same holds for any current-parameter downweighted subset with  $w_s \in [0, 1]$  (including an unrescaled subsample). A *stale* window that reuses Jacobians from an older parameter is *not* covered; see (iii).

**(iii) Periodic refresh (heuristic; stale operator).** Recompute the curvature operator every  $M$  rounds on a cached buffer and reuse it for  $M$  consecutive actions. Because the reused operator evaluates Jacobians at an *older* parameter, it is no longer a subset of the current-parameter outer products, so the  $\hat{\mathbf{C}}_t \preceq \mathbf{C}_t$  argument of Proposition 28 does not apply and the stale operator can both over- and under-estimate  $\mathbf{C}_t$ . This remains a heuristic: Theorem 29 covers it only when a one-sided factor  $\kappa_{+,t}$  (an explicit bound  $\hat{\mathbf{C}}_t \preceq \kappa_{+,t}\mathbf{C}_t$ ) is separately verified, which we evaluate empirically. The rescaled subsample of (i), with weights  $w_s = (t-1)/n > 1$ , is likewise not covered by Proposition 28; it instead uses the two-sided matrix Bernstein bound of Lemma 18.

## A.4 Computational Cost Model

**Primitive assumptions.** CC-UCB requires a single computational primitive: the *curvature-vector product* (CVP)  $\mathbf{v} \mapsto \mathbf{C}_t\mathbf{v}$ , which decomposes into per-sample operations  $\mathbf{v} \mapsto \mathbf{g}_t(\mathbf{x}_s, a_s)[\mathbf{g}_t(\mathbf{x}_s, a_s)^\top \mathbf{v}]$ . Each such operation consists of:

1. a Jacobian-vector product (JVP)  $\mathbf{v} \mapsto \mathbf{g}_t^\top \mathbf{v}$  (forward-mode AD, cost  $\approx 1 \times$  forward pass);
2. a vector-Jacobian product (VJP)  $\alpha \mapsto \alpha \mathbf{g}_t$  (reverse-mode AD, cost  $\approx 1 \times$  backward pass).

Sharded autodiff can distribute JVP/VJP evaluation, but its communication pattern and sample parallelism are implementation-dependent and are not benchmarked here. For  $m_t := |\mathcal{S}_t|$ , one matrix-free oracle application consists of  $m_t$  sample-CVPs plus  $O(d)$  damping and accumulation work. Thus  $I_{t,a}$  CG recurrences for action  $a$



use  $I_{t,a}m_t$  sample-CVPs; an explicit final check of the original-system residual  $\mathbf{g}_t(a) - \mathcal{C}_t\tilde{\mathbf{u}}$  uses another  $m_t$ . The sequential Krylov state is  $O(d)$  ( $O(Kd)$  if right-hand sides are batched), excluding model and optimizer state, activations, and raw replay. Caching an explicit  $m_t \times d$  Jacobian matrix instead uses  $O(m_t d)$  additional state and is not the on-the-fly JVP/VJP implementation analyzed by this operation count.

Krylov methods for Hessian spectral analysis have been run on networks with  $> 10^8$  parameters under distributed training (Granzio and Juarev, 2026), which indicates that curvature–vector primitives of this general form are implementable at scale. We do *not* claim foundation-model feasibility from this: that work times Hessian–vector products, not the exact per-example GGN JVP/VJP used here. Our `real-torch.func` pipeline has only smoke and development-pilot runs, not a completed evaluation-seed systems study. Warm-starting CG from the previous round’s solution can reduce the iteration count in practice.

**Cost accounting.** Let  $C_F(W_{\text{net}})$ ,  $C_B(W_{\text{net}})$ ,  $C_J(W_{\text{net}})$ , and  $C_V(W_{\text{net}})$  denote the one-example forward, reverse, JVP, and VJP costs for model-compute scale  $W_{\text{net}}$ , and define

$$C_{\text{op},t} := m_t[C_J(W_{\text{net}}) + C_V(W_{\text{net}})] + O(d).$$

Parameter count  $d$  alone does not determine these costs when architecture, input length, or weight sharing changes. Let  $J_{t+1}^{\text{opt}}$  be the number of full-batch first-order steps after round  $t$ ;  $q_{t+1}^{\text{opt}} \in \{0, 1\}$  records a separate full-gradient optimizer certificate, and  $q_{t,a}^{\text{CG}} \in \{0, 1\}$  records an explicit final CG residual application. Table 3 separates the training history  $N_t$  from the curvature buffer  $m_t$ .

Table 3: Per-round arithmetic cost.  $N_t = t - 1$  is the training-history size;  $m_t = |\mathcal{S}_t|$  is the curvature-buffer size.

| Component                    | Cost                                                                           |
|------------------------------|--------------------------------------------------------------------------------|
| Candidate means/features     | $O(K[C_F + C_B])$                                                              |
| CG recurrences/checks        | $C_{\text{op},t} \sum_{a=1}^K (I_{t,a} + q_{t,a}^{\text{CG}})$                 |
| Parameter update/certificate | $(J_{t+1}^{\text{opt}} + q_{t+1}^{\text{opt}})\{(N_t + 1)[C_F + C_B] + O(d)\}$ |
| Path statistics/scoring      | $O(d + K)$                                                                     |

This is an operation count, not a wall-clock benchmark. For illustration,  $K = 10$ ,  $I = 25$ , and  $m_t = 256$  give 64,000 sample-CVPs for the recurrences and 66,560 when every solve receives one explicit final residual application. Both numbers exclude candidate gradients, full-history model training, and optimizer certification. Backtracking adds full-history objective passes, and second-order optimization has additional Hessian costs. A linear sufficient-statistic update may substitute a lower special-case cost, but must still supply the theorem’s optimizer certificate.

**Cost of the analyzed configuration.** The full-history instantiation  $\mathcal{C}_t = \mathbf{C}_t$ ,  $\mathcal{S}_t = \{1, \dots, t - 1\}$  with adaptive CG stopping has the certificate  $\bar{\kappa}_t = 1 + (t - 1)G^2/(\lambda\sigma^2)$ , so maintaining  $\max_a \varepsilon_{t,a} \leq \bar{\varepsilon}_t$  is certified by  $I_t = O(\sqrt{t}G/(\sigma\sqrt{\lambda}) \log(1/\bar{\varepsilon}_t))$  iterations per solve (Eq. (85)), so the per-round CG cost is  $K I_t (t - 1) = O(K t^{3/2})$  sample-CVPs, i.e.  $O(K T^{5/2})$  cumulatively—the worst-case width-solve sample-CVP count under the stated CG certificate. End-to-end work additionally contains  $\sum_{t \leq T} (J_t^{\text{opt}} + q_t^{\text{opt}})N_t$  full-data sample-gradient evaluations. This is  $O((J + q)T^2)$  only when those counts are uniformly bounded, and curvature windowing does not reduce it. Table 3 and the numerical example describe fixed-buffer operators ( $m_t \ll N_t$ ), which reduce width-solve work but are covered by Theorem 1 only under the corrected window/subsample conditions of §A.3: an unrescaled current-parameter buffer has the deterministic approximate-to-full factor  $\kappa_{+,t}=1$  (Proposition 28)—the overall guarantee remaining conditional on the full-to-frozen transfer and the other certificates of Theorem 1—while a rescaled or stale buffer additionally requires the verified transfer factor of Theorem 29.

For the growing window of Theorem 5, one solve costs  $O(m_t^{3/2} \log(1/\bar{\varepsilon}))$  sample-CVPs under the analytic condition-number bound. The following table uses the sufficient fixed-burn-in setting  $\tau_w = O(1)$ ,  $H_T = \tilde{O}(1)$ , and  $E_T = \tilde{O}(T^{1-q/2})$ ;  $E_T = o(T)$  alone is insufficient. The sample-CVP column counts only width solves. With uniformly bounded full-batch optimizer/certificate passes, the final column adds the  $T^2$  full-history training term. These are achieved upper bounds for this implementation, not a Pareto frontier, lower bound, or wall-clock measurement (logarithmic factors suppressed):

---

| $q$           | conditional regret | width-solve CVPs   | full-policy sample work |
|---------------|--------------------|--------------------|-------------------------|
| $\frac{1}{2}$ | $T^{\frac{3}{4}}$  | $KT^{\frac{7}{4}}$ | $T^2$                   |
| $\frac{2}{3}$ | $T^{\frac{2}{3}}$  | $KT^2$             | $T^2$                   |
| 1             | $T^{\frac{1}{2}}$  | $KT^{\frac{5}{2}}$ | $T^{\frac{5}{2}}$       |

---

### A.5 An Exact Log-Determinant Bookkeeping Identity

The regret bound is driven not by  $\gamma_T$  but by the realized *width complexity* of the widths the algorithm actually uses,

$$\Lambda_T^{\mathcal{C}} := \sum_{t=1}^T \log(1 + \sigma^{-2} s_{\mathcal{C},t}^2(a_t)). \quad (56)$$

Unlike  $\bar{\mathbf{C}}_t$ , the algorithmic curvature  $\mathcal{C}_t$  need *not* be monotone in  $t$ : relinearization, subsampling, windowing, and periodic refresh can all *shrink* it, so the elementary determinant telescoping of (28) no longer applies. The following lemma is an exact pathwise bookkeeping identity for a general predictable SPD sequence. It is not a standalone rate result: its upper bound is informative only when the terminal log determinant and determinant-decreasing refresh charge are independently controlled, as in Lemma 20.

**Lemma 19** (Dynamic log-determinant potential). *Let  $\{\mathcal{C}_t\}_{t \geq 1}$  satisfy (45) with  $\mathcal{C}_1 = \lambda \mathbf{I}$ . For the played feature  $\mathbf{h}_t := \mathbf{g}_t(\mathbf{x}_t, a_t)$  define the rank-one update  $\mathcal{C}_t^+ := \mathcal{C}_t + \sigma^{-2} \mathbf{h}_t \mathbf{h}_t^\top$  and, using the symmetric (principal) root  $(\mathcal{C}_t^+)^{-1/2}$ , the normalized perturbation*

$$\Xi_t := (\mathcal{C}_t^+)^{-1/2} (\mathcal{C}_{t+1} - \mathcal{C}_t^+) (\mathcal{C}_t^+)^{-1/2}, \quad (57)$$

so that  $\mathbf{I} + \Xi_t = (\mathcal{C}_t^+)^{-1/2} \mathcal{C}_{t+1} (\mathcal{C}_t^+)^{-1/2}$  is SPD; equivalently, every eigenvalue of  $\Xi_t$  exceeds  $-1$  and every eigenvalue of  $\mathbf{I} + \Xi_t$  is positive (with no smallness condition on the jump  $\mathcal{C}_{t+1} - \mathcal{C}_t^+$ ). Fix any end-of-round- $T$  terminal operator  $\mathcal{C}_{T+1}$  constructed only from data and randomness available through round  $T$  ( $\mathcal{F}_T$ -measurable, hence not a function of  $\mathbf{x}_{T+1}$ ,  $\mathcal{A}_{T+1}$ , or future rewards; e.g. the canonical rank-one update  $\mathcal{C}_{T+1} = \mathcal{C}_T^+$ , or the operator the algorithm commits for round  $T+1$ ). The identity below is then an a posteriori statement about the completed sequence  $\mathcal{C}_1, \dots, \mathcal{C}_{T+1}$ ; the quantities  $\Xi_T$ ,  $V_T^{\mathcal{C}}$ , and  $\Gamma_T^{\text{dyn}}$  are defined relative to this choice, and the bound  $\Lambda_T^{\mathcal{C}} \leq \Gamma_T^{\text{dyn}}$  holds for every such valid choice. Then

$$\Lambda_T^{\mathcal{C}} = \log \frac{\det(\mathcal{C}_{T+1})}{\det(\lambda \mathbf{I})} - \sum_{t=1}^T \log \det(\mathbf{I} + \Xi_t). \quad (58)$$

With the determinant-destroying variation charge and the dynamic potential

$$V_T^{\mathcal{C}} := \sum_{t=1}^T [-\log \det(\mathbf{I} + \Xi_t)]_+, \quad \Gamma_T^{\text{dyn}} := \log \frac{\det(\mathcal{C}_{T+1})}{\det(\lambda \mathbf{I})} + V_T^{\mathcal{C}}, \quad (59)$$

we have  $\Lambda_T^{\mathcal{C}} \leq \Gamma_T^{\text{dyn}}$ , and, using  $s_{\mathcal{C},t}^2(a_t) \leq G^2/\lambda$  and  $\log(1+x) \geq x/(1+x)$ ,

$$\sum_{t=1}^T s_{\mathcal{C},t}^2(a_t) \leq (\sigma^2 + G^2/\lambda) \Lambda_T^{\mathcal{C}} \leq (\sigma^2 + G^2/\lambda) \Gamma_T^{\text{dyn}}. \quad (60)$$

*Proof.* By the matrix determinant lemma applied to  $\mathcal{C}_t^+ = \mathcal{C}_t + \sigma^{-2} \mathbf{h}_t \mathbf{h}_t^\top$ ,

$$\frac{\det(\mathcal{C}_t^+)}{\det(\mathcal{C}_t)} = 1 + \sigma^{-2} \mathbf{h}_t^\top \mathcal{C}_t^{-1} \mathbf{h}_t = 1 + \sigma^{-2} s_{\mathcal{C},t}^2(a_t). \quad (61)$$

Writing  $B_t := (\mathcal{C}_t^+)^{-1/2}$ , we have  $\mathbf{I} + \Xi_t = B_t \mathcal{C}_t^+ B_t + B_t (\mathcal{C}_{t+1} - \mathcal{C}_t^+) B_t = B_t \mathcal{C}_{t+1} B_t$ , a congruence of the SPD matrix  $\mathcal{C}_{t+1}$  by the invertible symmetric  $B_t$ , hence SPD (in particular its eigenvalues exceed  $-1$  with no smallness condition on the jump  $\mathcal{C}_{t+1} - \mathcal{C}_t^+$ ), and

$$\frac{\det(\mathcal{C}_{t+1})}{\det(\mathcal{C}_t^+)} = \det(B_t^{-1} (\mathbf{I} + \Xi_t) B_t^{-1}) / \det(\mathcal{C}_t^+) = \det(\mathbf{I} + \Xi_t), \quad (62)$$

using  $\det(\mathcal{C}_{t+1}) = \det(\mathcal{C}_t^+) \det(\mathbf{I} + \Xi_t)$ . Taking logs of (61)–(62) and adding,

$$\log(1 + \sigma^{-2} s_{\mathcal{C},t}^2(a_t)) = \log \det(\mathcal{C}_{t+1}) - \log \det(\mathcal{C}_t) - \log \det(\mathbf{I} + \Xi_t).$$

Summing over  $t = 1, \dots, T$  telescopes the first difference to  $\log \det(\mathcal{C}_{T+1}) - \log \det(\mathcal{C}_1)$ ; with  $\mathcal{C}_1 = \lambda \mathbf{I}$  this is  $\log[\det(\mathcal{C}_{T+1})/\det(\lambda \mathbf{I})]$ , proving (58). Since  $-\log \det(\mathbf{I} + \Xi_t) \leq [-\log \det(\mathbf{I} + \Xi_t)]_+$  term by term,  $-\sum_t \log \det(\mathbf{I} + \Xi_t) \leq V_T^{\mathcal{C}}$ , so  $\Lambda_T^{\mathcal{C}} \leq \Gamma_T^{\text{dyn}}$ . For (60):  $\mathcal{C}_t \succeq \lambda \mathbf{I}$  gives  $s_{\mathcal{C},t}^2(a_t) = \mathbf{h}_t^\top \mathcal{C}_t^{-1} \mathbf{h}_t \leq \|\mathbf{h}_t\|^2/\lambda \leq G^2/\lambda =: c_{\max} \sigma^2$ . With  $x_t := \sigma^{-2} s_{\mathcal{C},t}^2(a_t) \in [0, c_{\max}]$  and  $\log(1+x) \geq x/(1+x) \geq x/(1+c_{\max})$ ,  $\Lambda_T^{\mathcal{C}} \geq (1+c_{\max})^{-1} \sum_t x_t$ , i.e.  $\sum_t s_{\mathcal{C},t}^2(a_t) \leq (\sigma^2 + G^2/\lambda) \Lambda_T^{\mathcal{C}}$ ; the second inequality is  $\Lambda_T^{\mathcal{C}} \leq \Gamma_T^{\text{dyn}}$ .  $\square$

**Remark** (Interpretation and scope of the dynamic potential). *Identity (58) is pathwise.  $V_T^{\mathcal{C}}$  charges only changes that reduce determinant relative to the rank-one update; determinant-increasing changes contribute zero. If  $\mathcal{C}_1 \neq \lambda \mathbf{I}$ , the denominator becomes  $\det(\mathcal{C}_1)$ . Neither  $V_T^{\mathcal{C}}$  nor the endpoint term is generally small. Bounding both is sufficient, but not necessary because the exact identity may cancel them. Thus the identity does not imply sublinear regret; outside Corollary 14 and Theorem 5, controlling  $\Lambda_T^{\mathcal{C}}$  for an adapting backbone remains open.*

**Lemma 20** (Rank-sensitive refresh bound). *Let  $\mathcal{R} = \{t : \mathcal{C}_{t+1} \neq \mathcal{C}_t^+\}$ . Suppose that for every  $t \in \mathcal{R}$ , the negative part  $(\Xi_t)_- = (-\Xi_t)_+$  has rank at most  $r_t$  and  $\lambda_{\min}(\mathbf{I} + \Xi_t) \geq 1 - \nu_t$  for  $0 \leq \nu_t < 1$ . Then*

$$V_T^{\mathcal{C}} \leq \sum_{t \in \mathcal{R}} r_t \log \frac{1}{1 - \nu_t}. \quad (63)$$

*If additionally  $\mathcal{C}_{T+1} = \lambda \mathbf{I} + A$ , where  $A \succeq 0$ ,  $\text{rank}(A) \leq r_{\text{end}}$ , and  $\text{tr}(A) \leq W_{\text{end}} G^2/\sigma^2$  for a specified total endpoint weight  $W_{\text{end}} \geq 0$ , then, for  $r_{\text{end}} > 0$ ,*

$$\Lambda_T^{\mathcal{C}} \leq r_{\text{end}} \log \left( 1 + \frac{W_{\text{end}} G^2}{r_{\text{end}} \lambda \sigma^2} \right) + \sum_{t \in \mathcal{R}} r_t \log \frac{1}{1 - \nu_t}. \quad (64)$$

*For  $r_{\text{end}} = 0$ , the first term is defined as zero.*

*Proof.* Let  $\xi_{t,i}$  be the eigenvalues of  $\Xi_t$ . At most  $r_t$  are negative, each is at least  $-\nu_t$ , and every nonnegative eigenvalue makes a nonpositive contribution to  $-\sum_i \log(1 + \xi_{t,i})$ . Hence  $[-\log \det(\mathbf{I} + \Xi_t)]_+ \leq r_t \log(1/(1 - \nu_t))$ , proving (63). List the nonzero eigenvalues of  $A$  and pad with zeros to obtain  $a_1, \dots, a_{r_{\text{end}}}$ . Concavity gives

$$\sum_i \log(1 + a_i/\lambda) \leq r_{\text{end}} \log \left( 1 + \frac{\sum_i a_i}{r_{\text{end}} \lambda} \right),$$

and the trace premise plus Lemma 19 proves (64).  $\square$

**Consequences.** Sequential Gram updates have  $\Xi_t = 0$ . If canonical rank-one updates are used outside  $J_T$  refresh rounds, and each refresh has  $r_t \leq r$ ,  $\nu_t \leq \nu < 1$ , then  $V_T^{\mathcal{C}} \leq J_T r \log(1/(1 - \nu))$ ; a geometric schedule is therefore polylogarithmic only when these rank and spectral-floor bounds hold. For a replacement  $\mathcal{C}_{t+1} - \mathcal{C}_t^+ = A_t^{\text{new}} - A_t^{\text{old}}$  with both terms PSD and  $\text{rank}(A_t^{\text{old}}) \leq r$ , Sylvester inertia gives  $r_t \leq r$ , but a spectral floor is still required.

For a fixed-feature window of length  $m$ , after warm-up each deletion has rank one. Writing  $N_{\text{drop}} = \max\{0, T - m\}$ , Sherman–Morrison gives the following exact charge for each  $t > m$ . Let  $v_t^{\text{drop}} = \sigma^{-1} \mathbf{g}_{t-m}(\mathbf{x}_{t-m}, a_{t-m})$  and let  $D_t = \mathcal{C}_t^+ - v_t^{\text{drop}} (v_t^{\text{drop}})^\top = \mathcal{C}_{t+1}$  be the remaining-window matrix. Then  $-\log \det(\mathbf{I} + \Xi_t) = \log[1 + (v_t^{\text{drop}})^\top D_t^{-1} v_t^{\text{drop}}]$ , and therefore

$$V_T^{\mathcal{C}} \leq N_{\text{drop}} \log \left( 1 + \frac{G^2}{\lambda \sigma^2} \right).$$

This can be linear in  $T$ . A relinearized window need not be a rank-one replacement. Likewise, merely holding a stale operator fixed does not make non-refresh transitions canonical rank-one updates, so geometric refresh timing alone does not imply a logarithmic charge.

**Corollary 21** (Observable CG upper bound on the width complexity). *If the played-action CG solve satisfies  $\varepsilon_{t,a_t} \leq \bar{\varepsilon}_t < 1$  so that  $\tilde{s}_t^2(a_t) \geq (1 - \bar{\varepsilon}_t) s_{\mathcal{C},t}^2(a_t)$  (Lemma 22), then*

$$\Lambda_T^{\mathcal{C}} \leq \widehat{\Lambda}_T^{\mathcal{C}} := \sum_{t=1}^T \log \left( 1 + \frac{\tilde{s}_t^2(a_t)}{\sigma^2 (1 - \bar{\varepsilon}_t)} \right). \quad (65)$$

*Proof.*  $s_{\mathcal{C},t}^2(a_t) \leq \tilde{s}_t^2(a_t)/(1 - \bar{\varepsilon}_t)$  and  $x \mapsto \log(1 + \sigma^{-2}x)$  is increasing, so each summand of  $\Lambda_T^{\mathcal{C}}$  is dominated term by term.  $\square$

This gives an observable upper bound on the dynamic width-complexity term, conditional on the certified CG tolerances;  $V_T^{\mathcal{C}}$  itself is *not* observable in general (it requires determinant estimation of the curvature sequence).

## A.6 CG Approximation Lemma

**Lemma 22** (CG multiplicative accuracy). *Fix a round  $t$  and an action  $a$  with right-hand side  $\mathbf{g} := \mathbf{g}_t(a) \neq \mathbf{0}$ . Let  $\mathbf{u} = \mathcal{C}_t^{-1}\mathbf{g}$  and let  $\tilde{\mathbf{u}}$  be the iterate after  $I_{t,a}$  steps of CG on the SPD system  $\mathcal{C}_t\mathbf{u} = \mathbf{g}$  starting from  $\tilde{\mathbf{u}}_0 = \mathbf{0}$ , with  $\mathcal{C}_t$  held fixed across all iterations. The relative energy-norm error is the per-action quantity  $\varepsilon_{t,a}$  of (9), and*

$$\varepsilon_{t,a} = \frac{\|\mathbf{u} - \tilde{\mathbf{u}}\|_{\mathcal{C}_t}}{\|\mathbf{u}\|_{\mathcal{C}_t}} \leq 2 \left( \frac{\sqrt{\bar{\kappa}_t} - 1}{\sqrt{\bar{\kappa}_t} + 1} \right)^{I_{t,a}}, \quad (66)$$

where  $\kappa_t := \lambda_{\max}(\mathcal{C}_t)/\lambda_{\min}(\mathcal{C}_t)$  is the true condition number and  $\bar{\kappa}_t \geq \kappa_t$  is a certified  $\mathcal{H}_t$ -measurable upper bound. The convergence rate in (66) is right-hand-side independent, so it holds for every action with the same  $\bar{\kappa}_t$ . For a nonnegative weighted outer-product operator  $\mathcal{C}_t = \lambda\mathbf{I} + \sigma^{-2} \sum_{s \in \mathcal{S}_t} w_s \mathbf{g}_t(\mathbf{x}_s, a_s) \mathbf{g}_t(\mathbf{x}_s, a_s)^\top$  ( $w_s \geq 0$ , so  $\mathcal{C}_t \succeq \lambda\mathbf{I} \succ 0$  and CG is well-posed),

$$\kappa_t \leq \frac{\lambda_{\max}(\mathcal{C}_t)}{\lambda} \leq 1 + \frac{W_t G^2}{\lambda \sigma^2} =: \bar{\kappa}_t,$$

with  $W_t = \sum_{s \in \mathcal{S}_t} w_s$  and  $\|\mathbf{g}_t(\mathbf{x}_s, a_s)\| \leq G$  from Assumption 2; here  $W_t = n$  for unrescaled windowed/refresh buffers, giving  $\bar{\kappa}_t$  and  $\alpha_t$  independent of  $T$ , and  $W_t = t-1$  for full history and for rescaled uniform subsampling with  $w_s = (t-1)/n$ . For a general SPD operator that is not such an outer product,  $\bar{\kappa}_t$  must be certified separately. Write  $\tilde{s}_t^2(a) = \mathbf{g}^\top \tilde{\mathbf{u}}$  and  $s_{\mathcal{C},t}^2(a) = \mathbf{g}^\top \mathcal{C}_t^{-1} \mathbf{g}$ .

**Assume**  $\varepsilon_{t,a} < 1$  (see Appendix E for the minimum  $I_{t,a}$  ensuring this as a function of  $\bar{\kappa}_t$ ; Algorithm 1 enforces a known per-round tolerance  $\bar{\varepsilon}_t$  uniformly over the scored actions,  $\max_a \varepsilon_{t,a} \leq \bar{\varepsilon}_t < 1$ , via adaptive residual stopping or a fixed budget under a certified condition-number bound; see §3). Then

$$(1 - \varepsilon_{t,a}) s_{\mathcal{C},t}^2(a) \leq \tilde{s}_t^2(a) \leq (1 + \varepsilon_{t,a}) s_{\mathcal{C},t}^2(a). \quad (67)$$

For  $\mathbf{g}_t(a) = \mathbf{0}$  the algorithm sets  $\tilde{s}_t^2(a) = 0 = s_{\mathcal{C},t}^2(a)$  and (67) holds trivially.

*Proof.* The CG minimax property gives  $\|e_I\|_{\mathcal{C}_t}/\|e_0\|_{\mathcal{C}_t} \leq \min_{p \in \mathcal{P}_I, p(0)=1} \max_{x \in [\lambda_{\min}, \lambda_{\max}]} |p(x)|$ . Choosing the scaled Chebyshev polynomial bounds this maximum by  $2[(\sqrt{\bar{\kappa}_t} - 1)/(\sqrt{\bar{\kappa}_t} + 1)]^I$ ; this ratio is increasing in  $\kappa_t$ , so  $\bar{\kappa}_t \geq \kappa_t$  proves (66).  $\tilde{s}_t^2(a) - s_{\mathcal{C},t}^2(a) = \mathbf{g}^\top (\tilde{\mathbf{u}} - \mathbf{u}) = (\mathcal{C}_t \mathbf{u})^\top (\tilde{\mathbf{u}} - \mathbf{u}) = \mathbf{u}^\top \mathcal{C}_t (\tilde{\mathbf{u}} - \mathbf{u})$ . By Cauchy-Schwarz in the  $\mathcal{C}_t$ -norm,  $|\mathbf{u}^\top \mathcal{C}_t (\tilde{\mathbf{u}} - \mathbf{u})| \leq \|\mathbf{u}\|_{\mathcal{C}_t} \|\tilde{\mathbf{u}} - \mathbf{u}\|_{\mathcal{C}_t} \leq \varepsilon_{t,a} \|\mathbf{u}\|_{\mathcal{C}_t}^2 = \varepsilon_{t,a} s_{\mathcal{C},t}^2(a)$ . For the outer-product special case the condition-number bound follows from  $\lambda_{\max}(\mathcal{C}_t) \leq \lambda + \sigma^{-2} \sum_{s \in \mathcal{S}_t} w_s \|\mathbf{g}_t(\mathbf{x}_s, a_s)\|^2 \leq \lambda + W_t G^2 / \sigma^2$ .  $\square$

**The algorithmic bonus.** With the *per-round* certified tolerance  $\bar{\varepsilon}_t \geq \max_a \varepsilon_{t,a}$  and the one-sided transfer factor  $u_t(a) \geq 1$  of Assumption 5, Algorithm 1 plays the bonus

$$w_t(a) := \omega_t \sqrt{\frac{u_t(a)}{1 - \bar{\varepsilon}_t}} \sqrt{\tilde{s}_t^2(a)}, \quad (68)$$

where  $\bar{\beta}_t \geq \beta_t$  is the predictable confidence upper bound,  $\bar{\psi}_t$  the centering certificate (Assumption 4), and

$$\omega_t := \bar{\beta}_t + \bar{\psi}_t \quad (69)$$

is the combined confidence-plus-centering width factor. Define the *per-round* CG inflation factor  $\alpha_t := \sqrt{(1 + \bar{\varepsilon}_t)/(1 - \bar{\varepsilon}_t)}$ .

**Lemma 23** (Bonus sandwiching). *Under Assumption 5 and  $\max_{a \in \mathcal{A}_t} \varepsilon_{t,a} \leq \bar{\varepsilon}_t < 1$ , the bonus (68) satisfies, for every  $a \in \mathcal{A}_t$ ,*

$$\omega_t \tilde{s}_t(a) \leq w_t(a) \leq \alpha_t \omega_t \sqrt{u_t(a)} s_{\mathcal{C},t}(a). \quad (70)$$

*Proof.* For  $\mathbf{g}_t(a) = \mathbf{0}$  all three terms vanish and the bound is trivial. Otherwise Lemma 22 with  $\varepsilon_{t,a} \leq \bar{\varepsilon}_t$  gives  $(1 - \bar{\varepsilon}_t)s_{\mathcal{C},t}^2(a) \leq \tilde{s}_t^2(a) \leq (1 + \bar{\varepsilon}_t)s_{\mathcal{C},t}^2(a)$ . The lower bound and  $\omega_t \geq 0$  give  $\sqrt{\tilde{s}_t^2(a)} \geq \sqrt{1 - \bar{\varepsilon}_t} s_{\mathcal{C},t}(a)$ , so  $w_t(a) \geq \omega_t \sqrt{u_t(a)} s_{\mathcal{C},t}(a)$ ; then (47) gives  $\sqrt{u_t(a)} s_{\mathcal{C},t}(a) \geq \bar{s}_t(a)$ , i.e.  $w_t(a) \geq \omega_t \bar{s}_t(a)$ . The upper bound gives  $\sqrt{\tilde{s}_t^2(a)} \leq \sqrt{1 + \bar{\varepsilon}_t} s_{\mathcal{C},t}(a)$ , whence  $w_t(a) \leq \omega_t \sqrt{u_t(a)/(1 - \bar{\varepsilon}_t)} \sqrt{1 + \bar{\varepsilon}_t} s_{\mathcal{C},t}(a) = \alpha_t \omega_t \sqrt{u_t(a)} s_{\mathcal{C},t}(a)$ . Both bounds are per-round; no horizon maximum is taken before summing.  $\square$

## A.7 Confidence and Discrepancy Lemmas

**Lemma 24** (Confidence for the linearized ridge estimator). *Suppose Assumptions 1 to 3 hold and  $\mu^* = \mu_{\theta^*}$  with  $\|\theta^*\| \leq S$ . Define the linearized ridge estimator  $\hat{\theta}_t^{\text{lin}}$  by (80), the frozen-feature variance  $\bar{s}_t(a) = \sqrt{\mathbf{g}_t(a)^\top \bar{\mathbf{C}}_t^{-1} \mathbf{g}_t(a)}$ , and the cumulative squared linearization error  $F_t := \sum_{s=1}^{t-1} \varepsilon_{\text{lin}}(s)^2$ . Then with probability at least  $1 - \delta$ , for all  $t \geq 1$  and all  $a \in \mathcal{A}_t$  simultaneously,*

$$|\mathbf{g}_t(a)^\top (\hat{\theta}_t^{\text{lin}} - \theta^*)| \leq \bar{s}_t(a) \beta_t, \quad (71)$$

where

$$\beta_t := \sqrt{\gamma_{t-1} + 2 \log(1/\delta)} + \sqrt{\lambda} S + \frac{1}{\sigma} \sqrt{F_t}. \quad (72)$$

(Here  $\gamma_{t-1} = \log \det(\bar{\mathbf{C}}_t) / \det(\lambda \mathbf{I})$  is the information gain through round  $t-1$ .)

The quantity  $\beta_t$  in (72) is a theoretical radius. The algorithm may use any predictable upper bound  $\bar{\beta}_t \geq \beta_t$ ; all results below remain valid with  $\bar{\beta}_t$  in place of  $\beta_t$ . This is necessary because  $\beta_t$  depends on the unknown cumulative linearization term  $F_t$ .

*Proof.* From (79),  $y_s = \mathbf{g}_s^\top \theta^* + \rho_s + \eta_s$  with  $|\rho_s| \leq \varepsilon_{\text{lin}}(s)$ . The closed form of  $\hat{\theta}_t^{\text{lin}}$  gives

$$\hat{\theta}_t^{\text{lin}} - \theta^* = -\lambda \bar{\mathbf{C}}_t^{-1} \theta^* + \bar{\mathbf{C}}_t^{-1} \sigma^{-2} \sum_{s < t} \mathbf{g}_s \rho_s + \bar{\mathbf{C}}_t^{-1} \sigma^{-2} \sum_{s < t} \mathbf{g}_s \eta_s.$$

Taking the inner product with  $\mathbf{g}_t(a)$  and bounding in absolute value:

- (i) *Prior bias.*  $|\lambda \mathbf{g}_t(a)^\top \bar{\mathbf{C}}_t^{-1} \theta^*| \leq \lambda \bar{s}_t(a) \|\theta^*\|_{\bar{\mathbf{C}}_t^{-1}} \leq \bar{s}_t(a) \sqrt{\lambda} S$ . (Cauchy–Schwarz in the  $\bar{\mathbf{C}}_t$ -norm;  $\lambda \|\theta^*\|_{\bar{\mathbf{C}}_t^{-1}} \leq \lambda \|\theta^*\| / \sqrt{\lambda_{\min}(\bar{\mathbf{C}}_t)} \leq \sqrt{\lambda} S$  since  $\bar{\mathbf{C}}_t \succeq \lambda \mathbf{I}$ .)
- (ii) *Linearization bias.* Define  $\Phi = [\mathbf{g}_1/\sigma, \dots, \mathbf{g}_{t-1}/\sigma]$  (the  $d \times (t-1)$  design matrix) and  $\hat{\rho} = (\rho_1, \dots, \rho_{t-1})^\top$ . Then  $\sigma^{-2} \sum \mathbf{g}_s \rho_s = \sigma^{-1} \Phi \hat{\rho}$ . Since  $\bar{\mathbf{C}}_t \succeq \Phi \Phi^\top$ , we have  $\Phi^\top \bar{\mathbf{C}}_t^{-1} \Phi \preceq \mathbf{I}$ , so  $\|\Phi \hat{\rho}\|_{\bar{\mathbf{C}}_t^{-1}} = \|\hat{\rho}\|_{\Phi^\top \bar{\mathbf{C}}_t^{-1} \Phi} \leq \|\hat{\rho}\| = \sqrt{F_t}$ . Therefore  $|\mathbf{g}_t(a)^\top \bar{\mathbf{C}}_t^{-1} \sigma^{-1} \Phi \hat{\rho}| \leq \bar{s}_t(a) \sigma^{-1} \sqrt{F_t}$ .
- (iii) *Noise.*

$$\left\| \sum_{s < t} \frac{\mathbf{g}_s}{\sigma} \frac{\eta_s}{\sigma} \right\|_{\bar{\mathbf{C}}_t^{-1}} \leq \sqrt{\gamma_{t-1} + 2 \log(1/\delta)}$$

holds for all  $t \geq 1$  with probability at least  $1 - \delta$  by the time-uniform vector self-normalized martingale bound (Abbasi-Yadkori et al., 2011). Here  $x_s = \mathbf{g}_s/\sigma$  is predictable with respect to the pre-reward sigma-algebra  $\mathcal{G}_s$ , and  $\xi_s = \eta_s/\sigma$  is conditionally 1-sub-Gaussian given  $\mathcal{G}_s$ . For any candidate action  $a$ , Cauchy–Schwarz in the  $\bar{\mathbf{C}}_t$ -norm gives

$$\left| \mathbf{g}_t(a)^\top \bar{\mathbf{C}}_t^{-1} \sum_{s < t} \frac{\mathbf{g}_s \eta_s}{\sigma^2} \right| \leq \bar{s}_t(a) \sqrt{\gamma_{t-1} + 2 \log(1/\delta)}.$$

Summing gives (71)–(72).  $\square$

**Corollary 25** (Predictable confidence schedule). *Suppose a nonnegative  $\mathcal{H}_t$ -measurable sequence  $\{\bar{F}_t\}_{t \geq 1}$  satisfies  $F_t \leq \bar{F}_t$  for all  $t$ . Then Algorithm 1 may set*

$$\bar{\beta}_t := \sqrt{\gamma_{t-1} + 2 \log(1/\delta)} + \sqrt{\lambda} S + \sigma^{-1} \sqrt{\bar{F}_t},$$

which is predictable and satisfies  $\bar{\beta}_t \geq \beta_t$ .

*Proof.*  $\bar{F}_t \geq F_t$  term-by-term implies each summand in (72) is dominated, so  $\bar{\beta}_t \geq \beta_t$ . Predictability holds because  $\gamma_{t-1}$ ,  $S$ ,  $\bar{F}_t$ ,  $\lambda$ ,  $\sigma$ , and  $\delta$  are all determined before round  $t$ .  $\square$

Lemma 11 supplies the computable choice  $\bar{F}_t = \sum_{s < t} (\bar{\epsilon}_s^{\text{lin}})^2$ ; a tighter externally certified predictable sequence may replace it.

**Corollary 26** (Observable information-gain surrogate via the one-sided transfer factor). *Run Algorithm 1 with a predictable algorithmic curvature sequence  $C_s$ , per-round CG tolerance with  $\epsilon_{s,a_s} \leq \bar{\epsilon}_s < 1$  for the played action, and an  $\mathcal{H}_s$ -measurable one-sided optimism-transfer factor  $u_s(a) \geq 1$  satisfying  $\bar{s}_s^2(a) \leq u_s(a) s_{C,s}^2(a)$  (Assumption 5). Suppose also that the nonnegative predictable sequence  $\bar{F}_t$  satisfies  $\bar{F}_t \geq F_t$ . Define the observable information-gain surrogate*

$$\hat{\gamma}_{t-1} := \sum_{s=1}^{t-1} \log\left(1 + \sigma^{-2} \frac{u_s(a_s)}{1 - \bar{\epsilon}_s} \bar{s}_s^2(a_s)\right),$$

where  $\bar{s}_s^2(a_s) = \mathbf{g}_s(a_s)^\top \tilde{\mathbf{u}}_s$  is the CG variance proxy computed at round  $s$  for the played action. Then  $\hat{\gamma}_{t-1}$  is  $\mathcal{H}_t$ -measurable and  $\hat{\gamma}_{t-1} \geq \gamma_{t-1}$ , so

$$\hat{\beta}_t := \sqrt{\hat{\gamma}_{t-1} + 2\log(1/\delta)} + \sqrt{\lambda}S + \sigma^{-1}\sqrt{\bar{F}_t}$$

is a predictable schedule with  $\hat{\beta}_t \geq \beta_t$ . Played-action CG accuracy is sufficient for this information comparison. If the same tolerance also holds uniformly for every candidate action before selection, then Theorem 1 holds with  $\bar{\beta}_t = \hat{\beta}_t$ .

*Proof.* The transfer factor gives  $\bar{s}_s^2(a_s) \leq u_s(a_s) s_{C,s}^2(a_s)$ , and Lemma 22 with  $\epsilon_{s,a_s} \leq \bar{\epsilon}_s$  gives  $s_{C,s}^2(a_s) \leq \bar{s}_s^2(a_s)/(1 - \bar{\epsilon}_s)$ ; hence  $\bar{s}_s^2(a_s) \leq u_s(a_s) \bar{s}_s^2(a_s)/(1 - \bar{\epsilon}_s)$ . By monotonicity of  $x \mapsto \log(1 + \sigma^{-2}x)$ , each summand of  $\gamma_{t-1}$  in (28) is dominated by the corresponding summand of  $\hat{\gamma}_{t-1}$ . Measurability holds because  $\bar{s}_s^2(a_s)$ ,  $u_s(a_s)$ , and  $\bar{\epsilon}_s$  are all  $\mathcal{H}_s$ -measurable. The schedule property then follows term by term as in Corollary 25.  $\square$

This makes the information-gain component observable when  $u_s(a_s)$  is certified. Together with Lemma 11, it renders the full confidence and centering schedules operational; it does not make them tight.

**Remark** (Computability and dimension dependence of  $\beta_t$ ). *Two caveats about  $\beta_t$  deserve emphasis. First,  $\gamma_{t-1}$  is defined through the frozen-feature Gram matrix  $\bar{\mathbf{C}}_t$ , whose per-sample Jacobians the matrix-free implementation deliberately does not store;  $\gamma_{t-1}$  is therefore not directly computable in the stated memory model. Second, the worst-case substitute  $\gamma_{t-1} \leq d \log(1 + (t-1)G^2/(d\lambda\sigma^2))$  does not exploit low realized information gain and can be order  $t$  when  $d \gg t$ , rendering the resulting bonus and theorem bound vacuous. Corollary 26 addresses both issues with an observable surrogate built from quantities the algorithm already computes; the resulting guarantee is data-dependent, and its tightness still depends on the realized surrogate together with the other theorem factors.*

**Lemma 27** (Action-scaled centering discrepancy). *Under Assumption 4,  $|\mathbf{g}_t(a)^\top(\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{\psi}_t \bar{s}_t(a)$  for all  $a \in \mathcal{A}_t$ .*

*Proof.* Immediate from (36).  $\square$

*Remark (status of the centering certificate  $\bar{\psi}_t$ ).* Lemma 11 computes a predictable  $\bar{\psi}_t \geq \psi_t$  from the certified optimizer residual, Welford path moments, observed collection residuals, and analytic smoothness constants. The schedule is operational but can be loose; the corrected-center variant (Corollary 8) removes it at the cost of frozen-feature access.

In the convex last-layer regime, under exact optimization,  $\boldsymbol{\theta}_t = \hat{\boldsymbol{\theta}}_t^{\text{lin}}$  and  $\bar{\psi}_t = 0$ . The full prediction error at  $(\mathbf{x}_t, a)$  is bounded by the triangle inequality, with the centering discrepancy *action-scaled*:

$$|\mu_{\boldsymbol{\theta}^*}(\mathbf{x}_t, a) - \mu_{\boldsymbol{\theta}_t}(\mathbf{x}_t, a)| \leq (\beta_t + \bar{\psi}_t) \bar{s}_t(a) + \epsilon_{\text{lin}}(t) \leq \omega_t \bar{s}_t(a) + \epsilon_{\text{lin}}(t), \quad (73)$$

using (71), Lemma 27, the per-round linearization remainder  $|\rho_t| \leq \epsilon_{\text{lin}}(t)$ , and  $\bar{\beta}_t \geq \beta_t$ .

## A.8 Cost of Approximate Curvature

**Proposition 28** (Unrescaled current-parameter subset). *Let  $\mathcal{S}_t \subseteq \{1, \dots, t-1\}$  and  $0 \leq w_s \leq 1$ , and define the current-parameter operator*

$$\hat{\mathbf{C}}_t = \lambda \mathbf{I} + \sigma^{-2} \sum_{s \in \mathcal{S}_t} w_s \mathbf{g}_t(\mathbf{x}_s, a_s) \mathbf{g}_t(\mathbf{x}_s, a_s)^\top,$$

*using the same relinearized Jacobians  $\mathbf{g}_t(\mathbf{x}_s, a_s) = \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}_t}(\mathbf{x}_s, a_s)$  that define  $\mathbf{C}_t$  (5). Then  $\hat{\mathbf{C}}_t \preceq \mathbf{C}_t$ , so Theorem 29 applies with  $\kappa_{+,t} = 1$  and, in the exact-transfer case,  $u_t = (1 + \bar{\chi}_t)^2$  (Lemma 16).*

*Proof.*  $\mathbf{C}_t - \hat{\mathbf{C}}_t = \sigma^{-2} \sum_{s \in \mathcal{S}_t} (1 - w_s) \mathbf{g}_t \mathbf{g}_t^\top + \sigma^{-2} \sum_{s \notin \mathcal{S}_t} \mathbf{g}_t \mathbf{g}_t^\top$ , a sum of PSD outer products (each  $1 - w_s \geq 0$ ), hence  $\mathbf{C}_t - \hat{\mathbf{C}}_t \succeq 0$ . With  $\hat{\mathbf{C}}_t \preceq \mathbf{C}_t$  (so  $\kappa_{+,t} = 1$  in (74)) and  $\hat{\mathbf{C}}_t \succeq \lambda \mathbf{I}$ , the hypotheses of Theorem 29 hold.  $\square$

This covers an unrescaled current-parameter sliding window, an unrescaled current-parameter subsample, and any current-parameter downweighted subset with weights in  $[0, 1]$ . It does *not* cover rescaled subsampling with weights  $(t-1)/n > 1$  (Lemma 18), a stale periodic-refresh operator evaluated at an older parameter, or an arbitrary randomized sketch (e.g. an eigenspace Lanczos surrogate), for which a factor  $\kappa_{+,t}$  must be certified separately.

**Theorem 29** (Approximate curvature via one-sided transfer). *Adopt the hypotheses of Theorem 1 with the algorithmic curvature  $\mathcal{C}_t = \hat{\mathbf{C}}_t$ , an approximate operator that is symmetric,  $\mathcal{H}_t$ -measurable, and satisfies  $\hat{\mathbf{C}}_t \succeq \lambda \mathbf{I}$ . Suppose the two one-sided bounds*

$$\hat{\mathbf{C}}_t \preceq \kappa_{+,t} \mathbf{C}_t, \quad \mathbf{C}_t \preceq (1 + \bar{\chi}_t)^2 \bar{\mathbf{C}}_t \quad (74)$$

*hold for every  $t \leq T$ , where  $\kappa_{+,t} \geq 1$  is  $\mathcal{H}_t$ -measurable and  $\bar{\chi}_t \geq \chi_t$  is the whitened-drift certificate of Lemma 16. Then, for every action  $a$ ,*

$$\bar{s}_t^2(a) \leq \kappa_{+,t} (1 + \bar{\chi}_t)^2 \hat{s}_t^2(a), \quad (75)$$

*so Assumption 5 holds with the uniform factor  $u_t = \kappa_{+,t} (1 + \bar{\chi}_t)^2$  and Theorem 1 applies verbatim with  $\mathcal{C}_t = \hat{\mathbf{C}}_t$  and  $s_{\mathcal{C},t} = \hat{s}_t$ :*

$$R_T \leq 2 \sqrt{(\sigma^2 + G^2/\lambda) \Lambda_T^{\mathcal{C}} \sum_{t=1}^T \alpha_t^2 u_t(a_t) \omega_t^2} + 2E_T.$$

*No lower spectral factor on  $\hat{\mathbf{C}}_t$  is required for optimism; the cost of an excessively small, stale, or nonmonotone approximate operator is charged through its realized width complexity  $\Lambda_T^{\mathcal{C}}$  (and its variation  $V_T^{\mathcal{C}}$  inside  $\Gamma_T^{\text{dyn}}$ ), not through a lower Loewner bound.*

*Proof.*  $\hat{\mathbf{C}}_t \preceq \kappa_{+,t} \mathbf{C}_t$  with both sides SPD inverts to  $\hat{\mathbf{C}}_t^{-1} \succeq \kappa_{+,t}^{-1} \mathbf{C}_t^{-1}$ , so  $\hat{s}_t^2(a) \geq \kappa_{+,t}^{-1} s_t^2(a)$ ; and (51) gives  $\bar{s}_t^2(a) \leq (1 + \bar{\chi}_t)^2 s_t^2(a)$ . Chaining,  $\bar{s}_t^2(a) \leq (1 + \bar{\chi}_t)^2 \kappa_{+,t} \hat{s}_t^2(a)$ , which is (75) with  $s_{\mathcal{C},t} = \hat{s}_t$ . Theorem 1 applies with  $\mathcal{C}_t = \hat{\mathbf{C}}_t$ . The two-sided  $\sqrt{\rho_+/\rho_-}$  inflation—which instead assumes a lower spectral factor—is retained for comparison in Appendix F; the two results trade different assumptions (a lower spectral bound versus dynamic width complexity), and neither numerically dominates the other in general.  $\square$

**Proposition 30** (Roundwise scalar-width invariance). *If two width functions satisfy  $\hat{s}_t(a) = c_t s_t(a)$  for every candidate action, where  $c_t > 0$  is action independent, then replacing a scalar bonus coefficient  $b_t$  by  $b_t/c_t$  gives identical UCB scores and therefore the same selected action under the same tie rule.*

*Proof.*  $\hat{\mu}_t(a) + (b_t/c_t) \hat{s}_t(a) = \hat{\mu}_t(a) + b_t s_t(a)$  for every  $a$ .  $\square$

This is a roundwise statement. If two policies also have identical centers, histories, optimization, and coupled exogenous randomness, if  $c_t$  is predictable, and if proportionality holds at every subsequent common history, induction gives identical actions and rewards after the coefficient adjustment. Absent those conditions the proposition makes no regret claim. Accordingly, regret differences can reflect anisotropic width distortion or a difference in damping, center, adaptive path, optimization, calibration, or solve accuracy.

**Proposition 31** (Off-diagonal linear-Gram witness). *Let  $\sigma = 1$ ,  $\mathbf{u} = (1, 1)/\sqrt{2}$ ,  $\mathbf{v} = (1, -1)/\sqrt{2}$ , the action set be  $\{\mathbf{u}, \mathbf{v}\}$ , and  $\boldsymbol{\theta}^* = \epsilon \mathbf{u} + \Delta \mathbf{v}$  with  $\Delta > \epsilon > 0$ . Rewards are noiseless, the ridge is  $\lambda \mathbf{I}$ , all policies use the same full-ridge mean center, and ties select  $\mathbf{u}$ . If*

$$\max \left\{ \sqrt{\lambda(\epsilon^2 + \Delta^2)}, \epsilon \sqrt{\lambda} \left( 1 + \sqrt{\frac{\lambda}{\lambda + 1}} \right) \right\} < b \leq \left( \Delta - \frac{\epsilon}{\lambda + 1} \right) \sqrt{\lambda + 1}, \quad (76)$$

*then full-Gram UCB plays  $\mathbf{u}$  once and  $\mathbf{v}$  thereafter, so  $R_T^{\text{full}} = \Delta - \epsilon$ . The raw diagonal policy and the policy inflated by the uniform analytic factor  $\kappa_n = (\lambda + n/2)/\lambda$  play  $\mathbf{u}$  forever and have  $R_T^{\text{diag}} = T(\Delta - \epsilon)$ .*

*Proof.* After  $n$  pulls of  $\mathbf{u}$  and no pull of  $\mathbf{v}$ ,

$$C_n = \lambda \mathbf{I} + n \mathbf{u} \mathbf{u}^\top, \quad \hat{\boldsymbol{\theta}}_n = \frac{n\epsilon}{\lambda + n} \mathbf{u}.$$

Thus the full scores are  $n\epsilon/(\lambda + n) + b/\sqrt{\lambda + n}$  and  $b/\sqrt{\lambda}$ . The strict lower inequality in (76) makes the second score larger at  $n = 1$ . The confidence lower bound also implies  $b > \Delta\sqrt{\lambda}$ ; after any number of  $\mathbf{v}$  pulls its score is strictly larger than  $\Delta$ , whereas the upper inequality makes the fixed  $\mathbf{u}$  score at most  $\Delta$ . Hence there is no return. In contrast,  $D_n = \text{diag}(C_n) = (\lambda + n/2)\mathbf{I}$ , so both diagonal widths are equal and the center gives  $\mathbf{u}$  a strict advantage after round one. Moreover  $D_n \preceq \kappa_n C_n$  analytically, but multiplying both diagonal widths by  $\sqrt{\kappa_n}$  leaves them equal; uniform certified inflation therefore does not change the diagonal action. For example,  $(\lambda, \epsilon, \Delta, b) = (1, 0.1, 1, 1.1)$  satisfies the interval.  $\square$

This is a two-dimensional linear-Gram existence witness. For a linear predictor with Gaussian squared loss, this Gram is exactly the GGN/Fisher operator, so the construction isolates a GGN-relevant geometric mechanism. It is neither a nonlinear-network result nor a uniform dominance result. A dense action-dependent factor recovers the full widths, but it uses the reference geometry that the diagonal approximation was meant to avoid.

**Composing subsampling with feature drift.** Lemma 18 guarantees a sandwich for the *relinearized*  $\mathbf{C}_t$ :  $(1 - \epsilon)\mathbf{C}_t \preceq \hat{\mathbf{C}}_t \preceq (1 + \epsilon)\mathbf{C}_t$ . Only the *upper* half is needed here:  $\hat{\mathbf{C}}_t \preceq (1 + \epsilon)\mathbf{C}_t$  gives  $\kappa_{+,t} = 1 + \epsilon$ , and composing with the one-sided whitened bound  $\mathbf{C}_t \preceq (1 + \bar{\chi}_t)^2 \bar{\mathbf{C}}_t$  (Lemma 16) yields  $u_t = (1 + \epsilon)(1 + \bar{\chi}_t)^2$  in Theorem 29. A union bound over  $t \leq T$  (replacing  $\delta'$  by  $\delta'/T$  in Lemma 18) makes the subsampling event time-uniform and contributes to  $\delta_{\text{cert}}$ .

For an *unrescaled current-parameter window* (or any current-parameter downweighted subset with weights in  $[0, 1]$ ), Proposition 28 gives the *deterministic* upper factor  $\kappa_{+,t} = 1$  with *no* distributional assumption, so Theorem 29 applies directly with  $u_t = (1 + \bar{\chi}_t)^2$ . Only the *stale periodic-refresh* operator (§A.3(iii)), which reuses Jacobians from an older parameter, lacks a deterministic  $\kappa_{+,t}$ ; there Theorem 29 provides a guarantee only when a one-sided bound  $\hat{\mathbf{C}}_t \preceq \kappa_{+,t} \mathbf{C}_t$  is separately verified, which we evaluate empirically.

**What remains open.** The path schedules close the computability gap, but their constants can be extremely loose. The approximate-rank result replaces exact rank by a realized or predictably bounded spectral tail, but does not show that this tail or feature drift remains small during unrestricted training. Useful control of these two quantities for a genuinely adapting backbone, tighter smoothness/residual certificates, and verified numerical enclosures for floating-point solver and spectral calculations remain open.

## B Related Work

**Neural and linear contextual bandits.** LinUCB (Abbasi-Yadkori et al., 2011) and Thompson sampling (Thompson, 1933; Agrawal and Goyal, 2013) provide the analytical backbone for contextual bandit exploration. Neural extensions—NeuralUCB (Zhou et al., 2020), NeuralTS (Zhang et al., 2021), neural-linear models (Riquelme et al., 2018)—use different neural feature covariances. NeuralUCB and NeuralTS accumulate full neural-gradient Grams in their stated algorithms (and diagonalize them in experiments), while NeuralLinear maintains Bayesian linear uncertainty over learned last-layer features. CC-UCB instead re-evaluates every replayed Jacobian at the current parameter and accesses the resulting full operator implicitly. Neural-linear designs that train a deep representation but explore only in the last layer (Xu et al., 2022a) occupy the opposite design point.



---

Recent alternatives sharpen different parts of this landscape. Salgia (2023) analyze overparameterized smooth-activation networks through non-asymptotic network–NTK approximation and give an efficient sublinear-regret algorithm. Deb et al. (2024) instead reduce contextual bandits to online neural regression, using perturbed predictions to obtain their regression guarantees. Bui et al. (2025) combine a learned deep representation with a variance-aware linear-UCB exploration layer. These are, respectively, an NTK analysis, a regression-oracle route, and a learned last-layer geometry; none analyzes the dynamically relinearized full GGN and its CG error. Our low-rank/excitation result is correspondingly narrower than generic neural regression, but controls that current-parameter operator.

**Online Bayesian and EKF-style neural bandits.** Durán-Martín et al. (2022) apply extended Kalman filter (EKF) updates in a learned or random low-dimensional affine parameter subspace. They maintain a dense or diagonal covariance in that subspace, inducing a low-rank covariance over parameters spanning all network layers, with memory independent of the number of rounds. This differs from applying exact curvature–vector products in the full parameter space. More directly, Chang et al. (2023) maintain a diagonal-plus-low-rank approximation to the full-parameter posterior precision through recursive EKF and online-SVD updates, and use it for Thompson sampling in an MNIST contextual bandit. Their LO-FI method is a one-pass compressed-posterior approach with fixed-rank cost linear in parameter count and no replay of past data. CC-UCB instead relinearizes the full replay history at the current parameter and accesses the uncompressed GGN through residual-checked CG. Their bandit study is empirical; our conditional UCB analysis explicitly accounts for CG truncation, but does not establish a practical advantage over LO-FI.

**Last-layer and structured Laplace approximations.** Last-layer Laplace fixes the backbone and restricts the approximate posterior to the final layer (Kristiadi et al., 2020), omitting backbone-weight uncertainty by construction. This is only one scalable option: all-layer block-diagonal Kronecker curvature has also been used (Ritter et al., 2018), and Daxberger et al. (2021) systematize all-weight, subnetwork, and last-layer variants with full, low-rank, KFAC, or diagonal curvature. CC-UCB uses the uncompressed current-parameter GGN as its reference operator without materializing the covariance matrix.

**Sampling-based full-parameter exploration.** Langevin Monte Carlo Thompson sampling (LMC-TS) (Xu et al., 2022b) explores with approximate samples of a loss-defined parameter posterior, obtained by noisy gradient steps on the regularized loss; like CC-UCB it never materializes a covariance matrix. Its regret guarantee is for linear contextual bandits; its neural-bandit evidence is empirical. The two routes pay for exploration differently: LMC-TS through the mixing of a per-round sampling chain, CC-UCB through CG truncation and spectral distortion of an explicit curvature metric, which yields deterministic widths, an optimism-style analysis, and a fixed curvature operator reused across separate per-action CG solves.

**Efficient nonlinear UCB.** Kassraie and Krause (2022) analyze NTK-UCB and propose NN-UCB and its convolutional variant CNN-UCB. NN-UCB combines a trained-network mean with confidence widths from the empirical NTK feature map at initialization. They prove sublinear regret for the SupNN-UCB variant under NTK-RKHS, sufficient-exploration, and network-width assumptions. Their fixed-feature uncertainty differs from CC-UCB’s replayed GGN/Fisher curvature at the current parameter  $\theta_t$ , accessed matrix-free via CG. CC-UCB thus targets the regime where the practitioner has a specific pretrained model and wants to use its curvature as-is, rather than constructing a surrogate NTK-based estimator.

**Spectral sketching for linear bandits.** Kuzborskij et al. (2019) analyze linear bandits where the Gram matrix is replaced by a Frequent Directions sketch, deriving regret bounds that depend on the spectral approximation quality. Wen et al. (2026) show that fixed-size sketch regret bounds can degenerate to the trivial linear rate under heavy spectral tails, and propose an adaptive dyadic block sketch. Our spectral-distortion theorem (Theorem 29) instantiates the same conceptual lever in the matrix-free GGN/CVP setting by exposing transfer and realized-width terms along the realized trajectory. We do not claim novelty for the sandwich-to-regret reduction itself; the contribution is adapting it to the CG/GGN context with explicit CG truncation and feature-drift factors.

**Laplace approximation in deep learning.** The Laplace approximation (MacKay, 1992; Ritter et al., 2018) provides a Gaussian posterior centered at a MAP estimate with covariance given by the inverse curvature. Scalable variants use diagonal, KFAC (Martens and Grosse, 2015), or subnetwork approximations (Daxberger et al., 2021).

Matrix-free full-Jacobian/Fisher predictive variance using implicit products and CG predates this work (Maddox et al., 2021); CC-UCB adapts that computational route to an online, current-relinearized GGN operator and supplies bandit-specific transfer and solver-error accounting.

**GGN–vector products, Hessian spectral analysis, and influence functions.** Pearlmutter’s algorithm (Pearlmutter, 1994) computes exact Hessian–vector products; the analogous GGN–vector product has the same Jacobian-transpose-times-Jacobian-times-vector structure (Schraudolph, 2002). Recent work on Krylov methods for Hessian spectral analysis (Granzio and Juarev, 2026) demonstrates related distributed HVP/Krylov infrastructure at  $> 10^8$ -parameter scale, but does not benchmark the exact per-example GGN operator or end-to-end CC-UCB latency. Influence functions (Koh and Liang, 2017) use inverse-curvature–vector products to attribute predictions to training points; the CG-based implicit solve we employ is the same computational primitive.

**Optimal experimental design.** Information-theoretic exploration objectives—such as maximizing  $\log \det(\bar{\mathbf{C}}_t)$  or minimizing the trace of  $\bar{\mathbf{C}}_t^{-1}$ —connect bandit exploration to the classical D-optimal and A-optimal design criteria (Valko et al., 2013). CC-UCB’s use of the full GGN/Fisher curvature enables stochastic estimates of these quantities via Hutchinson-type trace estimators (Ubaru et al., 2017), without forming the matrix.

**Filtration.** We separate *pre-action* randomness, revealed before the agent commits to  $a_t$ , from *post-reward* randomness, drawn only after  $r_t$  is observed. Let  $\Omega_t$  collect all pre-action algorithmic randomness used at round  $t$ : curvature subsampling indices, randomized sketches, and stochastic solver seeds. The randomized sketch or subsample that defines the curvature operator is drawn into  $\Omega_t$  and *fixed before action selection*; it is not re-randomized during any conjugate-gradient (CG) solve at round  $t$ . (Standard CG requires a fixed symmetric positive definite matrix–vector oracle across its iterations and does not apply to a stochastic oracle that changes between iterations.) Let  $U_{t+1}$  collect the *post-reward* randomness used to form the next parameter estimate  $\theta_{t+1}$  (minibatch order, dropout, perturbed warm starts). Define  $\mathcal{F}_0 = \sigma(\theta_1)$  and, for  $t \geq 1$ , the recursively built  $\sigma$ -algebra

$$\mathcal{F}_t := \mathcal{F}_{t-1} \vee \sigma(\mathbf{x}_t, \mathcal{A}_t, \Omega_t, a_t, r_t, U_{t+1}), \quad (77)$$

so that  $\theta_{t+1}$ —computed from the data through round  $t$  together with the post-reward randomness  $U_{t+1}$ —is  $\mathcal{F}_t$ -measurable. Define the *pre-action history*  $\sigma$ -algebra and the pre-reward  $\sigma$ -algebra

$$\mathcal{H}_t^- := \mathcal{F}_{t-1} \vee \sigma(\mathbf{x}_t, \mathcal{A}_t), \quad \mathcal{H}_t := \mathcal{H}_t^- \vee \sigma(\Omega_t), \quad \mathcal{G}_t := \mathcal{H}_t \vee \sigma(a_t). \quad (78)$$

At round  $t$  the context  $\mathbf{x}_t$ , the finite action set  $\mathcal{A}_t$ , and the pre-action randomness  $\Omega_t$  are revealed; the agent selects  $a_t$ ; then the reward  $r_t = \mu^*(\mathbf{x}_t, a_t) + \eta_t$  is observed; finally  $U_{t+1}$  is drawn and  $\theta_{t+1}$  is computed.  $\mathcal{H}_t^-$  is the history immediately before the randomized certificate draw; after that draw, its realized operator and schedules are  $\mathcal{H}_t$ -measurable. By construction of Algorithm 1 the selection rule is a deterministic function of  $\mathcal{H}_t$ -measurable quantities, so the chosen action  $a_t$  is itself  $\mathcal{H}_t$ -measurable and hence  $\mathcal{G}_t = \mathcal{H}_t$ ; we retain the symbol  $\mathcal{G}_t$  where the played action is used. The following are all  $\mathcal{H}_t$ -measurable (i.e. determined before the reward  $r_t$ ): the candidate action set  $\mathcal{A}_t$  and every candidate feature  $\mathbf{g}_t(\mathbf{x}_t, a)$ ,  $a \in \mathcal{A}_t$ ; the chosen action  $a_t$ ; the algorithmic curvature operator  $\mathcal{C}_t$  (Section 3); the confidence schedule  $\beta_t$ ; the centering certificate  $\psi_t$  (Assumption 4); the per-round CG tolerance  $\bar{\varepsilon}_t$  together with any condition-number certificate used by the CG stopping rule; and the one-sided optimism-transfer factor  $u_t(a)$  (Assumption 5). Since  $\theta_t$  is  $\mathcal{F}_{t-1} \subseteq \mathcal{H}_t$ -measurable, each candidate design vector  $\mathbf{g}_t(\mathbf{x}_t, a)$  is  $\mathcal{H}_t$ -measurable and the realized design vector  $\mathbf{g}_t(\mathbf{x}_t, a_t)$  is  $\mathcal{G}_t$ -measurable. We assume  $\{U_s\}_{s \geq 1}$  is independent of the reward noise  $\{\eta_s\}_{s \geq 1}$ , so that revealing  $U_{t+1}$  and forming  $\theta_{t+1}$  after  $r_t$  preserves the conditional sub-Gaussian property of Assumption 1 for the future noise  $\{\eta_s\}_{s > t}$ . Index every randomized certificate source (operator sketch, transfer bound, condition bound, or other randomized oracle) by  $j \in \mathcal{J}$ . For its round- $t$  validity event  $\mathcal{E}_{j,t}$  choose in advance  $\delta_{j,t} \geq 0$  such that its stated conditional guarantee gives  $\Pr(\mathcal{E}_{j,t}^c \mid \mathcal{H}_t^-) \leq \delta_{j,t}$  and

$$\sum_{j \in \mathcal{J}} \sum_{t=1}^T \delta_{j,t} \leq \delta_{\text{cert}}, \quad \mathcal{E}_{\text{cert}} := \bigcap_{j \in \mathcal{J}} \bigcap_{t=1}^T \mathcal{E}_{j,t}.$$

The tower property and union bound yield  $\Pr(\mathcal{E}_{\text{cert}}) \geq 1 - \delta_{\text{cert}}$ . Deterministic or almost-sure certificates use  $\delta_{j,t} = 0$ . The self-normalized confidence event is separate and receives probability  $\delta$ . When Lemma 10 supplies the residual-energy envelope, include its time-uniform event  $\mathcal{E}_R$  as one certificate source and reserve  $\delta_R$  inside

$\delta_{\text{cert}}$  (equivalently, require  $\delta_R + \sum_{j,t} \delta_{j,t} \leq \delta_{\text{cert}}$  for the remaining sources). The residual event is never counted again in the self-normalized failure probability  $\delta$ .

**Auxiliary linearized ridge estimator.** For the analysis we define a *linearized ridge estimator* based on frozen features. Set the predictable intercept  $b_s := \mu_{\theta_s}(\mathbf{x}_s, a_s) - \mathbf{g}_s(\mathbf{x}_s, a_s)^\top \theta_s$  and the pseudo-response  $y_s := r_s - b_s$ . Under Assumption 3 and realizability ( $\mu^* = \mu_{\theta^*}$ ),

$$y_s = \mathbf{g}_s(\mathbf{x}_s, a_s)^\top \theta^* + \rho_s(\mathbf{x}_s, a_s, \theta^*) + \eta_s, \quad (79)$$

where  $|\rho_s| \leq \varepsilon_{\text{lin}}(s)$ . The linearized estimator is

$$\hat{\theta}_t^{\text{lin}} := \arg \min_{\theta} \left[ \frac{1}{2\sigma^2} \sum_{s=1}^{t-1} (y_s - \mathbf{g}_s(\mathbf{x}_s, a_s)^\top \theta)^2 + \frac{\lambda}{2} \|\theta\|^2 \right] = \bar{\mathbf{C}}_t^{-1} \frac{1}{\sigma^2} \sum_{s=1}^{t-1} \mathbf{g}_s(\mathbf{x}_s, a_s) y_s. \quad (80)$$

This is a standard linear ridge regression with predictable design vectors  $\mathbf{g}_s/\sigma$  and Gram matrix  $\bar{\mathbf{C}}_t$ , so the self-normalized martingale bound (Abbasi-Yadkori et al., 2011) applies to it directly (Lemma 24).

**GGN versus Hessian (context only).** The regret theorem uses the GGN/Fisher curvature as the exploration metric. It does not require the GGN to be an accurate Hessian approximation. For squared loss (3) the full Hessian of the data-fit term is

$$\mathbf{H}_t = \frac{1}{\sigma^2} \sum_{s=1}^{t-1} \left[ \mathbf{g}_t(\mathbf{x}_s, a_s) \mathbf{g}_t(\mathbf{x}_s, a_s)^\top - (r_s - \mu_{\theta_t}(\mathbf{x}_s, a_s)) \nabla^2 \mu_{\theta_t}(\mathbf{x}_s, a_s) \right] + \lambda \mathbf{I}.$$

The curvature matrix  $\mathbf{C}_t$  drops the residual-times-second-derivative term, retaining only the outer-product part that is guaranteed PSD. The dropped residual-times-second-derivative term  $D_t = \sigma^{-2} \sum_{s < t} (r_s - \mu_{\theta_t}) \nabla^2 \mu_{\theta_t}$  has operator norm at most  $(t-1)\epsilon_r H/\sigma^2$ , where  $\epsilon_r$  bounds the residuals and  $H$  the Hessian norms. When this quantity is small relative to  $\lambda$ ,  $\mathbf{C}_t$  approximates the full Hessian in the Loewner sense (a formal statement with proof is given in Appendix G). This comparison is *not* invoked in the regret analysis (Theorem 1 uses  $\mathbf{C}_t$  directly as the exploration metric); it serves only to contextualize  $\mathbf{C}_t$  as a proxy for the full Hessian. Outside the small-residual regime, the GGN is not an exact Hessian approximation; it is the PSD Fisher/GGN metric used by the algorithm.

*Remark (Fisher, empirical Fisher, and GGN).* Three curvature objects appear in the literature and must be distinguished. The *empirical Fisher* is the realized score outer product  $\sum_s \nabla_{\theta} \log p(r_s | \mathbf{x}_s, a_s; \theta) \nabla_{\theta} \log p(r_s | \mathbf{x}_s, a_s; \theta)^\top$ ; it depends on observed rewards and is *not* the operator used by CC-UCB. The *true (model) Fisher* is the model expectation of the score outer product,  $F(\theta) = \mathbb{E}_{r \sim p(\cdot | \mathbf{x}, a; \theta)} [\nabla_{\theta} \log p \nabla_{\theta} \log p^\top]$ . The *generalized Gauss-Newton* (GGN) drops the residual-times-Hessian term from the full Hessian, retaining only the outer-product Jacobian part that is guaranteed PSD (Martens, 2020; Schraudolph, 2002). For Gaussian regression with fixed variance (the squared-loss setting of this paper), the true Fisher and GGN are identical, and both reduce to the outer-product curvature  $\sigma^{-2} \sum_s \mathbf{g}_s \mathbf{g}_s^\top$  used in (5). More generally, for exponential-family negative log-likelihoods parameterized by natural parameters  $z_{\theta}(\mathbf{x}, a)$ , the true Fisher equals the GGN:  $J_{\theta}^\top \nabla^2 A(z) J_{\theta}$ , where  $J_{\theta} = \nabla_{\theta} z_{\theta}$  and  $\nabla^2 A$  is the Hessian of the log-partition function (Martens, 2020; Pascanu and Bengio, 2013). Bernoulli rewards yield scalar weight  $p(1-p)$ ; categorical rewards yield  $\text{Diag}(\mathbf{p}) - \mathbf{p}\mathbf{p}^\top$ . See Appendix H for the extension.

We access  $\mathbf{C}_t$  only through curvature-vector products  $\mathbf{v} \mapsto \mathbf{C}_t \mathbf{v}$ , computed via forward-mode Jacobian-vector products composed with reverse-mode vector-Jacobian products (Pearlmutter, 1994; Schraudolph, 2002).

**Extension to exponential-family likelihoods.** The same matrix-free CG/CVP machinery extends algorithmically to exponential-family likelihoods (e.g. Bernoulli or categorical rewards) via the Fisher/GGN curvature  $J_t^\top (\nabla^2 A) J_t$ , but no regret guarantee beyond Gaussian squared loss is claimed in this draft. Details of the curvature definitions are in Appendix H.

**Per-round CG tolerance  $\bar{\varepsilon}_t$ .** The algorithm uses a *known* per-round tolerance  $\bar{\varepsilon}_t \in (0, 1)$  in the UCB bonus (line 8), not the true per-action energy-norm errors  $\varepsilon_{t,a}$  (which are unobservable without the exact solves). It suffices that  $\max_{a \in \mathcal{A}_t} \varepsilon_{t,a} \leq \bar{\varepsilon}_t$  at every round ((10); zero-gradient actions have  $\varepsilon_{t,a} = 0$  by convention and are trivially within tolerance); the bonus is then conservative. Two strategies ensure this:

- 
1. *Adaptive residual stopping (recommended).* In exact arithmetic, for each action monitor the computable Euclidean relative residual  $\|r_I\|/\|\mathbf{g}_t(a)\|$  at each CG step (skipping actions with  $\mathbf{g}_t(a) = \mathbf{0}$ ). Since  $\varepsilon_{t,a} \leq \sqrt{\bar{\kappa}_t} \cdot \|r_I\|/\|\mathbf{g}_t(a)\|$  (standard CG bound, with  $\bar{\kappa}_t$  a certified upper bound on the condition number of  $\mathcal{C}_t$ ), stopping that solve when  $\|r_I\|/\|\mathbf{g}_t(a)\| \leq \bar{\varepsilon}_t/\sqrt{\bar{\kappa}_t}$  guarantees  $\varepsilon_{t,a} \leq \bar{\varepsilon}_t$ , hence (10). This is conservative by the factor  $\sqrt{\bar{\kappa}_t}$  but requires no access to  $\mathcal{C}_t^{-1}$  and adapts automatically to the per-round condition number. A floating-point implementation needs a verified residual enclosure to inherit this certification literally.
  2. *Fixed budget (requires a uniform condition-number bound).* If an a priori bound  $\bar{\kappa}_t \leq \kappa_{\max}$  holds for all rounds  $t \leq T$ , choose  $I$  via Eq. (85) at  $\kappa_{\max}$  and a fixed target  $\bar{\varepsilon}_t \equiv \bar{\varepsilon}$ . For a *nonnegative weighted outer-product* operator  $\mathcal{C}_t = \lambda \mathbf{I} + \sigma^{-2} \sum_{s \in \mathcal{S}_t} w_s \mathbf{g}_t(\mathbf{x}_s, a_s) \mathbf{g}_t(\mathbf{x}_s, a_s)^\top$  the condition number is bounded in closed form,  $\bar{\kappa}_t \leq 1 + W_t G^2 / (\lambda \sigma^2)$  with  $W_t = \sum_{s \in \mathcal{S}_t} w_s$ ; a fixed buffer of size  $n$  ( $W_t = n$ , unrescaled) makes this independent of  $T$ . For the full-history operator ( $W_t = t-1$ ), and for rescaled uniform subsampling with  $w_s = (t-1)/n$  (so  $W_t = \sum_{s \in \mathcal{S}_t} w_s = t-1$ ), the bound grows with  $t$ . For a known finite horizon, a fixed budget at  $\kappa_{\max} = 1 + (T-1)G^2/(\lambda \sigma^2)$  is valid but scales as  $O(\sqrt{T} \log(1/\bar{\varepsilon}))$ ; a horizon-independent fixed budget is not uniform, so adaptive stopping or an increasing  $I_t$  is preferable (see Appendix E). For a *general SPD* operator that is not a nonnegative weighted outer product (e.g. a randomized sketch), the closed-form bound  $1 + W_t G^2 / (\lambda \sigma^2)$  does not apply, and a separately certified  $\mathcal{H}_t$ -measurable condition-number upper bound  $\bar{\kappa}_t$  must be supplied.

In the regret bound (Theorem 1), the CG inflation enters through the per-round factor  $\alpha_t = \sqrt{(1+\bar{\varepsilon}_t)/(1-\bar{\varepsilon}_t)}$ , inside the sum of squares  $\sum_t \alpha_t^2 u_t(a_t) \omega_t^2$ ; no horizon maximum over  $\bar{\varepsilon}_t$  is taken.

**Action screening (practical heuristic).** When the action set  $\mathcal{A}_t$  is large, computing the CG solve for every action is prohibitive ( $|\mathcal{A}_t| \times I$  CG steps per round). A practical heuristic is to screen the top- $K$  actions by mean prediction  $\hat{\mu}_t(a)$  (forward passes only), then compute UCB widths for the  $K$  candidates. This reduces cost to  $K \times I$  CG solves. *The regret guarantee in Theorem 1 analyzes the full-enumeration case ( $K = |\mathcal{A}_t|$ ); screening can drop the optimal arm and is not covered by the theory unless one adds an assumption that  $a_t^* \in \mathcal{A}_t^K$  for all  $t$ .*

---

**Algorithm 2** Operational CC-UCB with path certificates

---

**Require:** Reward predictor  $\mu_\theta$ , damping  $\lambda > 0$ , noise proxy  $\sigma$ , confidence level  $\delta$ , radius  $S$ , analytic constants  $G, L_g, L_\mu$ ; algorithmic curvature mode  $\mathcal{C}_t$  (exact full-history  $\mathbf{C}_t$  or approximate  $\hat{\mathbf{C}}_t$ ) with certified one-sided factor  $\kappa_{+,t}$  and condition upper bound  $\bar{\kappa}_t$ ; per-round CG tolerance  $\bar{\varepsilon}_t < 1$

- 1: Initialize  $\theta_1$ , certified  $\zeta_1 \geq \|\nabla \mathcal{L}_1(\theta_1)\|$ ,  $n \leftarrow 0$ ,  $\bar{\theta} \leftarrow \mathbf{0}$ ,  $J \leftarrow 0$ ,  $R^{\text{col}} \leftarrow 0$ ,  $\bar{F} \leftarrow 0$ ,  $\hat{\gamma} \leftarrow 0$ , and  $\mathcal{D}_0 \leftarrow \emptyset$
- 2: **for**  $t = 1, 2, \dots, T$  **do**
- 3:   Observe context  $\mathbf{x}_t$  and action set  $\mathcal{A}_t$ ; draw pre-action randomness  $\Omega_t$  { $\mathcal{H}_t$  of (78)}
- 4:    $Q_t \leftarrow J + n\|\theta_t - \bar{\theta}\|^2$ ;  $\bar{\chi}_t \leftarrow L_g\sqrt{Q_t}/(\sigma\sqrt{\lambda})$ ;  $u_t \leftarrow \kappa_{+,t}(1 + \bar{\chi}_t)^2$
- 5:    $\bar{M}_t \leftarrow \sigma^{-2}[L_g\sqrt{R^{\text{col}}Q_t} + (3GL_g/2)Q_t + (L_g^2/2)Q_t^{3/2}]$ ;  $\bar{\psi}_t \leftarrow (\zeta_t + \bar{M}_t)/\sqrt{\lambda}$  {in a radius- $R$  region use  $L_g^2RQ_t$  for the last term}
- 6:    $\bar{\epsilon}_t^{\text{lin}} \leftarrow (L_\mu/2)(S + \|\theta_t\|)^2$ ;  $\bar{\beta}_t \leftarrow \sqrt{\hat{\gamma} + 2\log(1/\delta)} + \sqrt{\lambda}S + \sigma^{-1}\sqrt{\bar{F}}$ ;  $\omega_t \leftarrow \bar{\beta}_t + \bar{\psi}_t$  {corrected center uses  $\omega_t = \bar{\beta}_t$ }
- 7:   Form or receive a fixed  $\mathcal{H}_t$ -measurable SPD matrix-vector oracle  $\mathbf{v} \mapsto \mathcal{C}_t\mathbf{v}$ ,  $\mathcal{C}_t \succeq \lambda\mathbf{I}$ , held fixed across all CG solves this round {exact/buffered special case  $\mathcal{C}_t\mathbf{v} = \lambda\mathbf{v} + \sigma^{-2}\sum_{s \in \mathcal{S}_t} w_s \mathbf{g}_t(\mathbf{x}_s, a_s)[\mathbf{g}_t(\mathbf{x}_s, a_s)^\top \mathbf{v}]$ , weights  $w_s$  per §A.3}
- 8:   **for** each  $a \in \mathcal{A}_t$  **do**
- 9:      $\hat{\mu}_t(a) \leftarrow \mu_{\theta_t}(\mathbf{x}_t, a)$  {mean prediction (or corrected center (32))}
- 10:     $\mathbf{g}_t(a) \leftarrow \nabla_{\theta} \mu_{\theta_t}(\mathbf{x}_t, a)$  {gradient feature}
- 11:    **if**  $\mathbf{g}_t(a) = \mathbf{0}$  **then**
- 12:      $\tilde{s}_t^2(a) \leftarrow 0$  {skip CG; no residual ratio}
- 13:    **else**
- 14:      $\tilde{\mathbf{u}} \leftarrow \text{CG}(\mathcal{C}_t, \mathbf{g}_t(a))$ ; stop only once  $\|r_I\|/\|\mathbf{g}_t(a)\| \leq \bar{\varepsilon}_t/\sqrt{\bar{\kappa}_t}$ , or use a certified budget  $I_t$  from (85) {recompute the original residual in floating point; a point check is not an enclosure}
- 15:      $\tilde{s}_t^2(a) \leftarrow \mathbf{g}_t(a)^\top \tilde{\mathbf{u}}$  {approx. predictive variance}
- 16:    **end if**
- 17:   **end for**
- 18:    $a_t \leftarrow \arg \max_{a \in \mathcal{A}_t} \left[ \hat{\mu}_t(a) + \omega_t \sqrt{\frac{u_t(a)}{1 - \bar{\varepsilon}_t}} \sqrt{\tilde{s}_t^2(a)} \right]$  {bonus (68); ties broken by a fixed rule (e.g. lowest index) or by  $\Omega_t$ }
- 19:   Play  $a_t$ , observe  $r_t$
- 20:    $\mathcal{D}_t \leftarrow \mathcal{D}_{t-1} \cup \{(\mathbf{x}_t, a_t, r_t)\}$
- 21:    $c_t \leftarrow \mu_{\theta_t}(\mathbf{x}_t, a_t) - r_t$ ;  $R^{\text{col}} \leftarrow R^{\text{col}} + c_t^2$ ;  $\bar{F} \leftarrow \bar{F} + (\bar{\epsilon}_t^{\text{lin}})^2$ ;  $\hat{\gamma} \leftarrow \hat{\gamma} + \log[1 + \sigma^{-2}u_t\tilde{s}_t^2(a_t)/(1 - \bar{\varepsilon}_t)]$
- 22:    $n \leftarrow n + 1$ ;  $d \leftarrow \theta_t - \bar{\theta}$ ;  $\bar{\theta} \leftarrow \bar{\theta} + d/n$ ;  $J \leftarrow J + d^\top(\theta_t - \bar{\theta})$  {vector Welford update}
- 23:   Run  $J_{t+1}^{\text{opt}}$  warm-started optimizer steps on the *full* data set  $\mathcal{D}_t$  to obtain  $\theta_{t+1}$ , projecting to the certified region if required
- 24:   Certify  $\zeta_{t+1} \geq \|\nabla \mathcal{L}_{t+1}(\theta_{t+1})\|$ ; unless the terminal optimizer evaluation supplies this quantity, this requires one additional full-data gradient evaluation
- 25: **end for**

---

For the exact-curvature theorem, Algorithm 1 is run with  $\mathcal{S}_t = \{1, \dots, t-1\}$ ,  $w_s = 1$ , full action enumeration,  $\mathcal{C}_t = \mathbf{C}_t$ , and the transfer factor  $u_t = (1 + \bar{\chi}_t)^2$  from Lemma 16. For an approximate operator  $\mathcal{C}_t = \hat{\mathbf{C}}_t$ , use  $u_t = \kappa_{+,t}(1 + \bar{\chi}_t)^2$  from Theorem 29. The selection rule needs only the *one-sided* upper transfer factor  $u_t(a)$  for optimism; no lower spectral factor is required (the cost of the operator is charged through the realized dynamic complexity  $\Lambda_T^{\mathcal{C}}$ ). Algorithm 2 computes  $u_t$ ,  $\bar{\psi}_t$ , and  $\bar{\beta}_t$  from pre-action state. The operator-specific inputs are the proved factor  $\kappa_{+,t}$  and condition upper bound  $\bar{\kappa}_t$ ; exact full and unrescaled-subset operators use  $\kappa_{+,t} = 1$ . Windowed/refresh curvature and action screening are practical heuristics unless the corresponding one-sided transfer factor or optimal-arm-survival condition is separately verified.

**Corrected-center remark.** Replacing the center  $\hat{\mu}_t(a)$  by the corrected center  $\hat{\mu}_t^{\text{corr}}(a)$  of (32) sets  $\bar{\psi}_t = 0$  (so  $\omega_t = \bar{\beta}_t$ ) and removes the centering-stability assumption exactly (Corollary 8), at the cost of frozen-feature access for  $\hat{\theta}_t^{\text{lin}}$ ; this departs from the  $O(d)$  additional curvature-state-memory, no-stored-feature design unless a frozen-feature oracle, replay mechanism, or certified sketch is supplied (see the remark after Corollary 8).

CC-UCB is the analyzed method throughout.

**How to read the theorem.** The theorem is a transfer result for a general predictable SPD curvature sequence. Lemma 11 and Corollary 4 instantiate its nonlinear confidence, centering, and feature-transfer inputs from  $O(d)$  path state, observed collection residuals, analytic smoothness constants, and a certified optimizer residual. These schedules are executable but can be very loose. A sublinear guarantee from this bound still requires control of the realized dynamic width complexity and a restricted parameter path; unrestricted end-to-end fine-tuning is not covered. No horizon-uniform lower transfer is required by Theorem 1; the lower sandwich appears only in the explicit near-linear corollary.

### Discussion of terms.

- *Exploration term.* Write  $S_T := \sum_{t=1}^T \alpha_t^2 u_t(a_t) \omega_t^2$  for the sum of squares of the per-round factors. The exploration term is  $2\sqrt{(\sigma^2 + G^2/\lambda) \Lambda_T^C S_T}$ . A sufficient condition for sublinear regret is

$$\Lambda_T^C S_T = o(T^2), \quad E_T = o(T).$$

Concretely this holds whenever the per-round factors are polylogarithmically bounded ( $\sup_t \alpha_t^2 u_t(a_t) \omega_t^2 = O(\text{polylog } T)$ ), so  $S_T = O(T \text{polylog } T)$  and the realized dynamic complexity is small ( $\Lambda_T^C = O(\text{polylog } T)$ , e.g. via  $\Lambda_T^C \leq \Gamma_T^{\text{dyn}}$  with both the variation charge  $V_T^C$  and the endpoint  $\log \det(\mathcal{C}_{T+1}/\lambda \mathbf{I})$  polylogarithmically bounded). Note  $\omega_t^2$  contains the cumulative squared linearization error through  $\bar{F}_t$  (via  $\beta_t$ ) and the centering certificate  $\bar{\psi}_t$ , so a small  $\Lambda_T^C$  alone is insufficient.  $\alpha_t$  is bounded whenever the CG tolerance  $\bar{\varepsilon}_t$  is bounded away from 1; in the exact-curvature case  $u_t = (1 + \bar{\chi}_t)^2$  is bounded whenever the whitened-drift certificate  $\bar{\chi}_t$  is bounded (Lemma 16); the whitened  $\ell_2$  accumulation  $\chi_t \leq \|\mathbf{D}_t\|_F$  makes this a near-stationarity condition on replayed-Jacobian drift. The unstable-curvature cost is isolated in  $\Lambda_T^C$  (through  $V_T^C$ ), not hidden in a spectral constant. The factor  $\sqrt{\sigma^2 + G^2/\lambda}$  arises from the width-sum bound of the dynamic width-complexity lemma (Lemma 19, Eq. (60)); when  $G^2 \leq \lambda \sigma^2$  it is at most  $\sigma\sqrt{2}$ .

- *Linearization remainder.* The bound involves  $E_T = \sum_t \varepsilon_{\text{lin}}(t)$ , not  $T\varepsilon_{\text{lin}}$ . Near-linear regimes also require the squared-error sum  $F_T = \sum_t \varepsilon_{\text{lin}}(t)^2$  (through  $\bar{F}_t$ , hence  $\omega_t$ ) to be controlled.  $\|\boldsymbol{\theta}^* - \boldsymbol{\theta}_t\| = O(1/t)$  gives  $E_T = O(1)$ ;  $\|\boldsymbol{\theta}^* - \boldsymbol{\theta}_t\| = O(t^{-1/4})$  gives  $E_T = O(\sqrt{T})$ . The additive role of  $E_T$  mirrors misspecified linear bandits, where a uniformly bounded per-round model error commonly contributes an additive regret term that scales linearly in the cumulative approximation error (Ghosh et al., 2017; Lattimore et al., 2020). The displayed bound requires  $E_T/T \rightarrow 0$ ; a nonvanishing uniform per-round error bound alone yields an  $O(T)$  additive term.
- *Centering.* The centering certificate enters *multiplicatively*, through  $\omega_t = \bar{\beta}_t + \bar{\psi}_t$  inside the sum of squares  $S_T$ , rather than as a separate additive term. It vanishes ( $\bar{\psi}_t = 0$ ) when  $\boldsymbol{\theta}_t = \hat{\boldsymbol{\theta}}_t^{\text{lin}}$  (convex last-layer models with exact optimization), and the corrected-center variant (Corollary 8) removes it in general at the cost of a frozen-feature solve.
- *Regret improvement over diagonal/last-layer surrogates.* The bound itself does not prove that full curvature yields strictly lower regret than diagonal or last-layer methods in any specific regime; it provides an upper bound with explicit per-round factors. *Demonstrating regret improvement is an empirical hypothesis* (Section 5). The theorem exposes the transfer factor  $u_t$  and realized width complexity  $\Lambda_T^C$  as operator-dependent terms along the realized trajectory; it does not identify a causal regret cost because the adaptive path and other factors may also change.
- *Scope.* The bound is proved for Gaussian/squared-loss rewards. Extending it to exponential-family likelihoods (Bernoulli, categorical) would require a weighted frozen-feature analysis accounting for the data-dependent Fisher weights  $\nabla^2 A(z)$  and is left as future work (Appendix H).

**Corollary 32** (Sanity-check: frozen features recover linear UCB). *Suppose the backbone is frozen and the reward head is linear in its parameters, the ridge objective is optimized exactly, and take the algorithmic curvature to be the sequential frozen Gram matrix,  $\mathcal{C}_t = \bar{\mathcal{C}}_t$ . Then:  $\varepsilon_{\text{lin}}(t) = 0$  for all  $t$  ( $L = 0$  in Assumption 3), so  $E_T = F_T = 0$ ;  $\bar{\psi}_t = 0$  (under exact optimization the MAP and linearized estimators coincide, so Assumption 4 holds with equality), giving  $\omega_t = \bar{\beta}_t$ ;  $u_t(a) = 1$  for all  $a$  (indeed  $\bar{s}_t^2(a) = s_{\bar{\mathcal{C}}_t}^2(a)$ , so Assumption 5 holds with equality); and, since  $\bar{\mathcal{C}}_{t+1} = \bar{\mathcal{C}}_t + \sigma^{-2} \mathbf{g}_t(\mathbf{x}_t, a_t) \mathbf{g}_t(\mathbf{x}_t, a_t)^\top$  by (7), the update is exactly the rank-one step, i.e.  $\mathcal{C}_{t+1} = \mathcal{C}_t^+$ , so  $\Xi_t = 0$ ,  $V_T^C = 0$ , and  $\Lambda_T^C = \gamma_T$ . Then Theorem 1 gives*

$$R_T \leq 2\sqrt{(\sigma^2 + G^2/\lambda) \gamma_T \sum_{t=1}^T \alpha_t^2 \bar{\beta}_t^2}.$$




Figure 3: Supplementary certificate-tightness and curvature-mechanism overview. Panel (a) compares final-round float64 audit quantities with predictable schedules in the 50-seed tanh execution. Panel (b) gives paired regret differences in the preregistered curvature phase map. These are descriptive point audits.

If in addition  $\bar{\varepsilon}_t \leq \bar{\varepsilon}$  so that  $\alpha_t \leq \alpha_I := \sqrt{(1 + \bar{\varepsilon})/(1 - \bar{\varepsilon})}$ , the looser Corollary 33 recovers the familiar LinUCB form  $R_T \leq 2 \alpha_I \bar{\beta}_{\max, T} \sqrt{(\sigma^2 + G^2/\lambda) T \gamma_T}$ , where  $\bar{\beta}_{\max, T} := \max_{t \leq T} \bar{\beta}_t$ .

*Proof.* Substitute  $\varepsilon_{\text{lin}}(t) = 0$ ,  $\bar{\psi}_t = 0$ ,  $u_t(a) = 1$ , and  $\Lambda_T^C = \gamma_T$ ,  $V_T^C = 0$  into (15); the frozen Gram is monotone, so  $\Xi_t = 0$  and the dynamic potential collapses to the information gain  $\gamma_T$ . The LinUCB form is Corollary 33 with  $\alpha_t \leq \alpha_I$ ,  $u_t = 1$ ,  $\omega_t = \bar{\beta}_t \leq \bar{\beta}_{\max, T}$ .  $\square$

## C Executed-Policy Audit Protocols

**Interpretation and statistics.** Primary regret results are generated by online policies executed from scratch. Synthetic evaluators know the teacher only to compute pseudo-regret, optimism failures, and post-hoc theorem-event checks; the action rule never receives it. Hyperparameter selection, where used, is performed on the declared tuning seeds before evaluation. Common random numbers pair methods within an evaluation seed. We report arithmetic means, marginal 95% Student- $t$  intervals, and paired Student- $t$  intervals for method differences. No multiplicity adjustment is applied, so the intervals are descriptive. Scaling-curve bands instead use pointwise percentile bootstrap intervals. Their slope intervals resample whole seed trajectories, reaverage the same five nested prefixes, and refit the specified unweighted OLS model; they measure seed sensitivity conditional on that finite grid, not whether a power law is appropriate. An offline common-trajectory diagnostic uses one logged action and checkpoint sequence for all operators and is never mixed with online regret.

**Predictable tanh execution.** The analytically bounded tanh instance is  $\mu_\theta(x, a) = \tanh(\phi(x, a)^\top \theta)$  with  $d = 6$ , four actions,  $\|\phi\| \leq 0.75$ , Gaussian noise  $\sigma = 0.20$ , teacher radius  $S = 0.35$ , and an optimizer projected to radius 0.50. We use the global constants  $G = 0.75$  and  $L_g = L_\mu = 4(0.75)^2/(3\sqrt{3})$ ,  $\lambda = 1$ ,  $\delta = 0.05$ , eight clipped gradient steps per round, and a CG energy target 0.05 with the analytic condition bound  $1 + (t - 1)G^2/(\lambda\sigma^2)$ . At every round the policy records  $Q_t$ ,  $R_t^{\text{col}}$ ,  $\bar{F}_t$ ,  $\hat{\gamma}_t$ ,  $\bar{\chi}_t$ ,  $\bar{M}_t$ ,  $\bar{\psi}_t$ ,  $\bar{\beta}_t$ , the optimizer gradient norm, all candidate residuals, and whether any post-hoc exact check fails. No post-hoc field is read by action selection. Seeds 60–69 run a preregistered eight-cell one-factor sensitivity grid; the fixed reference configuration is rerun from scratch on evaluation seeds 160–209 for both centers. The 50-seed means and intervals in Section 5.1 come only from those evaluation runs.

| $r$ | Method           | Regret slope [95% interval] | $R^2$  | Max. $ \log R - \widehat{\log R} $ | $\Lambda$ slope [95% interval] |
|-----|------------------|-----------------------------|--------|------------------------------------|--------------------------------|
| 4   | Exact current    | -0.000 [-0.000, 0.000]      | 1.0000 | 2.220e-16                          | 0.106 [0.105, 0.107]           |
| 4   | Full CG          | -0.000 [-0.000, 0.000]      | 1.0000 | 2.220e-16                          | 0.105 [0.104, 0.106]           |
| 4   | Window $q = 1/2$ | 0.851 [0.843, 0.859]        | 0.9937 | 8.445e-02                          | 0.465 [0.464, 0.465]           |
| 4   | Window $q = 2/3$ | 0.309 [0.292, 0.326]        | 0.8529 | 1.550e-01                          | 0.334 [0.333, 0.334]           |
| 4   | Window $q = 1$   | -0.000 [-0.000, 0.000]      | 1.0000 | 2.220e-16                          | 0.106 [0.105, 0.107]           |
| 4   | Frozen           | 1.000 [1.000, 1.000]        | 1.0000 | 4.441e-15                          | 0.106 [0.106, 0.106]           |
| 4   | Diagonal         | 1.000 [1.000, 1.000]        | 1.0000 | 4.441e-15                          | 0.673 [0.673, 0.673]           |
| 4   | Greedy           | 0.000 [-0.000, 0.000]       | 1.0000 | 5.551e-17                          | 0.115 [0.114, 0.115]           |
| 8   | Exact current    | 0.134 [0.115, 0.153]        | 0.6887 | 1.217e-01                          | 0.099 [0.098, 0.100]           |
| 8   | Full CG          | 0.141 [0.122, 0.161]        | 0.6969 | 1.247e-01                          | 0.099 [0.098, 0.099]           |
| 8   | Window $q = 1/2$ | 1.000 [1.000, 1.000]        | 1.0000 | 1.066e-14                          | 0.383 [0.383, 0.383]           |
| 8   | Window $q = 2/3$ | 0.994 [0.993, 0.995]        | 1.0000 | 7.659e-03                          | 0.304 [0.304, 0.304]           |
| 8   | Window $q = 1$   | 0.134 [0.116, 0.153]        | 0.6887 | 1.217e-01                          | 0.099 [0.098, 0.100]           |
| 8   | Frozen           | 1.000 [1.000, 1.000]        | 1.0000 | 1.066e-14                          | 0.115 [0.115, 0.115]           |
| 8   | Diagonal         | 1.000 [1.000, 1.000]        | 1.0000 | 1.066e-14                          | 0.671 [0.671, 0.671]           |
| 8   | Greedy           | 0.037 [0.018, 0.057]        | 0.5542 | 4.606e-02                          | 0.125 [0.124, 0.126]           |
| 16  | Exact current    | 0.818 [0.804, 0.831]        | 0.9827 | 1.452e-01                          | 0.117 [0.117, 0.118]           |
| 16  | Full CG          | 0.829 [0.815, 0.842]        | 0.9841 | 1.418e-01                          | 0.118 [0.117, 0.118]           |
| 16  | Window $q = 1/2$ | 1.000 [1.000, 1.000]        | 1.0000 | 2.665e-15                          | 0.406 [0.406, 0.406]           |
| 16  | Window $q = 2/3$ | 1.000 [1.000, 1.000]        | 1.0000 | 2.665e-15                          | 0.278 [0.278, 0.278]           |
| 16  | Window $q = 1$   | 0.818 [0.804, 0.831]        | 0.9827 | 1.452e-01                          | 0.117 [0.117, 0.118]           |
| 16  | Frozen           | 1.000 [1.000, 1.000]        | 1.0000 | 2.665e-15                          | 0.125 [0.125, 0.125]           |
| 16  | Diagonal         | 1.000 [1.000, 1.000]        | 1.0000 | 2.665e-15                          | 0.669 [0.669, 0.669]           |
| 16  | Greedy           | 0.273 [0.229, 0.315]        | 0.9024 | 1.086e-01                          | 0.140 [0.138, 0.141]           |

Table 4: Five-point OLS diagnostics for every method at  $d = 2048$ . Slope brackets are obtained from 2,000 paired resamples of whole seed trajectories.  $R^2$  and maximum absolute log residual describe the fit to the five selected prefix means; they are not tests of a power-law model or asymptotic exponents.

**Linear theorem-event audit.** The context is  $x \in \{-1, +1\}^8/\sqrt{8}$ , there are five actions, and  $\phi(x, a) = [x, e_a, x \otimes e_a]$  has dimension 53 and norm  $\sqrt{3}$ . The fixed teacher contains shared, arm-intercept, and context-arm interaction blocks; exact enumeration of all  $2^8$  contexts confirms that multiple arms are optimal. Noise is  $\mathcal{N}(0, 0.25^2)$ . Evaluation uses seeds 100–119 after tuning on seeds 0–9. The tuning grid is  $\lambda \in \{0.5, 1, 2\}$  and  $c_{\text{bonus}} \in \{1, 1.25\}$ ; a candidate is eligible only if every tuning run passes the recorded confidence and small-scale numerical checks. Fresh evaluation policies use dense full Gram, CG full Gram, diagonal, window 64, rescaled subsample 64, rank-16 Lanczos–Ritz, and refresh period 20. The CG target is 0.05 with at most 106 iterations. Fixed-reference and independently tuned comparisons are both saved.

**Linear policy checks and bound scale.** The implemented base radius is

$$\beta_t^{\text{base}} = \sqrt{d \log \left( 1 + \frac{(t-1)G^2}{d\lambda\sigma^2} \right) + 2 \log(1/\delta) + \sqrt{\lambda} S}, \quad \bar{\beta}_t = c_{\text{bonus}} \beta_t^{\text{base}}, \quad c_{\text{bonus}} \geq 1,$$

with  $d = 53$ ,  $G = \sqrt{3}$ ,  $S = \|\theta^*\|_2$ , and  $\delta = 0.05$ ; hence tuning  $c_{\text{bonus}} \in \{1, 1.25\}$  can only inflate the valid radius. Because  $\bar{\psi}_t = 0$ , the policy uses  $\omega_t = \bar{\beta}_t$ . Dense, CG, unrescaled window, and stale-refresh policies use the analytic factor  $u_t = 1$ . Diagonal, rescaled-subsample, and Lanczos policies materialize dense float64 small matrices before action selection and use  $u_t = \max\{1, (1 + 32\epsilon_{\text{mach}})\lambda_{\max}(\bar{C}_t^{-1/2}\mathcal{C}_t\bar{C}_t^{-1/2})\}$ . The multiplicative  $1 + 32\epsilon_{\text{mach}}$  is a heuristic inflation, not a verified rounding-error enclosure. The CG policy likewise sets  $\bar{\kappa}_t$  to a float64 point estimate of  $\text{cond}(\mathcal{C}_t)$  before action selection and enforces  $\bar{\epsilon}_t = 0.05$  using relative residual target  $\bar{\epsilon}_t/\sqrt{\bar{\kappa}_t}$ ; the other policies have  $\bar{\epsilon}_t = 0$ . These pre-action computations are deterministic conditional on each seeded operator draw, but predictability does not turn an unenclosed point estimate into an upper certificate. The analytic  $u_t = 1$  transfer statements for dense full Gram, unrescaled window, and stale-refresh are unaffected; diagonal, rescaled-subsample, Lanczos, and full-Gram CG remain numerical audits of theorem semantics rather than certified executions. Tuning-run event checks are a selection gate, not the mathematical basis of the evaluation checks. The machine-readable per-policy trace is `results/derived/certification_audit.json`.

**Nonlinear drift stress test.** The teacher and student are one-hidden-layer width-four tanh networks with four inputs, five outputs, and 45 displacement parameters. They share the initialized backbone, making the exact frozen-head ridge solve a 25-parameter linear regime. Full-network regimes use respectively (learning rate, updates, step cap) = (0.002, 2, 0.005), (0.004, 3, 0.020), (0.008, 4, 0.080). Each original- or corrected-center policy executes for 100 rounds on seeds 110–129 with a fixed, time-only schedule; no evaluation-seed



| Study          | Policy                | Regret [95% CI]      | Measured run time (s) | Event fail. | Status            |
|----------------|-----------------------|----------------------|-----------------------|-------------|-------------------|
| Tanh fixed     | Original center       | 7.86 [7.66, 8.07]    | 0.094                 | 0/50        | Post-hoc verified |
| Tanh fixed     | Corrected center      | 5.69 [5.46, 5.92]    | 0.095                 | 0/50        | Post-hoc verified |
| Balanced tuned | CC-UCB (full GGN-CG)  | 21.94 [20.64, 23.23] | 0.966                 | –           | Uncertified       |
| Balanced tuned | Diagonal full network | 16.11 [14.62, 17.61] | 0.024                 | –           | Uncertified       |
| Balanced tuned | LinUCB                | 10.58 [9.82, 11.33]  | 0.015                 | –           | Uncertified       |
| Balanced tuned | LinearTS              | 12.67 [12.02, 13.31] | 0.017                 | –           | Uncertified       |
| Balanced tuned | NeuralLinear          | 10.01 [9.33, 10.68]  | 0.014                 | –           | Uncertified       |
| Balanced tuned | NeuralUCB             | 23.04 [21.81, 24.27] | 0.027                 | –           | Uncertified       |
| Balanced tuned | NeuralTS              | 25.07 [23.40, 26.75] | 0.032                 | –           | Uncertified       |
| Balanced tuned | Frozen last layer     | 7.44 [6.90, 7.99]    | 0.012                 | –           | Uncertified       |
| Balanced tuned | Greedy                | 16.26 [14.44, 18.09] | 0.017                 | –           | Uncertified       |
| Balanced tuned | UCB1                  | 90.35 [87.95, 92.75] | 0.003                 | –           | Uncertified       |
| Balanced tuned | Context-free TS       | 89.40 [86.53, 92.26] | 0.004                 | –           | Uncertified       |

Table 5: Supplementary executed-policy summary. Tanh rows use one fixed theoretical configuration; balanced rows are selected only on tuning seeds and rerun on disjoint evaluation seeds. Float64 event checks are not verified enclosures, and balanced rows are uncertified performance comparisons. Measured run times are resource accounting under each executed implementation, not a controlled speed or matched-wall-clock comparison.




Figure 4: Off-diagonal linear-Gram witness. The analytic cell is constructive; the noisy cells are uncertified extensions. Paired differences are descriptive and establish no uniform ordering of curvature surrogates.

tuning occurs. Dense float64 replay computes  $E_T, F_T$ , the dense-reference centering ratio and primitive  $\psi_t, \chi_t$ , generalized transfer factors, condition numbers, dense-reference and residual-based CG errors,  $\Lambda_T^C$ , endpoint log determinant,  $V_T^C, \Gamma_T^{\text{dyn}}$ , and all-action optimism failures. Policy-available and post-hoc fields are separated in every JSONL record. The corrected-center identity is checked directly.

A secondary post-hoc analysis across the eight dependent regime-center cell means gives descriptive Spearman correlation 0.83 between the post-hoc decomposition score and regret, versus 0.57 for relative curvature change. With  $n = 8$ , reused seeds, and no seed-level sensitivity or inferential analysis, this is neither a primary result nor a causal or inferential claim. This stress test is secondary to the predictable tanh execution because its schedules fail.

**Preregistered curvature-mechanism map.** The fixed-gap bounded-linear family has  $d = 10$ , four actions,  $T = 30$ , and Gaussian noise 0.2. Its eight-cell fractional design crosses active-spectrum condition number  $\{3, 50\}$ , rotation  $\{0, 60^\circ\}$ , action gap  $\{0.15, 0.50\}$ , nuisance strength  $\{0.4, 1.2\}$ , effective rank  $\{3, 8\}$ , and damping  $\{0.5, 2\}$  according to the generators stored in the configuration; all cells are fixed before execution and no tuning is run. Exact full, residual-checked full CG, diagonal, block-diagonal, rank-three Lanczos, window-eight, and refresh-period-five policies execute independently on seeds 200–229. The exact-full trajectory is additionally

| Policy                          | $R_{10}$ | $R_{100}$ | $R_{1000}$ | $R_{10000}$ | Log-log slope |
|---------------------------------|----------|-----------|------------|-------------|---------------|
| Exact full                      | 0.90     | 0.90      | 0.90       | 0.90        | 0.000         |
| Residual-checked full CG        | 0.90     | 0.90      | 0.90       | 0.90        | 0.000         |
| Diagonal (raw)                  | 9.00     | 90.00     | 900.00     | 9000.00     | 1.000         |
| Diagonal (uniform transfer)     | 9.00     | 90.00     | 900.00     | 9000.00     | 1.000         |
| Diagonal (actionwise reference) | 0.90     | 0.90      | 0.90       | 0.90        | 0.000         |
| Greedy                          | 9.00     | 90.00     | 900.00     | 9000.00     | 1.000         |

Table 6: Executed analytic off-diagonal linear-Gram witness by horizon. Every value is generated from the provenance-validated derived artifact; fitted slopes are finite-horizon diagnostics.

| Method       | $T = 250$ | $T = 500$ | $T = 1000$ |
|--------------|-----------|-----------|------------|
| full CG      | 375.7     | 438.1     | 479.3      |
| full dense   | 361.5     | 422.4     | 461.7      |
| diagonal     | 4120.8    | 7710.5    | 13456.0    |
| Lanczos–Ritz | 495.7     | 496.4     | 509.3      |
| subsample 64 | 339.7     | 373.1     | 396.6      |
| refresh 20   | 370.3     | 465.1     | 526.7      |
| window 64    | 252.0     | 255.3     | 260.2      |

Table 7: Fixed-reference linear RHS-to-regret ratio  $\text{RHS}/R_T$ . The JSON artifact also records  $R_T$ , RHS,  $R_T/T$ , and  $\text{RHS}/T$  at every horizon for both fixed and tuned configurations. Every displayed RHS exceeds  $T$  times the exact maximum one-round pseudo-regret 0.8118 on the finite context support, so the reported RHS values are numerically vacuous. For the four affected policies they are not rigorously certified bounds because their spectral point estimates lack verified upper enclosures.

replayed through every operator. These common-trajectory records contain per-action width ratios, cross-action coefficient of variation, Spearman and Kendall ranks, raw and scalar-rescaled top-score disagreement, normalized decision-margin distortion, candidate-gradient alignment with retained and discarded eigenspaces, global and action-set transfer, effective rank, and condition number. Paired intervals are method-minus-exact-full; they are descriptive and unadjusted across cells.

**Coverage-matched extension.** The extension retains the same eight cells and operator settings, but adds tuning seeds 2000–2019 and evaluation seeds 2100–2149. Coefficients are pooled over every tuning cell, never selected per evaluation cell. The three protocols use an identical theoretical coefficient, the empirical 95th percentile of all one-step prediction-error/width ratios, or the ratio needed to match the full operator’s mean bonus on common tuning trajectories. The latter two coefficients are frozen before any evaluation run. Holm adjustment is over all  $3 \times 8 \times 6 = 144$  prespecified surrogate comparisons. Because this is a fixed-feature linear model, current and historical full-Gram rows are algebraically identical; LO-FI is not included because the repository has no compatible recursive precision implementation.

**Balanced contextual benchmark.** Normalized Rademacher contexts have dimension four, five available actions, and Gaussian noise 0.10; exact enumeration of the 16 contexts verifies that every arm is optimal somewhere. Each method selects  $(\lambda, c_{\text{bonus}})$  from  $\{0.5, 1\} \times \{0.5, 1, 2\}$  by mean pseudo-regret over tuning seeds 70–79 at  $T = 80$ , then runs from scratch for  $T = 200$  on seeds 260–289 with shared contexts and reward noise. Full-network methods use the same width-four tanh backbone, one clipped float64 update per round, learning rate 0.01, and model ridge 0.01. The CC-UCB row relinearizes all selected historical samples at the current parameter and solves one fixed full-GGN operator per action with explicit residual checks. The local NeuralUCB row is the historical-time full-network gradient-Gram baseline: it stores each selected gradient at its collection parameter rather than relinearizing it. NeuralLinear is a Bayesian Gaussian head on the frozen initialized representation; the frozen-last-layer row uses the same posterior with a UCB rule. NeuralUCB and NeuralTS are identified as local matched implementations rather than exact reproductions of every published training detail. The one-sample clipped updates do not optimize (3) or certify its full-gradient residual, so these benchmark rows are practical comparisons rather than theorem instantiations. LinUCB, linear Thompson sampling, diagonal full-network UCB, greedy, Gaussian UCB1, and Gaussian context-free Thompson sampling complete the comparison. Arithmetic means, marginal intervals, paired differences, selected hyperparameters, run times, and all seed-level records are retained.

| Operator               | Identical | 95% coverage | Mean bonus |
|------------------------|-----------|--------------|------------|
| block_kfac             | 1.000     | 0.222        | 1.137      |
| current_full_ggn       | 1.000     | 0.208        | 1.000      |
| diagonal               | 1.000     | 0.248        | 1.337      |
| historical_frozen_gram | 1.000     | 0.208        | 1.000      |
| rank_r                 | 1.000     | 0.154        | 0.421      |
| stale_refresh          | 1.000     | 0.198        | 0.685      |
| window                 | 1.000     | 0.163        | 0.665      |

Table 8: Tuning-selected scalar multipliers for the coverage-matched operator study. Evaluation outcomes are not used to select these values.




Figure 5: Full-grid evaluation coverage and top-action disagreement under all three frozen calibration protocols. These are independently executed policy diagnostics, not causal operator effects.

**Secondary operator and CG audits.** The operator factorial varies windows 16/64, uniform resamples 16/64, Lanczos ranks 4/16, and refresh periods 5/20 in addition to full, frozen, and diagonal operators. On both the bounded linear task and the medium-drift, corrected-center tanh task, twenty online evaluation seeds (120–139) and matching offline common trajectories record exact global and action-set transfer, widths, disagreement, determinant terms, regret, runtime, and memory. The nonlinear policies use the fixed schedule and remain post-hoc audits rather than certified executions. The CG audit uses  $d = 64$ ,  $n = 256$ ,  $K = 5$ , condition numbers 10/100/1000, five target errors, two initializations, and two preconditioners on seeds 130–139. Its solver audit contains 600,000 dense-reference action solves; its separate policy audit executes 200 full-enumeration linear-bandit policies for 200 rounds with the same five tolerances and start/preconditioner factorial. Unlike the one-step scaling construction, the rotated-SPD right-hand sides exercise nontrivial Krylov convergence. At target energy error 0.1 with zero start and no preconditioner, mean iteration counts are 6.026, 22.643, and 55.043 for condition numbers 10,  $10^2$ , and  $10^3$ ; mean measured energy errors are  $2.23 \times 10^{-2}$ ,  $6.87 \times 10^{-3}$ , and  $2.45 \times 10^{-3}$ , and mean relative width errors are  $2.54 \times 10^{-4}$ ,  $2.43 \times 10^{-5}$ , and  $3.28 \times 10^{-6}$ . The residual certificate is at most the prescribed 0.1 in every retained solve. These solver audits do not establish a method-speed advantage or a policy benefit.

| Policy              | $T = 200$     | $T = 500$     | $T = 1000$    | $T = 1500$     |
|---------------------|---------------|---------------|---------------|----------------|
| full GGN-CG         | 163.3 / 0.183 | 399.0 / 0.202 | 794.3 / 0.206 | 1189.7 / 0.207 |
| frozen Gram         | 164.0 / 0.180 | 401.1 / 0.198 | 797.5 / 0.202 | 1193.0 / 0.205 |
| diagonal full       | 166.4 / 0.168 | 407.7 / 0.185 | 807.2 / 0.193 | 1204.4 / 0.197 |
| last-layer full     | 170.7 / 0.146 | 420.9 / 0.158 | 849.4 / 0.151 | 1273.1 / 0.151 |
| last-layer diagonal | 171.1 / 0.144 | 426.0 / 0.148 | 851.6 / 0.148 | 1274.9 / 0.150 |
| greedy              | 162.3 / 0.189 | 393.9 / 0.212 | 781.9 / 0.218 | 1166.2 / 0.223 |
| UCB1                | 145.4 / 0.273 | 332.5 / 0.335 | 616.2 / 0.384 | 895.6 / 0.403  |
| Beta–Bernoulli TS   | 124.7 / 0.376 | 284.9 / 0.430 | 544.0 / 0.456 | 802.6 / 0.465  |

Table 9: Covertypes evaluation-seed means, shown as cumulative pseudo-regret / accuracy. Contextual hyperparameters are selected only on validation seeds; UCB1 and Thompson sampling are parameter-free. Per-horizon marginal and paired 95% Student- $t$  intervals, runtimes, policy types, and certification status are in `results/derived/covertypes_horizon_results.json`.

**Systems audit.** Five evaluation seeds (140–144) cross  $d \in \{32, 64, 128\}$ ,  $n \in \{32, 128, 512\}$ ,  $K \in \{5, 10\}$ , and  $I \in \{5, 15, 30\}$ . We measure dense exact, separate full CG, diagonal, final-coordinate block, Lanczos–Ritz, row-batched independent CG, and row-batched Jacobi-PCG. Every action retains its own original-system residual; batched methods share only the operator evaluation. An additional matrix-free CPU grid uses  $d \in \{512, 2048, 8192\}$ ,  $n = 32$ ,  $K = 4$ , and  $I = 16$  for the two batched solvers, with no dense allocation. Records include wall time, host memory, operator calls, equivalent sample CVPs, per-action iterations and residuals, and width error. These are synthetic float64 operator timings on one CPU host and use different computational primitives. The timed operator stores an  $n \times d$  random-feature matrix and applies an averaged Gram; it does not time model training or on-the-fly network JVP/VJP. Timings are descriptive resource accounting under the executed implementations, not matched-wall-clock comparisons; no trained-network or accelerator extrapolation is made. A Buck-managed Conda toolchain now runs the separate `real-torch.func` pipeline. Smoke checks and one preregistered development-pilot cell for each of the 131,841- and 9,972,737-parameter MLPs executed on an accelerator; the evaluation-seed grid, intervals, and main-paper systems artifacts remain unrun.

### C.1 Covertypes real-data baseline audit

We use UCI Covertypes DOI 10.24432/C50K5N via scikit-learn 1.9.0; the cached dataset checksum is `eeab5db05bb3a9c00c62ced4c7d13112a00cd3487f59ac1943aaf5983d7c3ddf`. A PCG64 permutation with seed 20260717 creates non-stratified train/validation/test counts 348,607/116,202/116,203. Training means and scales standardize columns 0–9; binary columns 10–53 are unchanged. A width-four tanh network is initialized from seeded fan-in normals and optimized by one online SGD step (learning rate 0.01, ridge 1) using only selected-arm feedback. Full-GGN, frozen-Gram, diagonal, and greedy policies train all 255 parameters; last-layer policies train 35 head parameters. This one-sample update is not the full-history objective (3) and supplies no optimizer-residual certificate. Damping  $\{1, 10\}$  and bonus  $\{1, 2, 4\}$  are selected on validation seeds 50–59; greedy fixes bonus zero. Evaluation seeds 150–159 use test contexts only after selection. CG uses zero initialization, relative residual  $10^{-3}$ , at most 80 iterations, and one fixed operator per solve. The selection artifact and its validation evidence are SHA-256 committed inside every evaluation manifest.

The exact test-split label counts are (42267, 56737, 7221, 536, 1902, 3444, 4096), giving a label-2 majority rate 0.488258. Uniform random has exact expected accuracy  $1/7 = 0.142857$ . The fixed majority arm is a whole-test-label oracle diagnostic, not a deployable policy. We add deployable context-free UCB1 (one forced pull per arm, then  $\hat{\mu}_a + \sqrt{2 \log t / N_a}$ ) and independent Beta(1, 1)–Bernoulli Thompson sampling. All eight executed methods share evaluation seeds and context order.

Both non-contextual policies beat every contextual method at every horizon. At  $T = 1500$ , UCB1 and Thompson sampling have regret 895.6 and 802.6 (accuracy 0.403 and 0.465), while full GGN-CG has 1189.7 (0.207). Their paired method-minus-full differences are  $-294.1$   $[-388.1, -200.1]$  and  $-387.1$   $[-479.3, -294.9]$ . Full and diagonal last-layer policies have regret 1273.1 and 1274.9 (accuracy 0.151 and 0.150). Mean complete-run times at  $T = 1500$  are 5.665 seconds for full GGN-CG, 0.602 for diagonal full-network, 0.007 for UCB1, and 0.017 for Thompson sampling. No method is theorem-certified: binary rewards are routed through Gaussian squared-loss curvature. The full-GGN CG calls reach their relative-residual rule, but no energy-error certificate is supplied, so they are not called solver-certified. This study is a failed baseline check, not external validation of curvature-aware exploration.

**Environment and artifacts.** All reported new runs used CPython 3.14.0 on one arm64 host with 16 logical CPUs, 64 GiB physical memory, Darwin 25.5.0, and no accelerator. Core versions are NumPy 2.5.1, SciPy 1.18.0, matplotlib 3.11.0, psutil 7.2.2, and scikit-learn 1.9.0. Each run directory contains the resolved configuration, seed, UTC timestamp, Git revision and dirty flag, package and hardware record, raw per-round JSONL, and a summary. Derived JSON, CSV, table, and figure artifacts carry SHA-256 sidecars linking them to the aggregate and raw inputs.

**Corollary 33** (Looser worst-case reduction). *If  $\alpha_t \leq \alpha_{\max}$ ,  $u_t(a_t) \leq u_{\max}$ , and  $\omega_t \leq \omega_{\max}$  for all  $t \leq T$ , then  $S_T \leq \alpha_{\max}^2 u_{\max} \omega_{\max}^2 T$  and*

$$R_T \leq 2 \alpha_{\max} \omega_{\max} \sqrt{u_{\max}} \sqrt{(\sigma^2 + G^2/\lambda) T \Lambda_T^{\mathcal{C}}} + 2E_T.$$

*This  $\sqrt{T \Lambda_T^{\mathcal{C}}}$  form recovers the familiar worst-case expression; it is weakly looser than (15) (no tighter, with equality when the per-round factors are constant) and is not the main result.*

## D Proof of Theorem 1

We recall the key objects:  $\mathbf{g}_s(\mathbf{x}, a) = \nabla_{\boldsymbol{\theta}} \mu_{\boldsymbol{\theta}_s}(\mathbf{x}, a)$  (frozen features),  $\bar{\mathbf{C}}_t = \lambda \mathbf{I} + \sigma^{-2} \sum_{s < t} \mathbf{g}_s(\mathbf{x}_s, a_s) \mathbf{g}_s(\mathbf{x}_s, a_s)^\top$  (frozen-feature Gram matrix),  $\bar{s}_t(a) = \sqrt{\mathbf{g}_t(a)^\top \bar{\mathbf{C}}_t^{-1} \mathbf{g}_t(a)}$ , the algorithmic curvature  $\mathcal{C}_t$  with width  $s_{\mathcal{C},t}(a) = \sqrt{\mathbf{g}_t(a)^\top \mathcal{C}_t^{-1} \mathbf{g}_t(a)}$ , the linearized ridge estimator  $\hat{\boldsymbol{\theta}}_t^{\text{lin}}$  (Eq. 80), the combined width factor  $\omega_t = \bar{\beta}_t + \bar{\psi}_t$ , the per-round CG inflation  $\alpha_t = \sqrt{(1 + \bar{\varepsilon}_t)/(1 - \bar{\varepsilon}_t)}$ , and the realized dynamic complexity  $\Lambda_T^{\mathcal{C}}$  (Eq. 56). Write  $\hat{\mu}_t(a) = \mu_{\boldsymbol{\theta}_t}(\mathbf{x}_t, a)$ .

**Step 1: Confidence, centering, and prediction-error bound.** By Lemma 24, with probability  $\geq 1 - \delta$ ,

$$|\mathbf{g}_t(a)^\top (\hat{\boldsymbol{\theta}}_t^{\text{lin}} - \boldsymbol{\theta}^*)| \leq \bar{s}_t(a) \beta_t, \quad \forall t, a, \quad (81)$$

where  $\beta_t$  is as in (72). The self-normalized martingale inequality (Abbasi-Yadkori et al., 2011) is applied to the 1-sub-Gaussian sequence  $\xi_s = \eta_s/\sigma$  with the predictable design  $\mathbf{g}_s/\sigma$  and Gram matrix  $\bar{\mathbf{C}}_t$ . Each realized design vector  $\mathbf{g}_s(\mathbf{x}_s, a_s)$  is  $\mathcal{G}_s$ -measurable (where  $\mathcal{G}_s = \mathcal{H}_s \vee \sigma(a_s)$  is the pre-reward sigma-algebra of (78)), and the sequential update (7) ensures the hypotheses of the self-normalized bound are satisfied. By Lemma 27 (Assumption 4),  $|\mathbf{g}_t(a)^\top (\boldsymbol{\theta}_t - \hat{\boldsymbol{\theta}}_t^{\text{lin}})| \leq \bar{\psi}_t \bar{s}_t(a)$ . Combined with  $|\rho_t| \leq \varepsilon_{\text{lin}}(t)$  and  $\bar{\beta}_t \geq \beta_t$ , the triangle inequality gives the prediction-error bound (73)

$$|\mu_{\boldsymbol{\theta}^*}(\mathbf{x}_t, a) - \hat{\mu}_t(a)| \leq (\beta_t + \bar{\psi}_t) \bar{s}_t(a) + \varepsilon_{\text{lin}}(t) \leq \omega_t \bar{s}_t(a) + \varepsilon_{\text{lin}}(t), \quad \forall a. \quad (82)$$

(The corrected-center variant of Corollary 8 reaches (82) with  $\omega_t = \bar{\beta}_t$  and without Assumption 4.)

**Step 2: Bonus sandwiching.** By Lemma 23 (which uses Lemma 22 with  $\max_{a \in \mathcal{A}_t} \varepsilon_{t,a} \leq \bar{\varepsilon}_t$  and the one-sided transfer (47)), for all  $a$ ,

$$\omega_t \bar{s}_t(a) \leq w_t(a) \leq \alpha_t \omega_t \sqrt{u_t(a)} s_{\mathcal{C},t}(a). \quad (83)$$

The lower bound and (82) give  $\mu^*(\mathbf{x}_t, a) \leq \hat{\mu}_t(a) + w_t(a) + \varepsilon_{\text{lin}}(t)$  for all  $a$ : the UCB covers  $\mu^*$  up to the additive  $\varepsilon_{\text{lin}}(t)$  model-mismatch term.

**Step 3: Instantaneous regret.** From the UCB rule  $\hat{\mu}_t(a_t) + w_t(a_t) \geq \hat{\mu}_t(a_t^*) + w_t(a_t^*)$ , and using  $\mu^*(a_t^*) \leq \hat{\mu}_t(a_t^*) + w_t(a_t^*) + \varepsilon_{\text{lin}}(t)$  (Step 2 at  $a_t^*$ , valid since  $a_t^* \in \mathcal{A}_t$ ),

$$\begin{aligned} \mu^*(a_t^*) - \mu^*(a_t) &\leq [\hat{\mu}_t(a_t^*) + w_t(a_t^*) + \varepsilon_{\text{lin}}(t)] - \mu^*(a_t) \\ &\leq [\hat{\mu}_t(a_t) + w_t(a_t) + \varepsilon_{\text{lin}}(t)] - \mu^*(a_t) \quad (\text{UCB selection}) \\ &= w_t(a_t) + [\hat{\mu}_t(a_t) - \mu^*(a_t)] + \varepsilon_{\text{lin}}(t) \\ &\leq w_t(a_t) + \omega_t \bar{s}_t(a_t) + 2\varepsilon_{\text{lin}}(t) \quad ((82) \text{ at } a_t) \\ &\leq 2w_t(a_t) + 2\varepsilon_{\text{lin}}(t) \quad (\text{lower bound in (83)}). \end{aligned}$$

---

**Step 4: Upper bonus bound.** The upper bound in (83) gives

$$\mu^*(a_t^*) - \mu^*(a_t) \leq 2\alpha_t \omega_t \sqrt{u_t(a_t)} s_{\mathcal{C},t}(a_t) + 2\varepsilon_{\text{lin}}(t). \quad (84)$$

**Step 5: Cauchy–Schwarz after summing.** Summing (84) over  $t = 1, \dots, T$  and applying Cauchy–Schwarz to the sum (*not* per round),

$$R_T \leq 2 \sum_{t=1}^T \alpha_t \omega_t \sqrt{u_t(a_t)} s_{\mathcal{C},t}(a_t) + 2E_T \leq 2 \sqrt{\sum_{t=1}^T \alpha_t^2 \omega_t^2 u_t(a_t)} \sqrt{\sum_{t=1}^T s_{\mathcal{C},t}^2(a_t)} + 2E_T.$$

**Step 6: Dynamic potential and collection.** Lemma 19 gives  $\sum_{t=1}^T s_{\mathcal{C},t}^2(a_t) \leq (\sigma^2 + G^2/\lambda) \Lambda_T^{\mathcal{C}}$ . Substituting,

$$R_T \leq 2 \sqrt{(\sigma^2 + G^2/\lambda) \Lambda_T^{\mathcal{C}} \sum_{t=1}^T \alpha_t^2 u_t(a_t) \omega_t^2} + 2E_T,$$

which is (15). The interpretable and observable versions (31) follow from  $\Lambda_T^{\mathcal{C}} \leq \Gamma_T^{\text{dyn}}$  (Lemma 19) and  $\Lambda_T^{\mathcal{C}} \leq \widehat{\Lambda}_T^{\mathcal{C}}$  (Corollary 21). The Cauchy–Schwarz step is applied only after summing, so the time-varying factors  $\alpha_t, u_t(a_t), \omega_t$  enter through a sum of squares; no horizon maximum is taken before Cauchy–Schwarz, and the usual  $\sqrt{T}$  scaling is recovered when the per-round factors are uniformly bounded (Corollary 33). The stated probability follows from a union bound over the self-normalized confidence event  $\mathcal{E}_{\text{conf}}$  ( $\geq 1 - \delta$ ) and the certificate event  $\mathcal{E}_{\text{cert}}$  ( $\geq 1 - \delta_{\text{cert}}$ ), giving  $\geq 1 - \delta - \delta_{\text{cert}}$ .  $\square$

## E CG Convergence Analysis

Fix a round  $t$  and an action  $a$  with  $\mathbf{g} := \mathbf{g}_t(a) \neq \mathbf{0}$ . After  $I_{t,a}$  iterations of CG on the SPD system  $\mathcal{C}_t \mathbf{u} = \mathbf{g}$  (starting from  $\tilde{\mathbf{u}}_0 = \mathbf{0}$ ), the standard energy-norm bound gives

$$\varepsilon_{t,a} = \frac{\|\tilde{\mathbf{u}}_{I_{t,a}} - \mathcal{C}_t^{-1} \mathbf{g}\|_{\mathcal{C}_t}}{\|\mathcal{C}_t^{-1} \mathbf{g}\|_{\mathcal{C}_t}} \leq 2 \left( \frac{\sqrt{\kappa_t} - 1}{\sqrt{\kappa_t} + 1} \right)^{I_{t,a}},$$

where  $\kappa_t = \lambda_{\max}(\mathcal{C}_t)/\lambda_{\min}(\mathcal{C}_t)$  is the true condition number. The bound is right-hand-side independent, so a common iteration budget bounds  $\max_a \varepsilon_{t,a}$ . This zero-start form is what the fixed-budget certificate (85) uses. For a *warm start*  $\tilde{\mathbf{u}}_0 \neq \mathbf{0}$ , energy-norm minimization over the shifted Krylov space instead gives

$$\frac{\|\tilde{\mathbf{u}}_{I_{t,a}} - \mathcal{C}_t^{-1} \mathbf{g}\|_{\mathcal{C}_t}}{\|\mathcal{C}_t^{-1} \mathbf{g}\|_{\mathcal{C}_t}} \leq 2q^{I_{t,a}} \frac{\|\tilde{\mathbf{u}}_0 - \mathcal{C}_t^{-1} \mathbf{g}\|_{\mathcal{C}_t}}{\|\mathcal{C}_t^{-1} \mathbf{g}\|_{\mathcal{C}_t}}, \quad q = \frac{\sqrt{\kappa_t} - 1}{\sqrt{\kappa_t} + 1},$$

i.e. the extra factor is the initial relative energy error  $\|e_0\|_{\mathcal{C}_t}/\|\mathcal{C}_t^{-1} \mathbf{g}\|_{\mathcal{C}_t}$  (equal to one at  $\tilde{\mathbf{u}}_0 = \mathbf{0}$ , recovering the zero-start bound). Under the same certified condition-number bound, adaptive residual stopping (line 5) remains valid with a warm start because it certifies the realized residual directly; but the fixed-budget count (85) is valid *as written* only for zero initialization, and a fixed-budget warm-start certificate additionally requires a bound on that initial relative energy error (the zero-start iteration count does not by itself certify a warm-started solve). Any certified upper bound  $\bar{\kappa}_t \geq \kappa_t$  may replace  $\kappa_t$ ; for a nonnegative weighted outer-product operator,  $\lambda_{\min}(\mathcal{C}_t) \geq \lambda$  and  $\lambda_{\max}(\mathcal{C}_t) \leq \lambda + W_t G^2/\sigma^2$  (where  $W_t = \sum_{s \in \mathcal{S}_t} w_s$ ) give

$$\kappa_t \leq \frac{\lambda_{\max}(\mathcal{C}_t)}{\lambda} \leq 1 + \frac{W_t G^2}{\lambda \sigma^2} =: \bar{\kappa}_t,$$

whereas a general SPD operator requires a separately certified  $\bar{\kappa}_t$ .

Two cases at value one must be distinguished. The true condition number  $\kappa_t = 1$  forces all eigenvalues equal, i.e.  $\mathcal{C}_t = c_t \mathbf{I}$  for some scalar  $c_t > 0$ , and one CG iteration from  $\tilde{\mathbf{u}}_0 = \mathbf{0}$  is *exact* ( $\varepsilon_{t,a} = 0$ ) for every  $\mathbf{g} \neq \mathbf{0}$ ; this does *not* by itself imply  $c_t = \lambda$ . With the specific certificate used here,  $\bar{\kappa}_t = 1$  additionally forces  $\lambda_{\max}(\mathcal{C}_t) = \lambda$ , which together with  $\mathcal{C}_t \succeq \lambda \mathbf{I}$  gives  $\mathcal{C}_t = \lambda \mathbf{I}$  exactly. In either reading the single-step solve is exact, so we keep

$\bar{\kappa}_t = 1$  as a separate branch and never evaluate the logarithmic budget below at its zero denominator. For  $\bar{\kappa}_t > 1$ , setting the right-hand side of the rate equal to  $\bar{\varepsilon}_{\text{CG}}$  and requiring  $\bar{\varepsilon}_{\text{CG}} < 1$  (where  $\bar{\varepsilon}_{\text{CG}}$  denotes the round- $t$  target tolerance  $\bar{\varepsilon}_t$ ) gives the per-solve budget

$$I_{\min}(\bar{\kappa}_t, \bar{\varepsilon}_{\text{CG}}) = \left\lceil \frac{\log(2/\bar{\varepsilon}_{\text{CG}})}{\log((\sqrt{\bar{\kappa}_t} + 1)/(\sqrt{\bar{\kappa}_t} - 1))} \right\rceil = O(\sqrt{\bar{\kappa}_t} \log(1/\bar{\varepsilon}_{\text{CG}})), \quad (85)$$

which, applied to every scored action, ensures  $\max_a \varepsilon_{t,a} \leq \bar{\varepsilon}_{\text{CG}} = \bar{\varepsilon}_t$ . For  $W_t = n$  (unrescaled fixed buffer size),  $\bar{\kappa}_t \leq 1 + nG^2/(\lambda\sigma^2)$  is bounded independently of  $T$ , so  $I_{\min}$  is constant. For full history, and for rescaled uniform subsampling with  $w_s = (t-1)/n$ , we have  $W_t = t-1$ , so the condition-number *bound*  $\bar{\kappa}_t = O(tG^2/(\lambda\sigma^2))$  grows with  $t$ , and the worst-case certificate is then obtained with  $I_{\min} = \Theta(G\sqrt{t/(\lambda\sigma^2)} \log(1/\bar{\varepsilon}_{\text{CG}}))$  iterations per solve—i.e.  $I = O(\sqrt{t} \log(1/\bar{\varepsilon}_{\text{CG}}))$  suffices under this bound (this is an upper-bound certificate, not a lower bound on the number of iterations CG actually needs). CG may converge faster than this worst-case certificate suggests under favorable spectral clustering of  $\mathcal{C}_t$ ; the realized iteration counts would be measured empirically.

## F Two-Sided Spectral-Distortion Corollary

The main-text Theorem 29 uses only a one-sided upper transfer factor. For comparison we record the earlier two-sided result, which instead assumes both an upper *and* a lower spectral factor and yields the familiar  $\sqrt{\rho_+/\rho_-}$  inflation. The two results trade different assumptions—the corollary below needs a lower spectral bound; the dynamic Theorem 1 needs the dynamic width complexity  $\Lambda_T^{\mathcal{C}}$  instead—and neither numerically dominates the other in general.

**Corollary 34** (Two-sided spectral-distortion inflation). *Adopt the hypotheses of Theorem 1 with  $\mathcal{C}_t = \hat{\mathcal{C}}_t$ , retaining the original nonlinear center (so  $\omega_t = \bar{\beta}_t + \bar{\psi}_t$ ; in the convex last-layer regime under exact optimization,  $\bar{\psi}_t = 0$  and  $\omega_t = \bar{\beta}_t$ ), and a uniform CG tolerance  $\bar{\varepsilon}_t \leq \bar{\varepsilon}$  so that  $\alpha_t \leq \alpha_I := \sqrt{(1+\bar{\varepsilon})/(1-\bar{\varepsilon})}$ . Suppose predictable factors  $\rho_{-,t}, \rho_{+,t}$  are certified before action selection and the two-sided sandwich  $\rho_{-,t}\bar{\mathbf{C}}_t \preceq \hat{\mathbf{C}}_t \preceq \rho_{+,t}\bar{\mathbf{C}}_t$  holds for all  $t \leq T$  with  $0 < \rho_{-,t} \leq \rho_{+,t} < \infty$ ; set  $\rho_- = \min_t \rho_{-,t}$ ,  $\rho_+ = \max_t \rho_{+,t}$ , and (to satisfy the convention  $u_t \geq 1$  of Assumption 5) take the transfer factor  $u_t := \max\{1, \rho_{+,t}\} \leq \max\{1, \rho_+\}$ . Write  $\omega_{\max,T} := \max_{t \leq T}(\bar{\beta}_t + \bar{\psi}_t)$  and  $\rho_+^* := \max\{1, \rho_+\}$ . Then, on the intersection of the confidence event  $\mathcal{E}_{\text{conf}}$  ( $\geq 1 - \delta$ ) and the certificate event  $\mathcal{E}_{\text{cert}}$  (which now also contains the two-sided sandwich event;  $\geq 1 - \delta_{\text{cert}}$ ),*

$$R_T \leq 2\alpha_I \sqrt{\rho_+^*/\rho_-} \omega_{\max,T} \sqrt{(\sigma^2 + G^2/\lambda)T\gamma_T} + 2E_T$$

with probability at least  $1 - \delta - \delta_{\text{cert}}$ , where  $\gamma_T$  is the frozen-feature information gain (28). Under the standard normalization  $\rho_+ \geq 1$  (achievable by rescaling the sandwich so its upper factor is at least one),  $\rho_+^* = \rho_+$  and the inflation is the familiar  $\sqrt{\rho_+/\rho_-}$ .

*Proof.* This is a *direct* argument through the monotone frozen Gram  $\bar{\mathbf{C}}_t$ ; it does *not* substitute  $\hat{s}_t$  for the dynamic width in Corollary 33. *Step 1 (variance sandwich).* Inverting the two-sided sandwich gives  $\rho_{+,t}^{-1}\bar{\mathbf{C}}_t^{-1} \preceq \hat{\mathbf{C}}_t^{-1} \preceq \rho_{-,t}^{-1}\bar{\mathbf{C}}_t^{-1}$ , so, evaluating the quadratic form at  $\mathbf{g}_t(a)$ ,  $\rho_{+,t}^{-1/2}\bar{s}_t(a) \leq \hat{s}_t(a) \leq \rho_{-,t}^{-1/2}\bar{s}_t(a)$ . The lower half gives  $\bar{s}_t^2(a) \leq \rho_{+,t}\hat{s}_t^2(a) \leq u_t\hat{s}_t^2(a)$ , so Assumption 5 holds with  $u_t = \max\{1, \rho_{+,t}\}$ . *Step 2 (per-round regret).* With  $s_{\mathcal{C},t} = \hat{s}_t$ , Theorem 1 Steps 1–4 give  $\text{regret}_t \leq 2\alpha_t \omega_t \sqrt{u_t} \hat{s}_t(a_t) + 2\varepsilon_{\text{lin}}(t)$ ; using  $\hat{s}_t(a_t) \leq \rho_{-,t}^{-1/2}\bar{s}_t(a_t)$ ,  $u_t = \max\{1, \rho_{+,t}\} \leq \rho_+^*$ ,  $\alpha_t \leq \alpha_I$ , and  $\omega_t \leq \omega_{\max,T}$ ,

$$\text{regret}_t \leq 2\alpha_I \omega_{\max,T} \sqrt{\rho_+^*/\rho_-} \bar{s}_t(a_t) + 2\varepsilon_{\text{lin}}(t) \leq 2\alpha_I \omega_{\max,T} \sqrt{\rho_+^*/\rho_-} \bar{s}_t(a_t) + 2\varepsilon_{\text{lin}}(t).$$

*Step 3 (frozen-feature elliptic potential).* Sum over  $t$  and apply Cauchy–Schwarz, then the classical elliptic-potential identity for the *monotone* Gram  $\bar{\mathbf{C}}_t$  ((7)),  $\sum_t \bar{s}_t^2(a_t) \leq (\sigma^2 + G^2/\lambda) \sum_t \log(1 + \sigma^{-2}\bar{s}_t^2(a_t)) = (\sigma^2 + G^2/\lambda) \gamma_T$ :

$$R_T \leq 2\alpha_I \omega_{\max,T} \sqrt{\rho_+^*/\rho_-} \sqrt{T \sum_t \bar{s}_t^2(a_t)} + 2E_T \leq 2\alpha_I \omega_{\max,T} \sqrt{\rho_+^*/\rho_-} \sqrt{(\sigma^2 + G^2/\lambda)T\gamma_T} + 2E_T.$$

The probability statement is the union bound over  $\mathcal{E}_{\text{conf}}$  and  $\mathcal{E}_{\text{cert}}$ .  $\square$

For Lemma 18, allocate  $\delta_{\text{cert}}/T$  per round and use the worst-case  $m = T-1$  in (55) to obtain the required time-uniform sandwich.

---

## G GGN Approximation Error

**Lemma 35** (GGN approximation error). *Suppose  $\mu_{\theta}(\mathbf{x}, a)$  is twice differentiable in  $\theta$ . Suppose  $|r_s - \mu_{\theta_t}(\mathbf{x}_s, a_s)| \leq \epsilon_r$  and  $\|\nabla^2 \mu_{\theta_t}(\mathbf{x}_s, a_s)\|_{\text{op}} \leq H$  for all  $s < t$ . Define  $\zeta := (t-1) \epsilon_r H / (\sigma^2 \lambda)$ . Then:*

- (i)  $\|\mathbf{H}_t - \mathbf{C}_t\|_{\text{op}} \leq (t-1) \epsilon_r H / \sigma^2 = \zeta \lambda$ .
- (ii) If  $\zeta < 1$ , then  $(1-\zeta)\mathbf{C}_t \preceq \mathbf{H}_t \preceq (1+\zeta)\mathbf{C}_t$ .

*Proof.* For  $D_t = \sigma^{-2} \sum_{s < t} (r_s - \mu_{\theta_t}) \nabla^2 \mu_{\theta_t}$ , the triangle inequality gives  $\|D_t\|_{\text{op}} \leq \zeta \lambda$ , proving (i). Combining  $-\|D_t\|\mathbf{I} \preceq -D_t \preceq \|D_t\|\mathbf{I}$  with  $\mathbf{C}_t \succeq \lambda \mathbf{I}$  proves (ii).  $\square$

## H Exponential-Family Curvature Definitions

For a regular exponential-family negative log-likelihood with natural parameter  $z_{\theta}(\mathbf{x}, a) \in \mathbb{R}^k$  (scalar for Bernoulli,  $k$ -dimensional for categorical) and log-partition function  $A$ , the GGN/Fisher curvature matrices are

$$\mathbf{C}_t^{\text{exp}} = \lambda \mathbf{I} + \sum_{s=1}^{t-1} J_t(\mathbf{x}_s, a_s)^{\top} \nabla^2 A(z_t(\mathbf{x}_s, a_s)) J_t(\mathbf{x}_s, a_s), \quad (86)$$

$$\bar{\mathbf{C}}_t^{\text{exp}} = \lambda \mathbf{I} + \sum_{s=1}^{t-1} J_s(\mathbf{x}_s, a_s)^{\top} \nabla^2 A(z_s(\mathbf{x}_s, a_s)) J_s(\mathbf{x}_s, a_s), \quad (87)$$

where  $J_t(\mathbf{x}, a) = \nabla_{\theta} z_{\theta_t}(\mathbf{x}, a) \in \mathbb{R}^{k \times d}$  is the Jacobian of the natural parameter and  $\nabla^2 A(z)$  is the output-space Hessian of the log-partition function. For Bernoulli rewards  $\nabla^2 A(z) = p(1-p)$  (scalar); for categorical rewards  $\nabla^2 A(z) = \text{Diag}(\mathbf{p}) - \mathbf{p}\mathbf{p}^{\top}$ . Setting  $k = 1$  and  $\nabla^2 A = \sigma^{-2}$  recovers the Gaussian curvature (5)–(6).

## I Broader Impact

Exploration can expose users to under-vetted actions and amplify feedback loops. Deployment requires safety constraints, action filters, offline evaluation, and online monitoring; the regret certificates here do not establish safety.

## J Data and Code Availability

The repository contains the Buck pipeline, tests, artifact scripts, validated legacy derived results, the retained  $d = 128, r = 4$  slice, and deterministic reproduction entry points for the new spectral-tail, premise-clean scaling, and real-autodiff studies. Their full evaluation grids are not committed and have not all run. The 2.1 GB legacy scaling tree and historical inputs outside Git are absent, so not every raw-derived chain is reproducible. Neither anonymous tier is built; the review builder lacks `experiments/data/sklearn/covertime/samples_py3` and `experiments/data/sklearn/covertime/targets_py3`. Once restored, the builder sanitizes identities and validates hashes/provenance. Covertime is CC BY 4.0 (DOI 10.24432/C50K5N) (Blackard, 1998). The autodiff study status is `development_pilot`; no evaluation result is claimed. The balanced-MNIST study is `smoke.validated`: canonical archives pass their published checksums and the real-data smoke pipeline completes, but the full grid has not run and the ignored local cache is not a clean-checkout fixture. No current bundle or identifying URL is claimed.